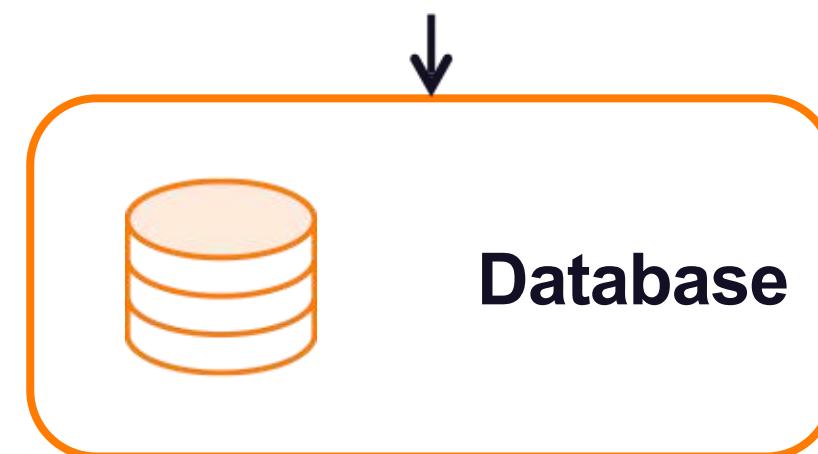
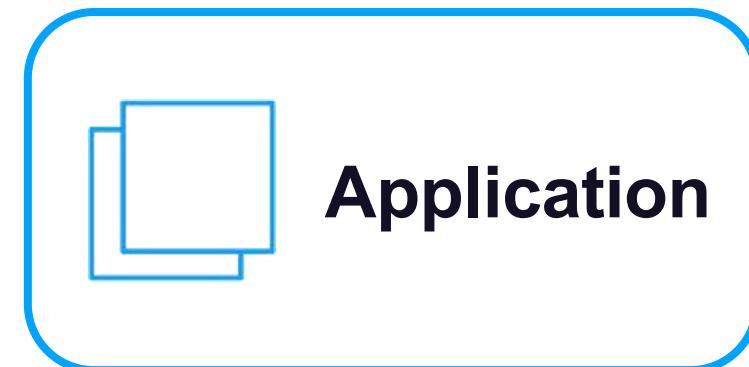
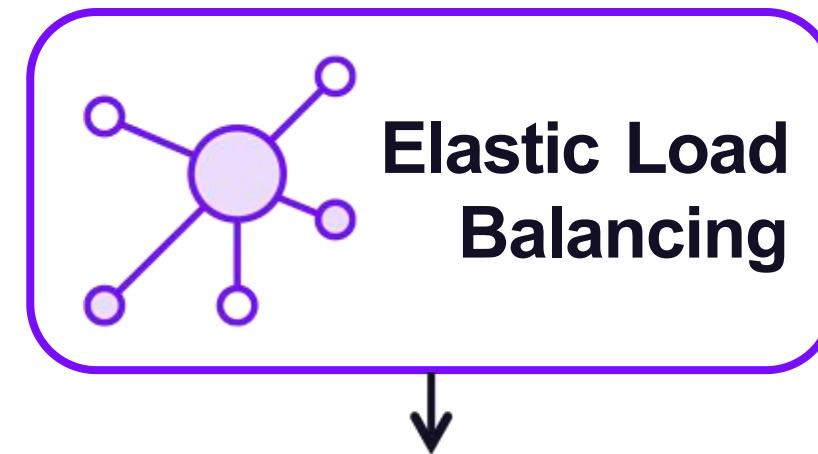
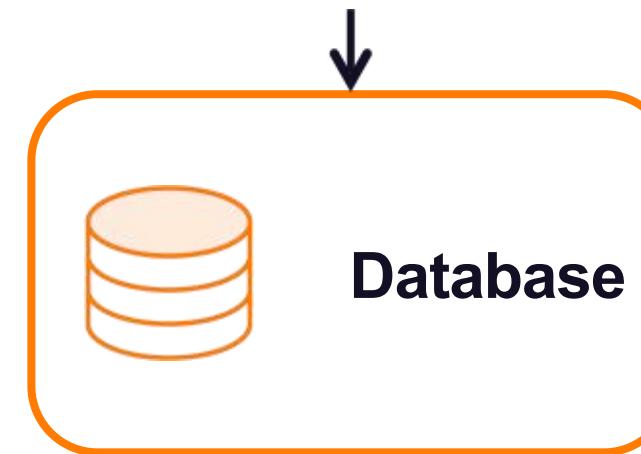
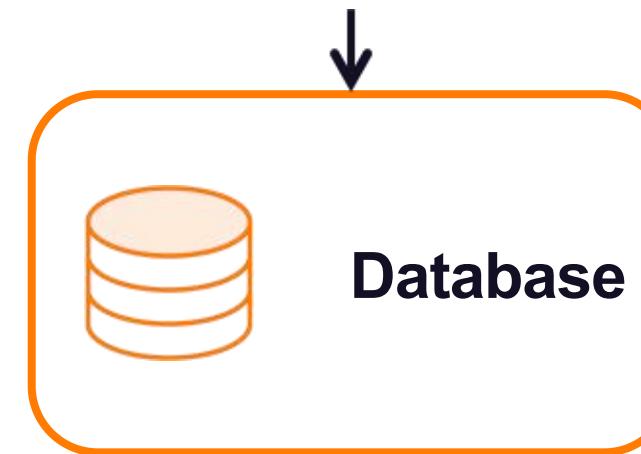
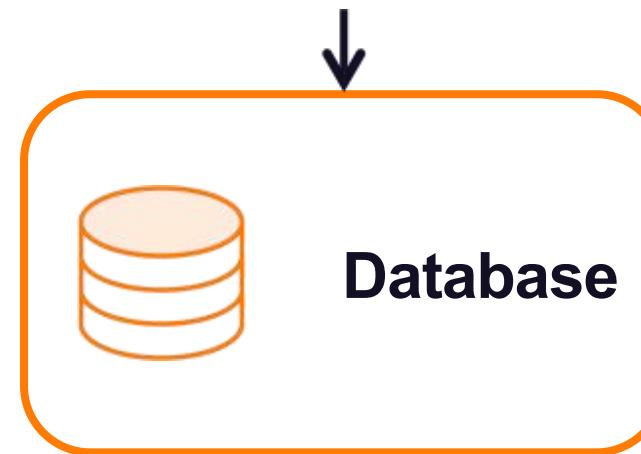
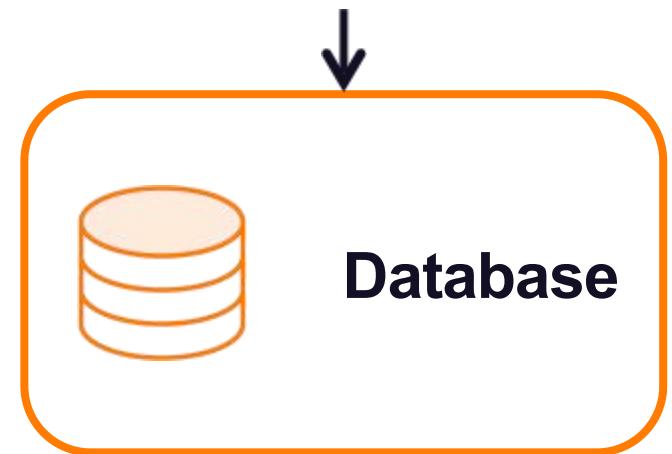
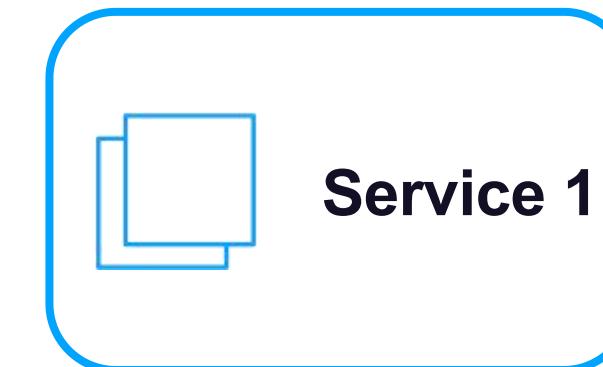
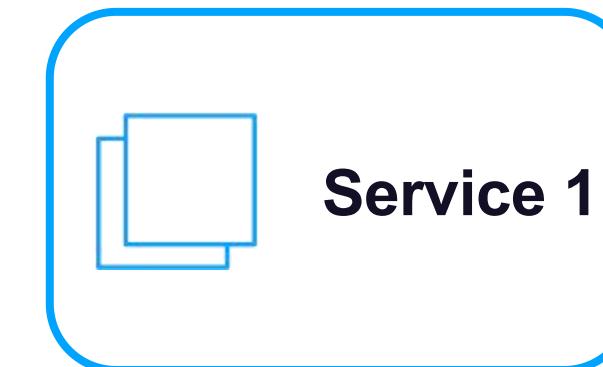
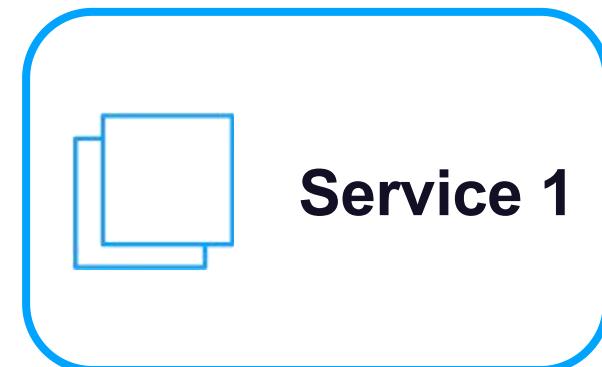
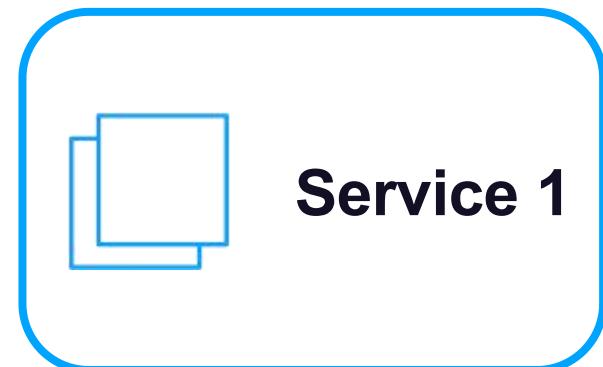
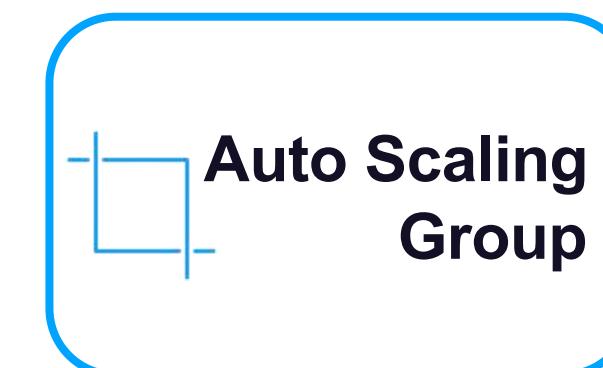
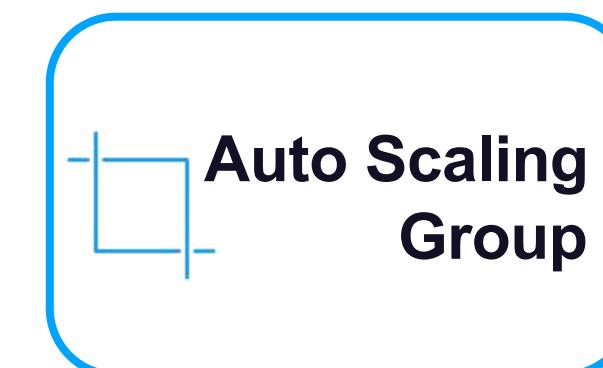
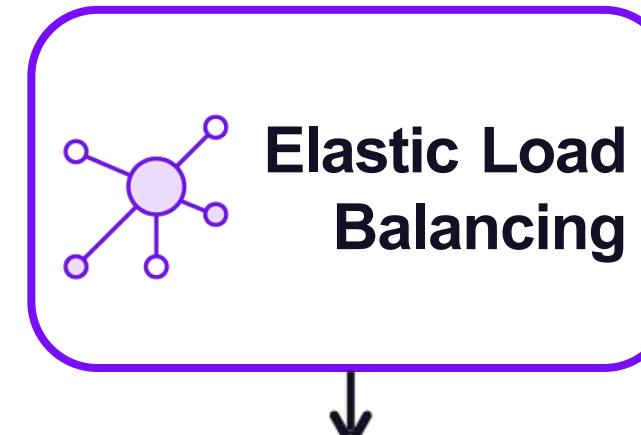
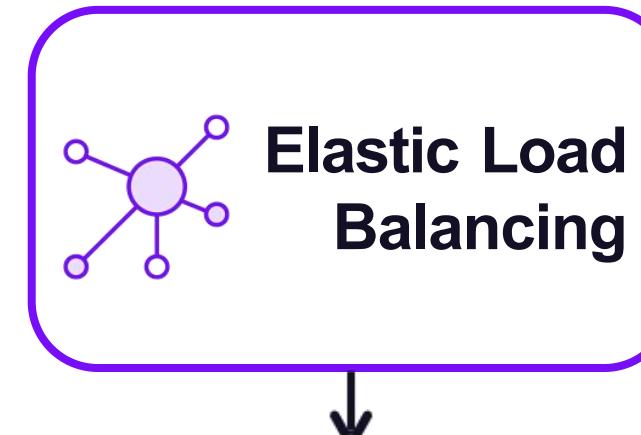
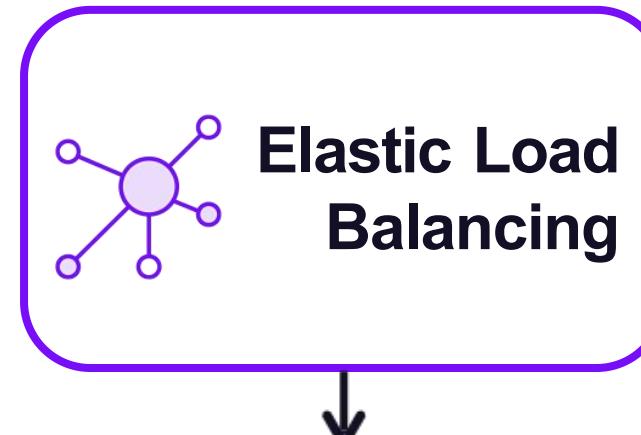
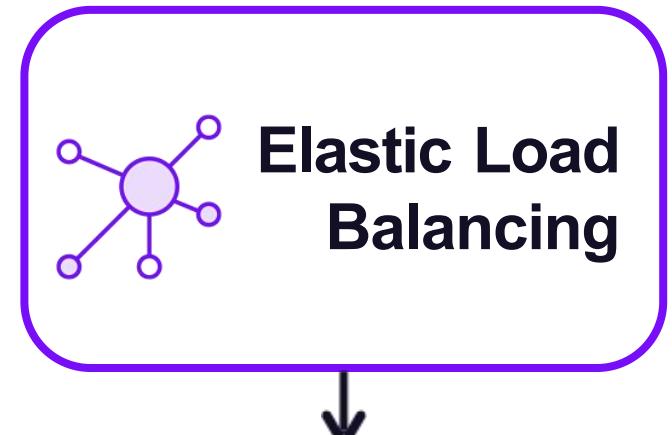


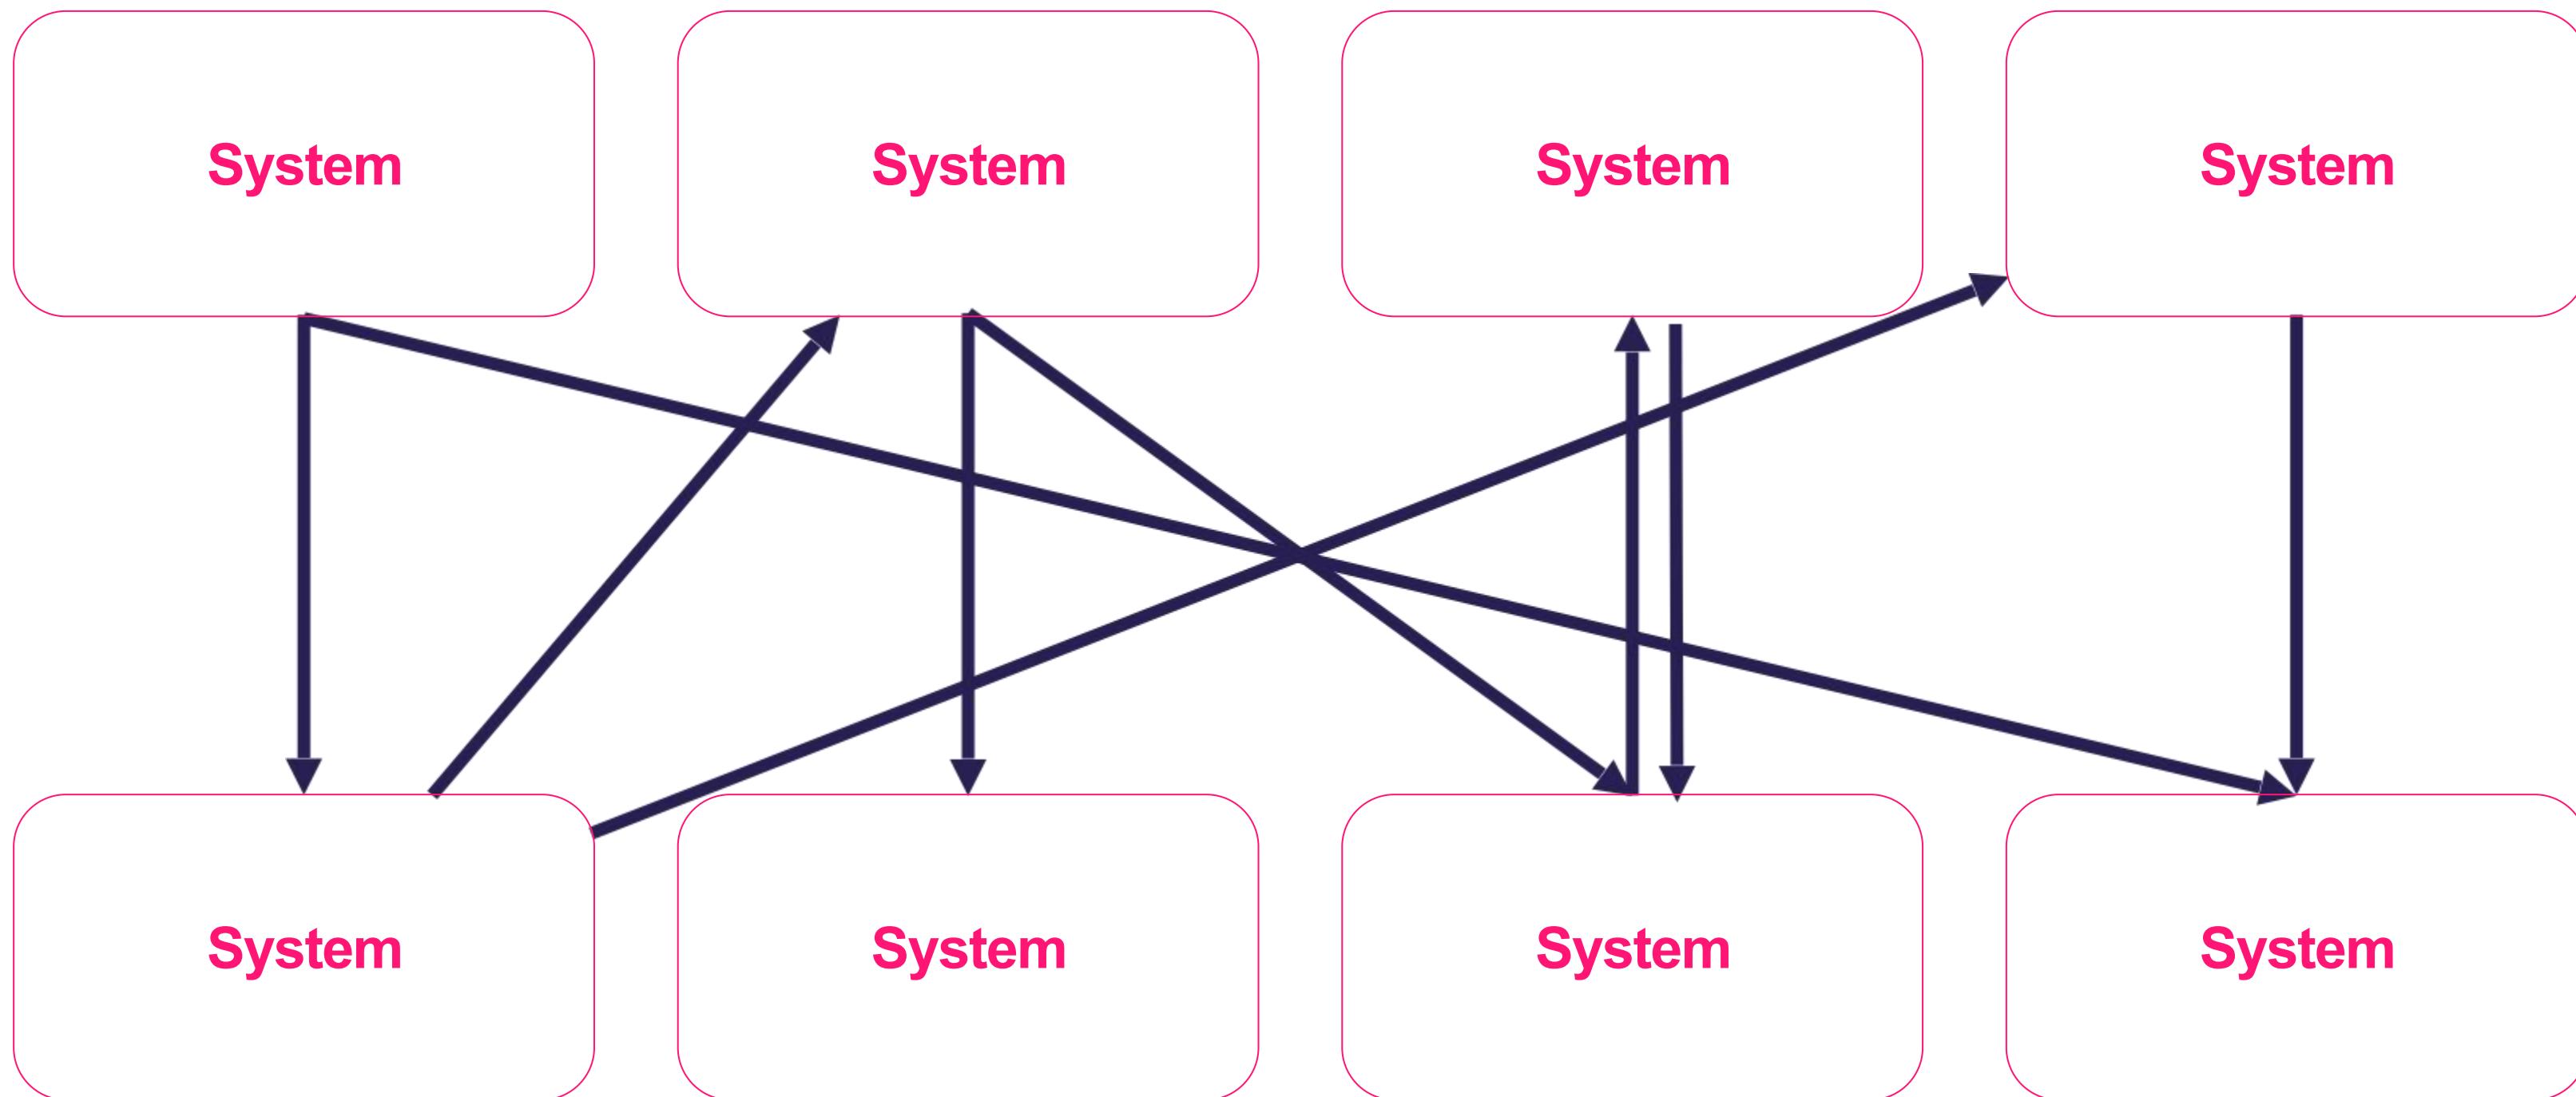
Introducing Pub Sub Systems

Meeting Pub-Sub Systems





Entangled System





Items Service

UI Service

We Call:



The entity/app creating a message, a publisher or producer



The entity/app consuming messages from a channel as a consumer



The system where the channels live and handle this requests as an Event Bus or, more recently, streaming platforms

The channel where messages flow as channel or topic

Back to Definitions



We say the pub-sub system is reliable when you ensure there is no message loss



Has at most once processing when you ensure there is no message duplication



And has exactly-once processing when you only process a message once ensuring it wasn't lost. Of course, the holy grail.

We Had Ton of Other Pub/Sub Systems Before:

RabbitMQ

Mulesoft

Redis

Amazon SQS

Amazon SNS

Azure HDInsight, Azure
Streaming Service, and
Google Pub/Sub

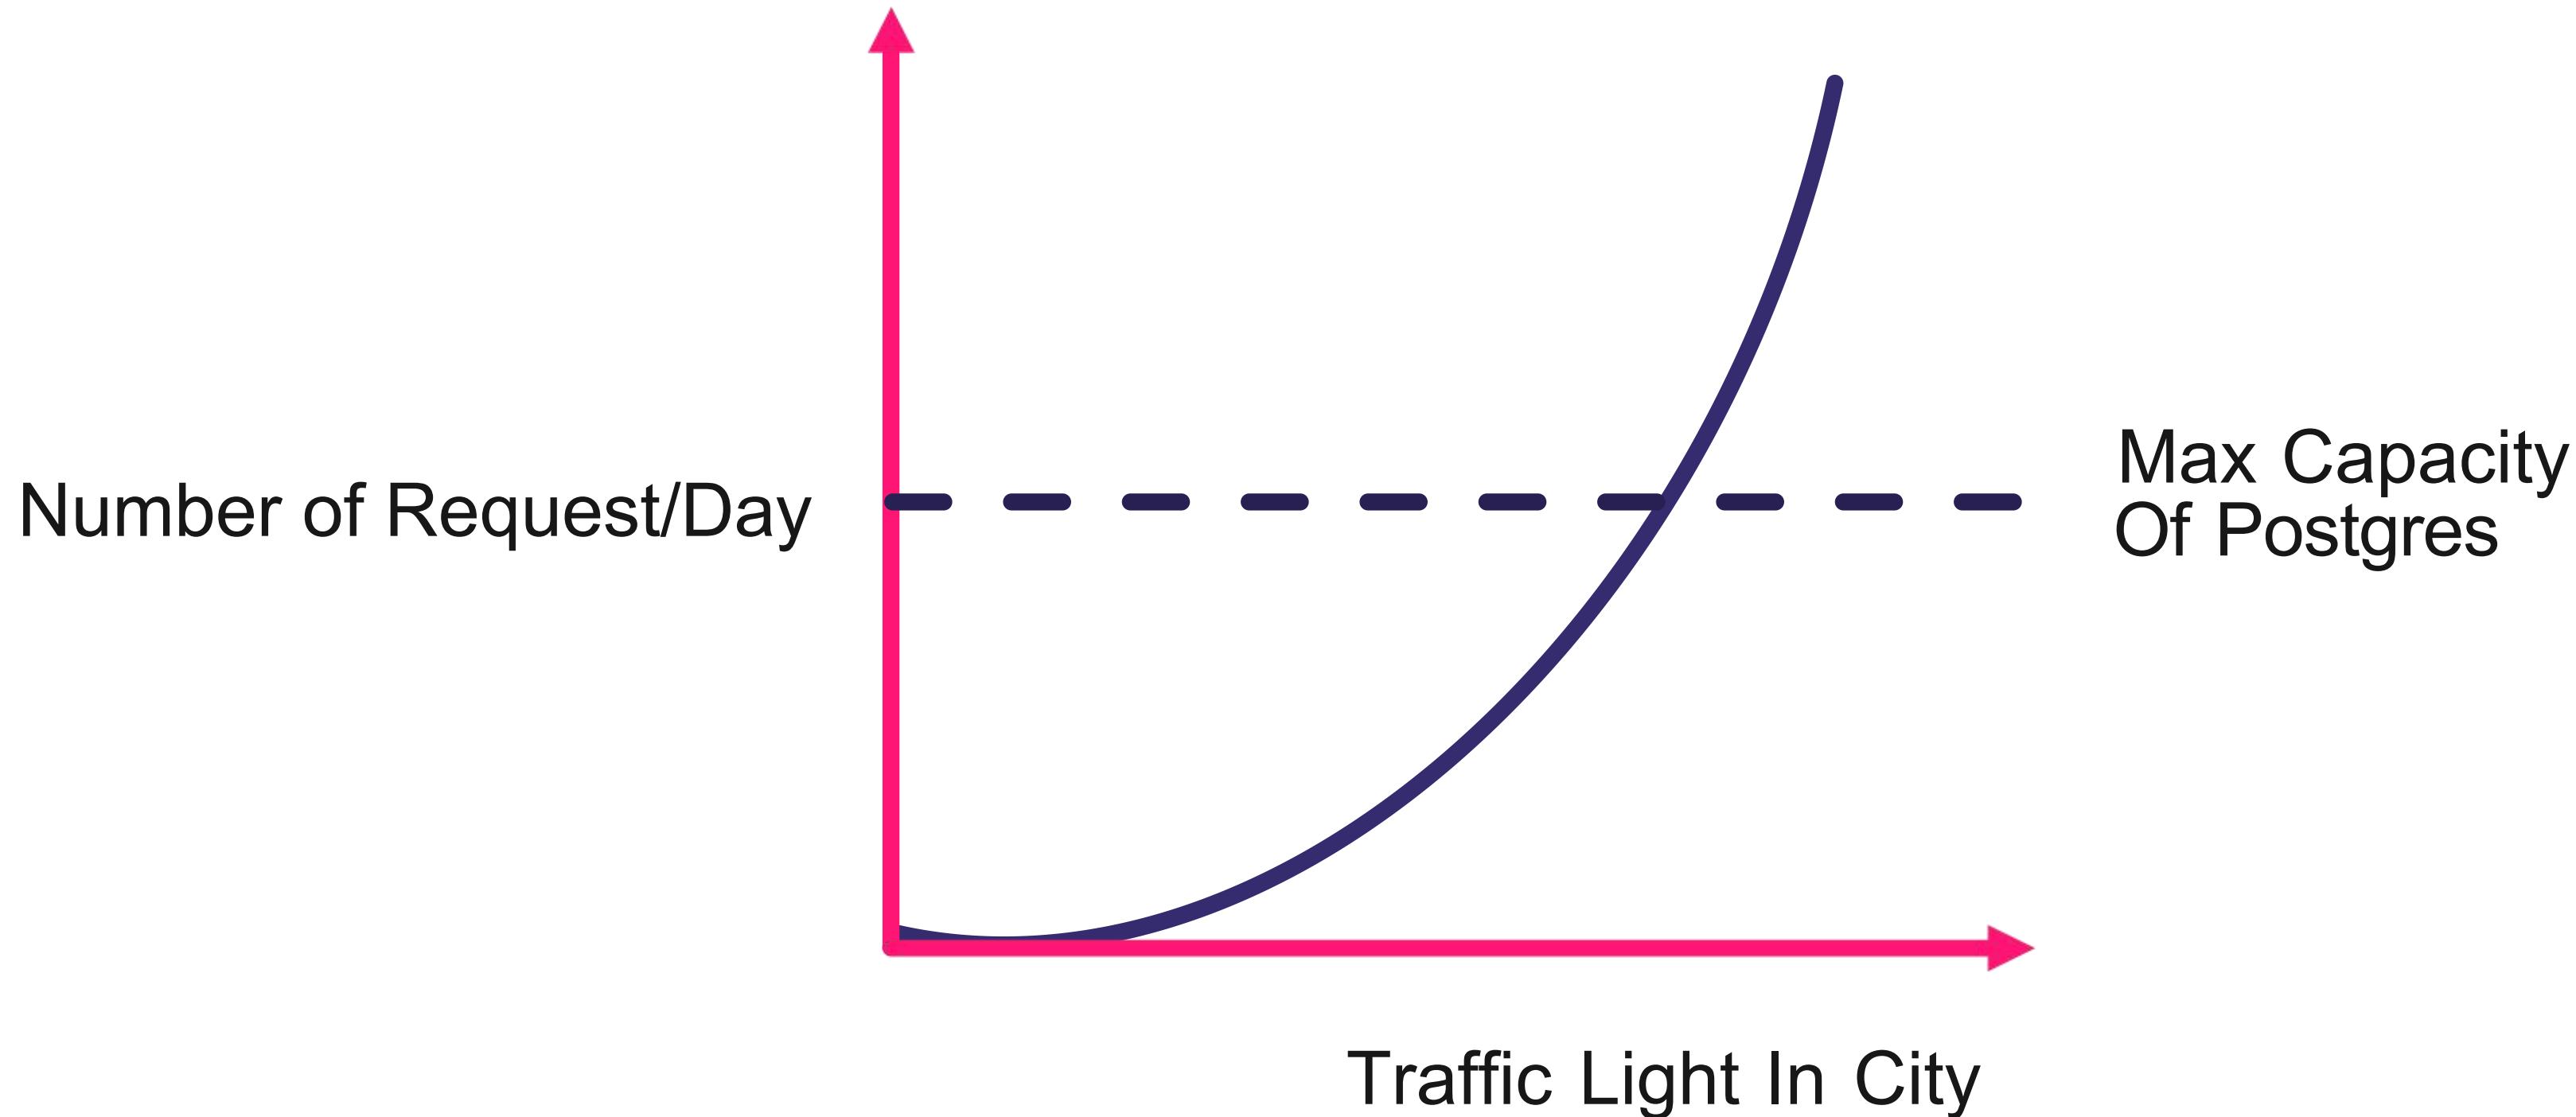


Why Kafka?

A Case Study: Traffic Lights



A Case Study: Traffic Lights



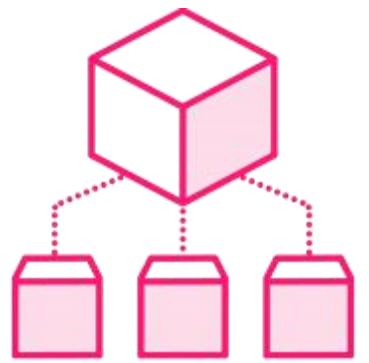
A Case Study: Traffic Lights



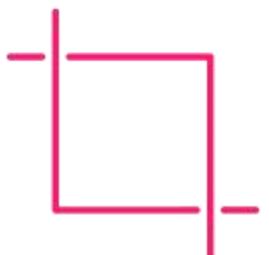
A Case Study: 911 Systems

Kafka comes to solve this!

What Is Kafka?



Event Streaming Platform



Scalable



High Throughput

Difference between Kafka and the Rest of Solutions



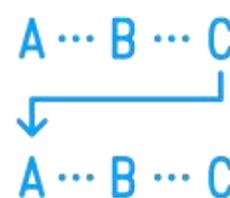
RabbitMQ: As Kafka uses topics in an unidirectional log without handshake, the **difference in performance is brutal, easily 1000x**



Mulesoft: Kafka **makes it simple** to create a channel of communication simply by producing a message the topic will be autocreated



Redis: Redis **does not handle streaming data** nor has a history past messages



Amazon SQS: It based on queues, way better tan RabbitMQ but still there is an **order of magnitude in throughput and speed between Kafka and SQS**. Plus SQS is not exactly cheap

Kafka Architecture

Kafka Streams

**Perform
administrative
tasks**

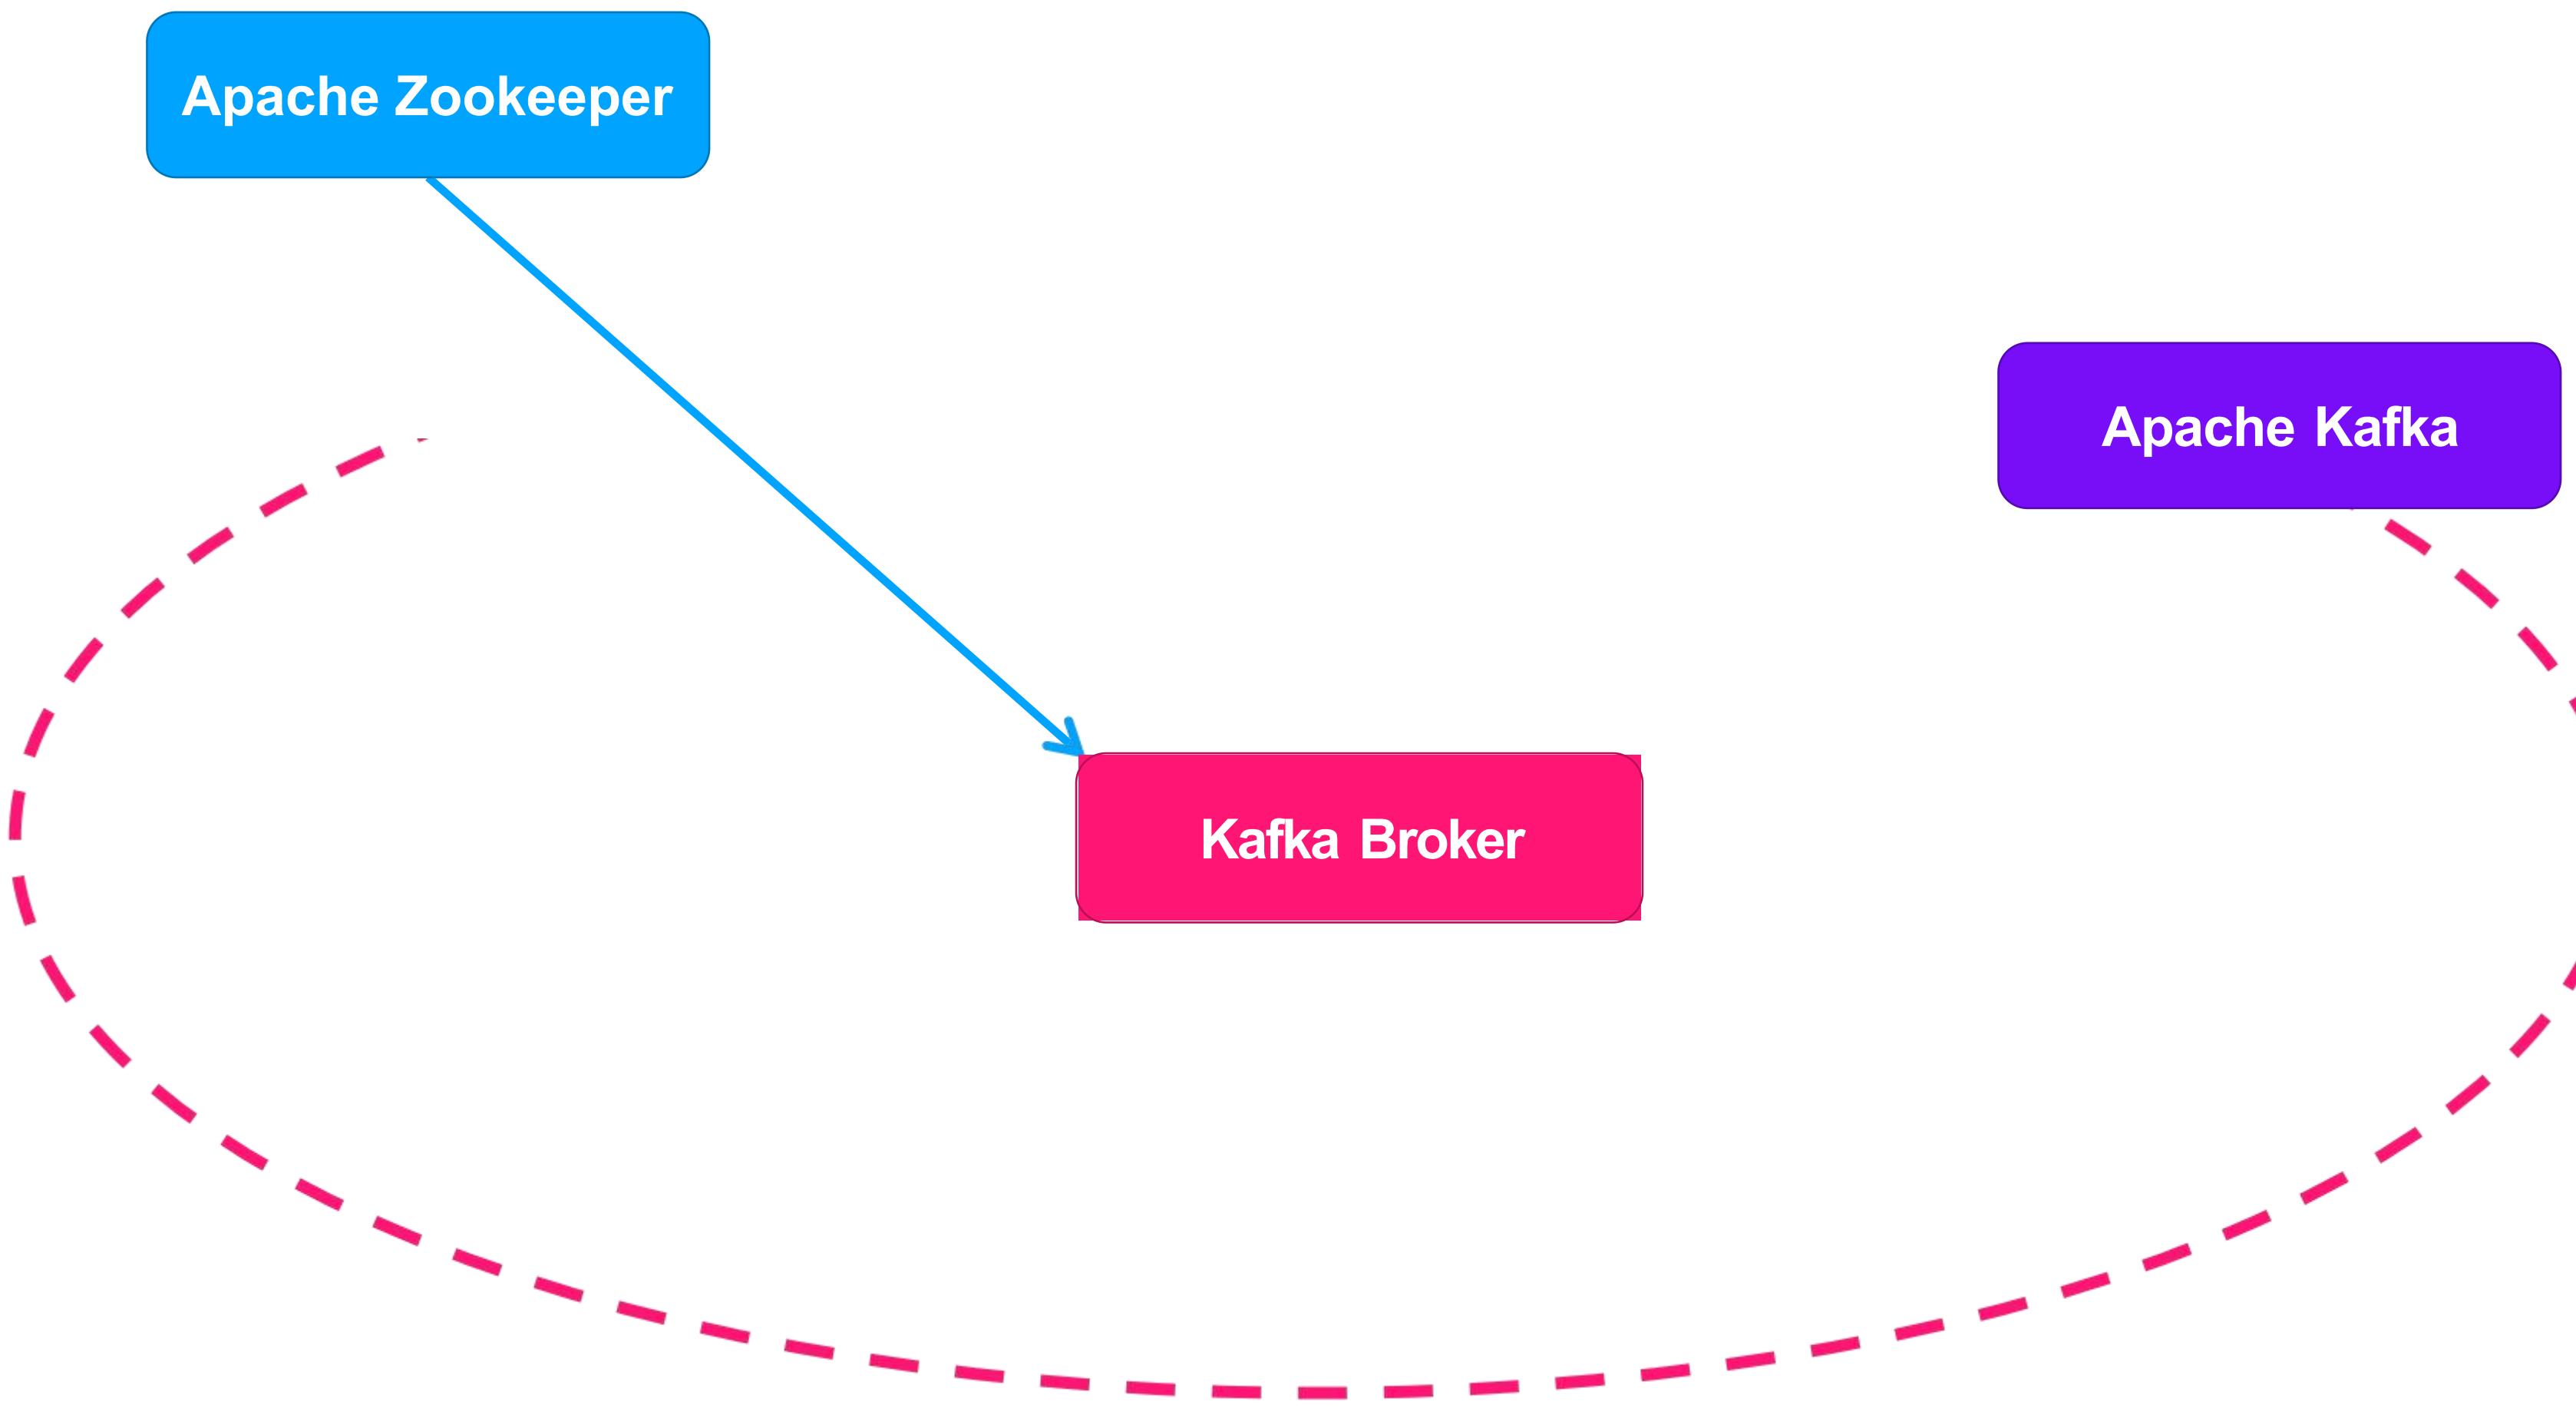
**Creating and Configuring
Producers and Consumers**

Kafka Connect

Meet Kafka

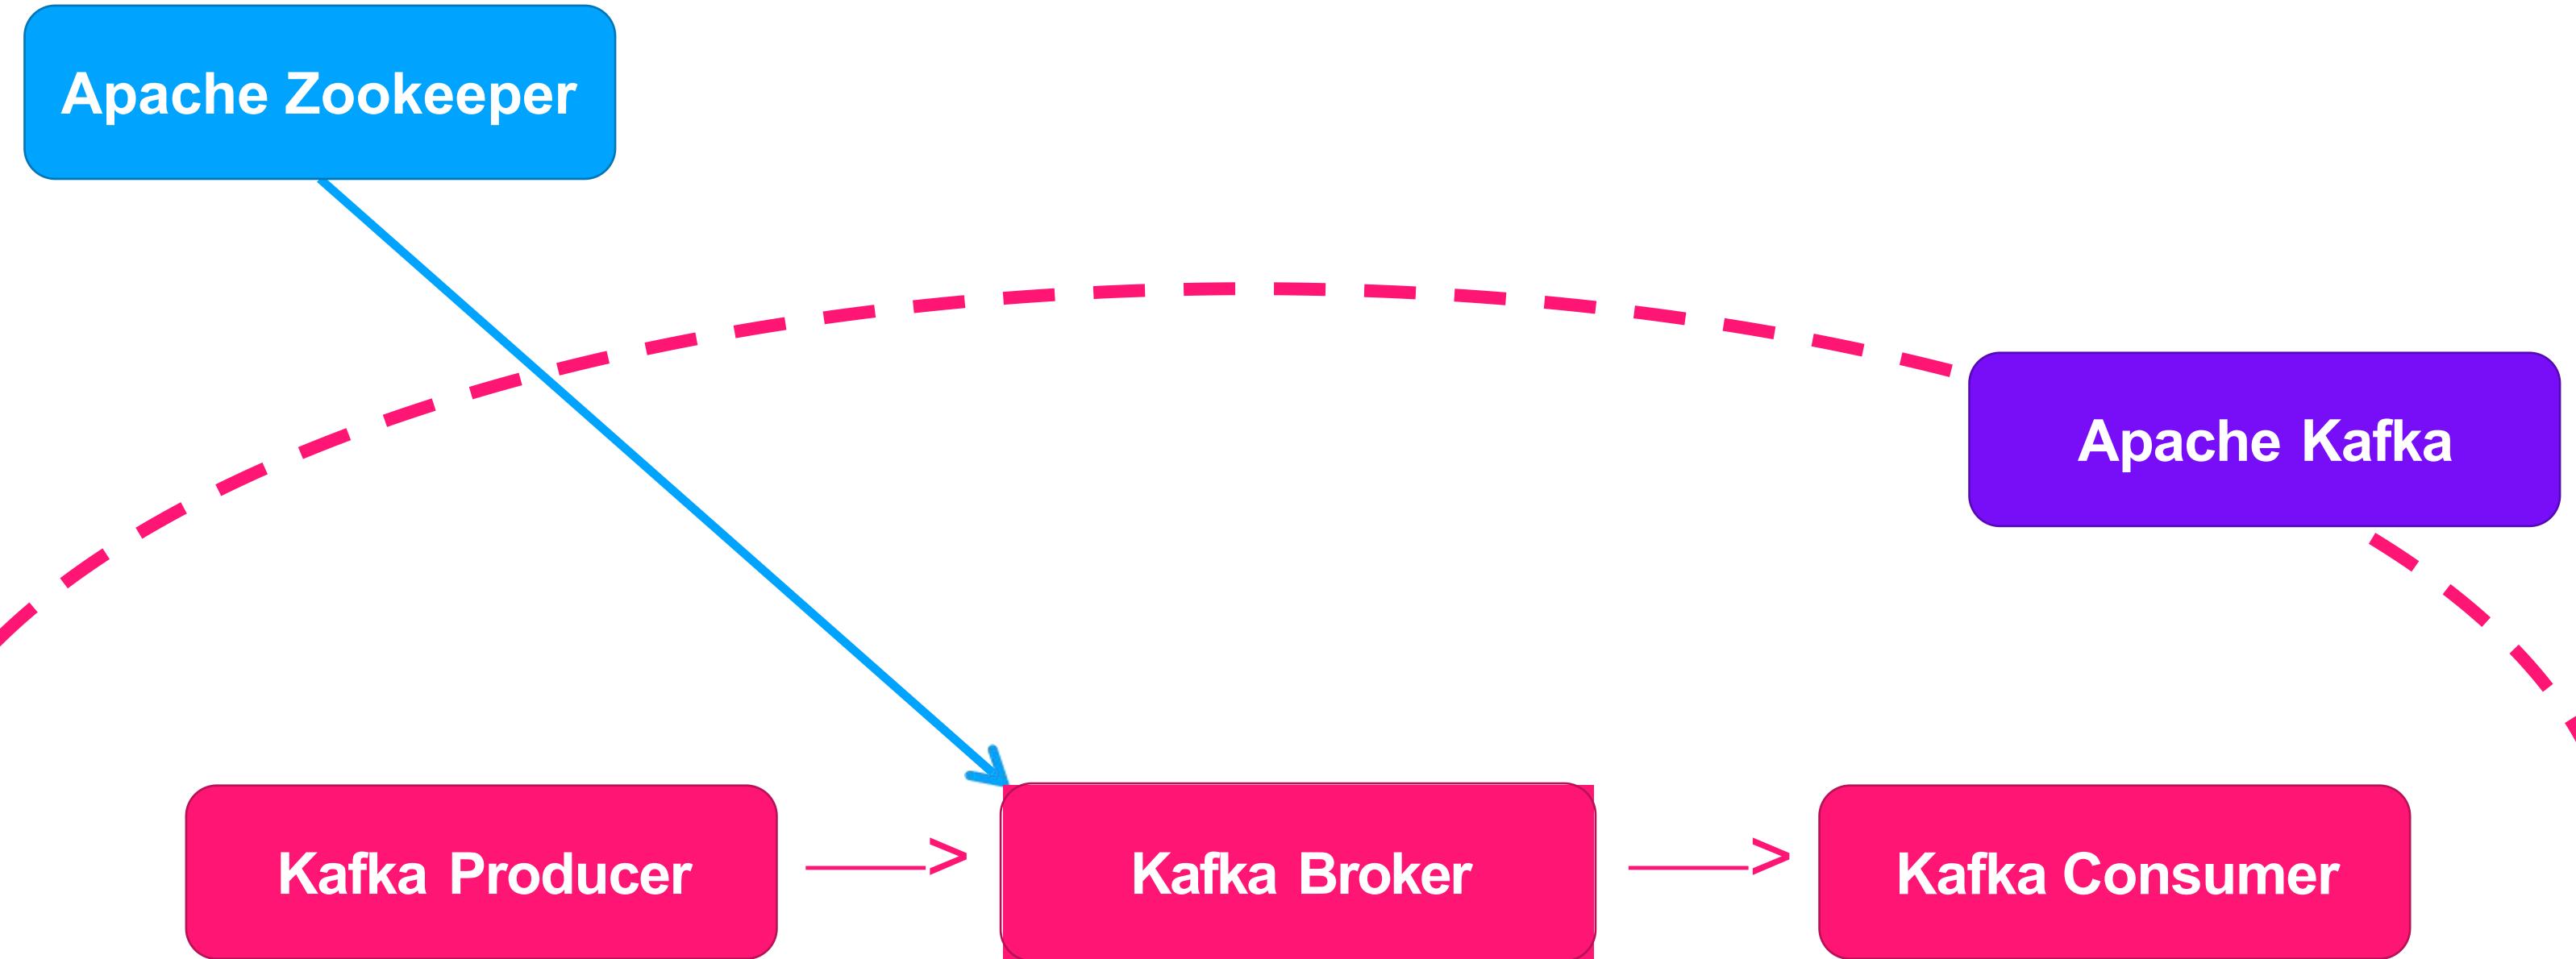
Apache Kafka

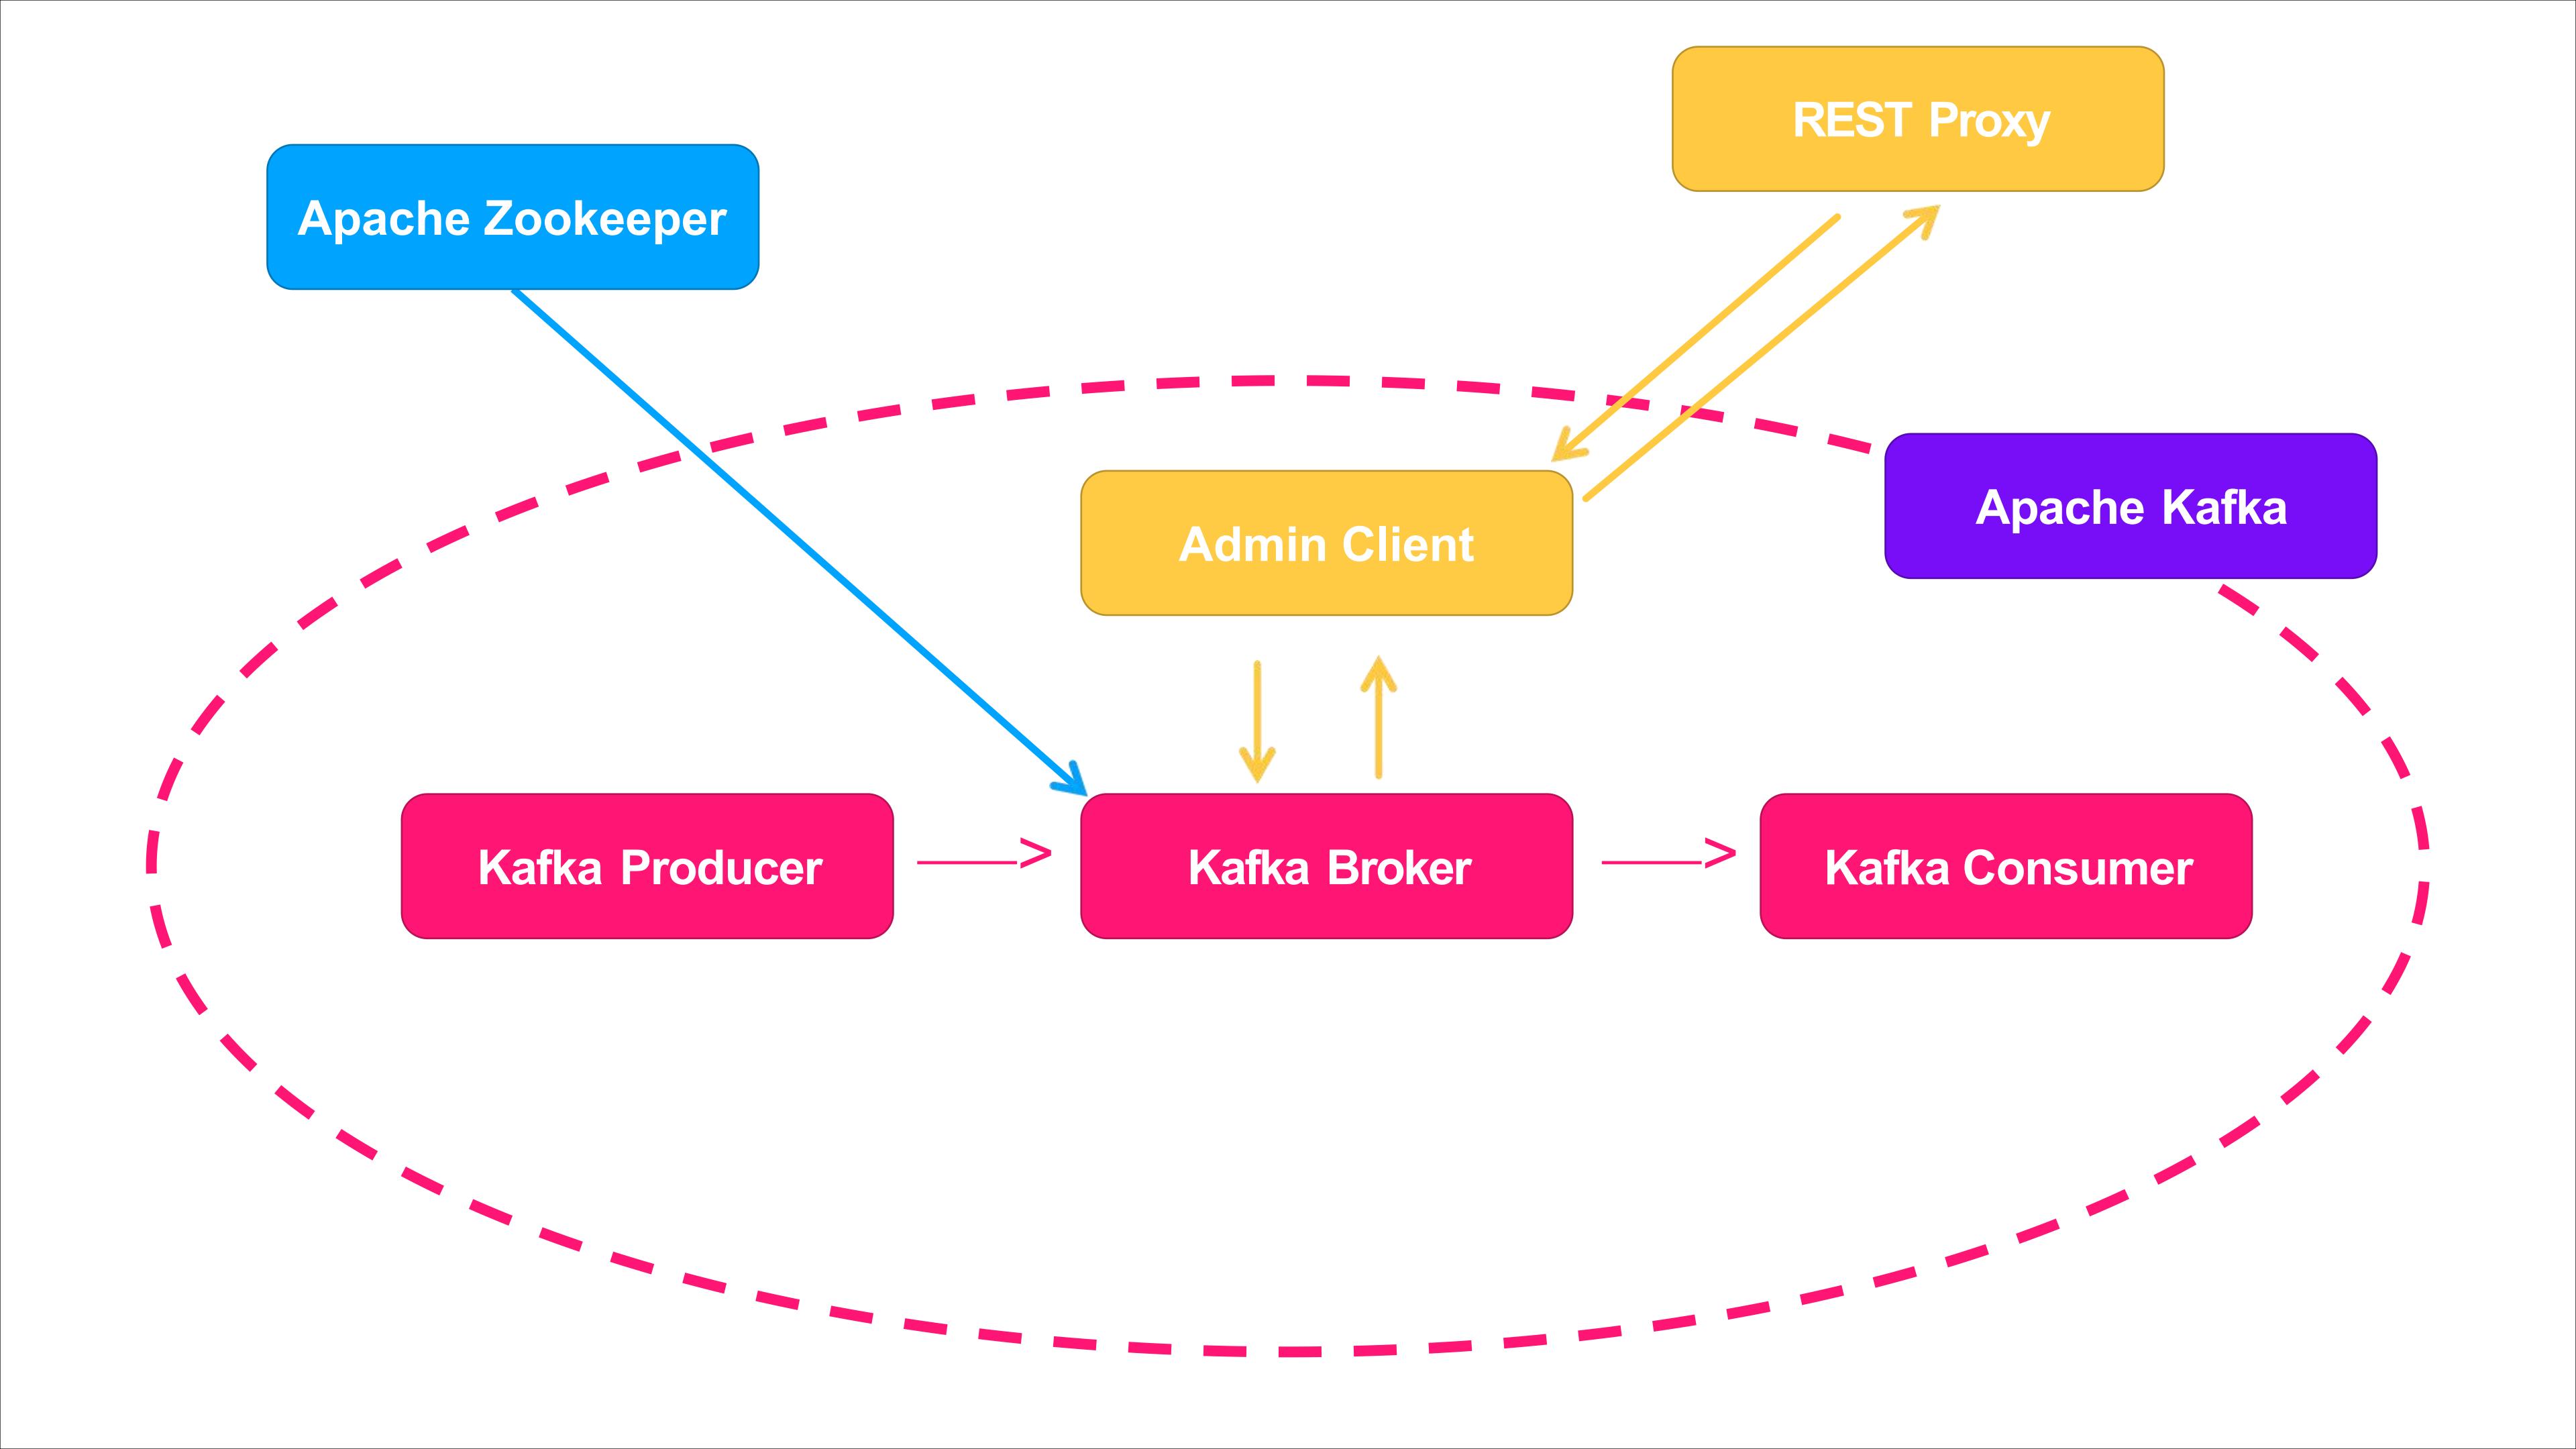
Kafka Broker

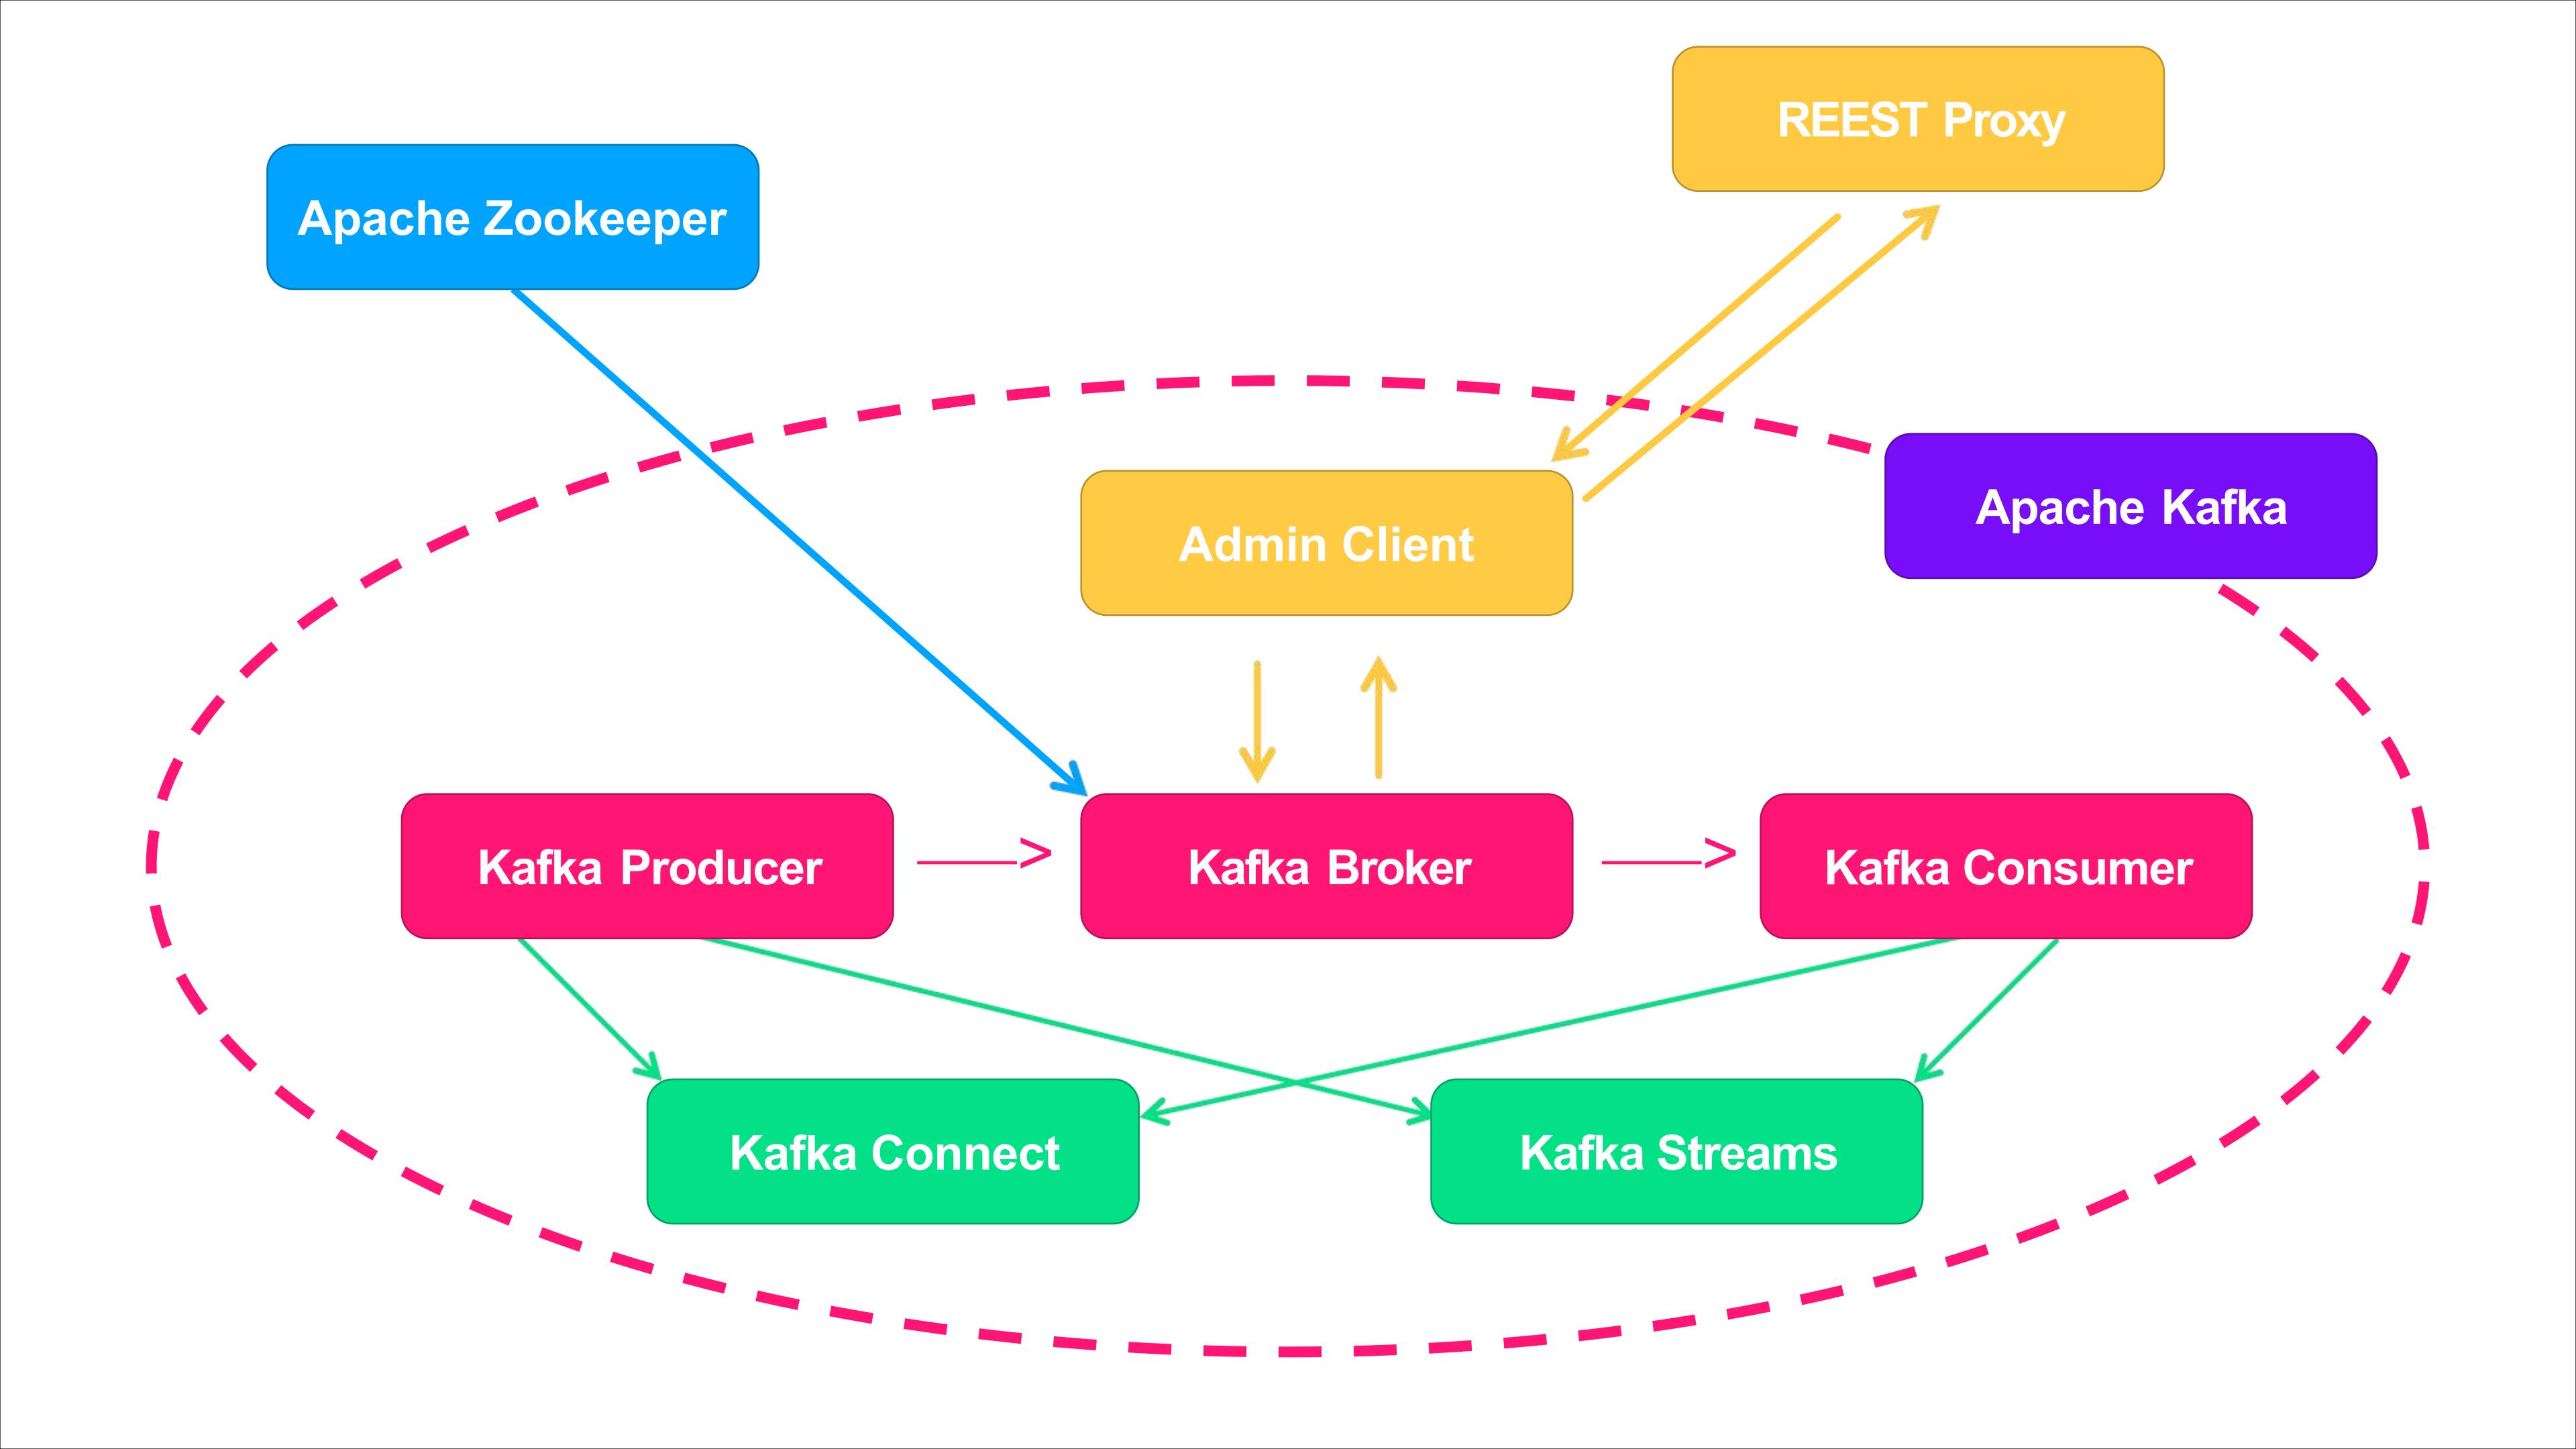


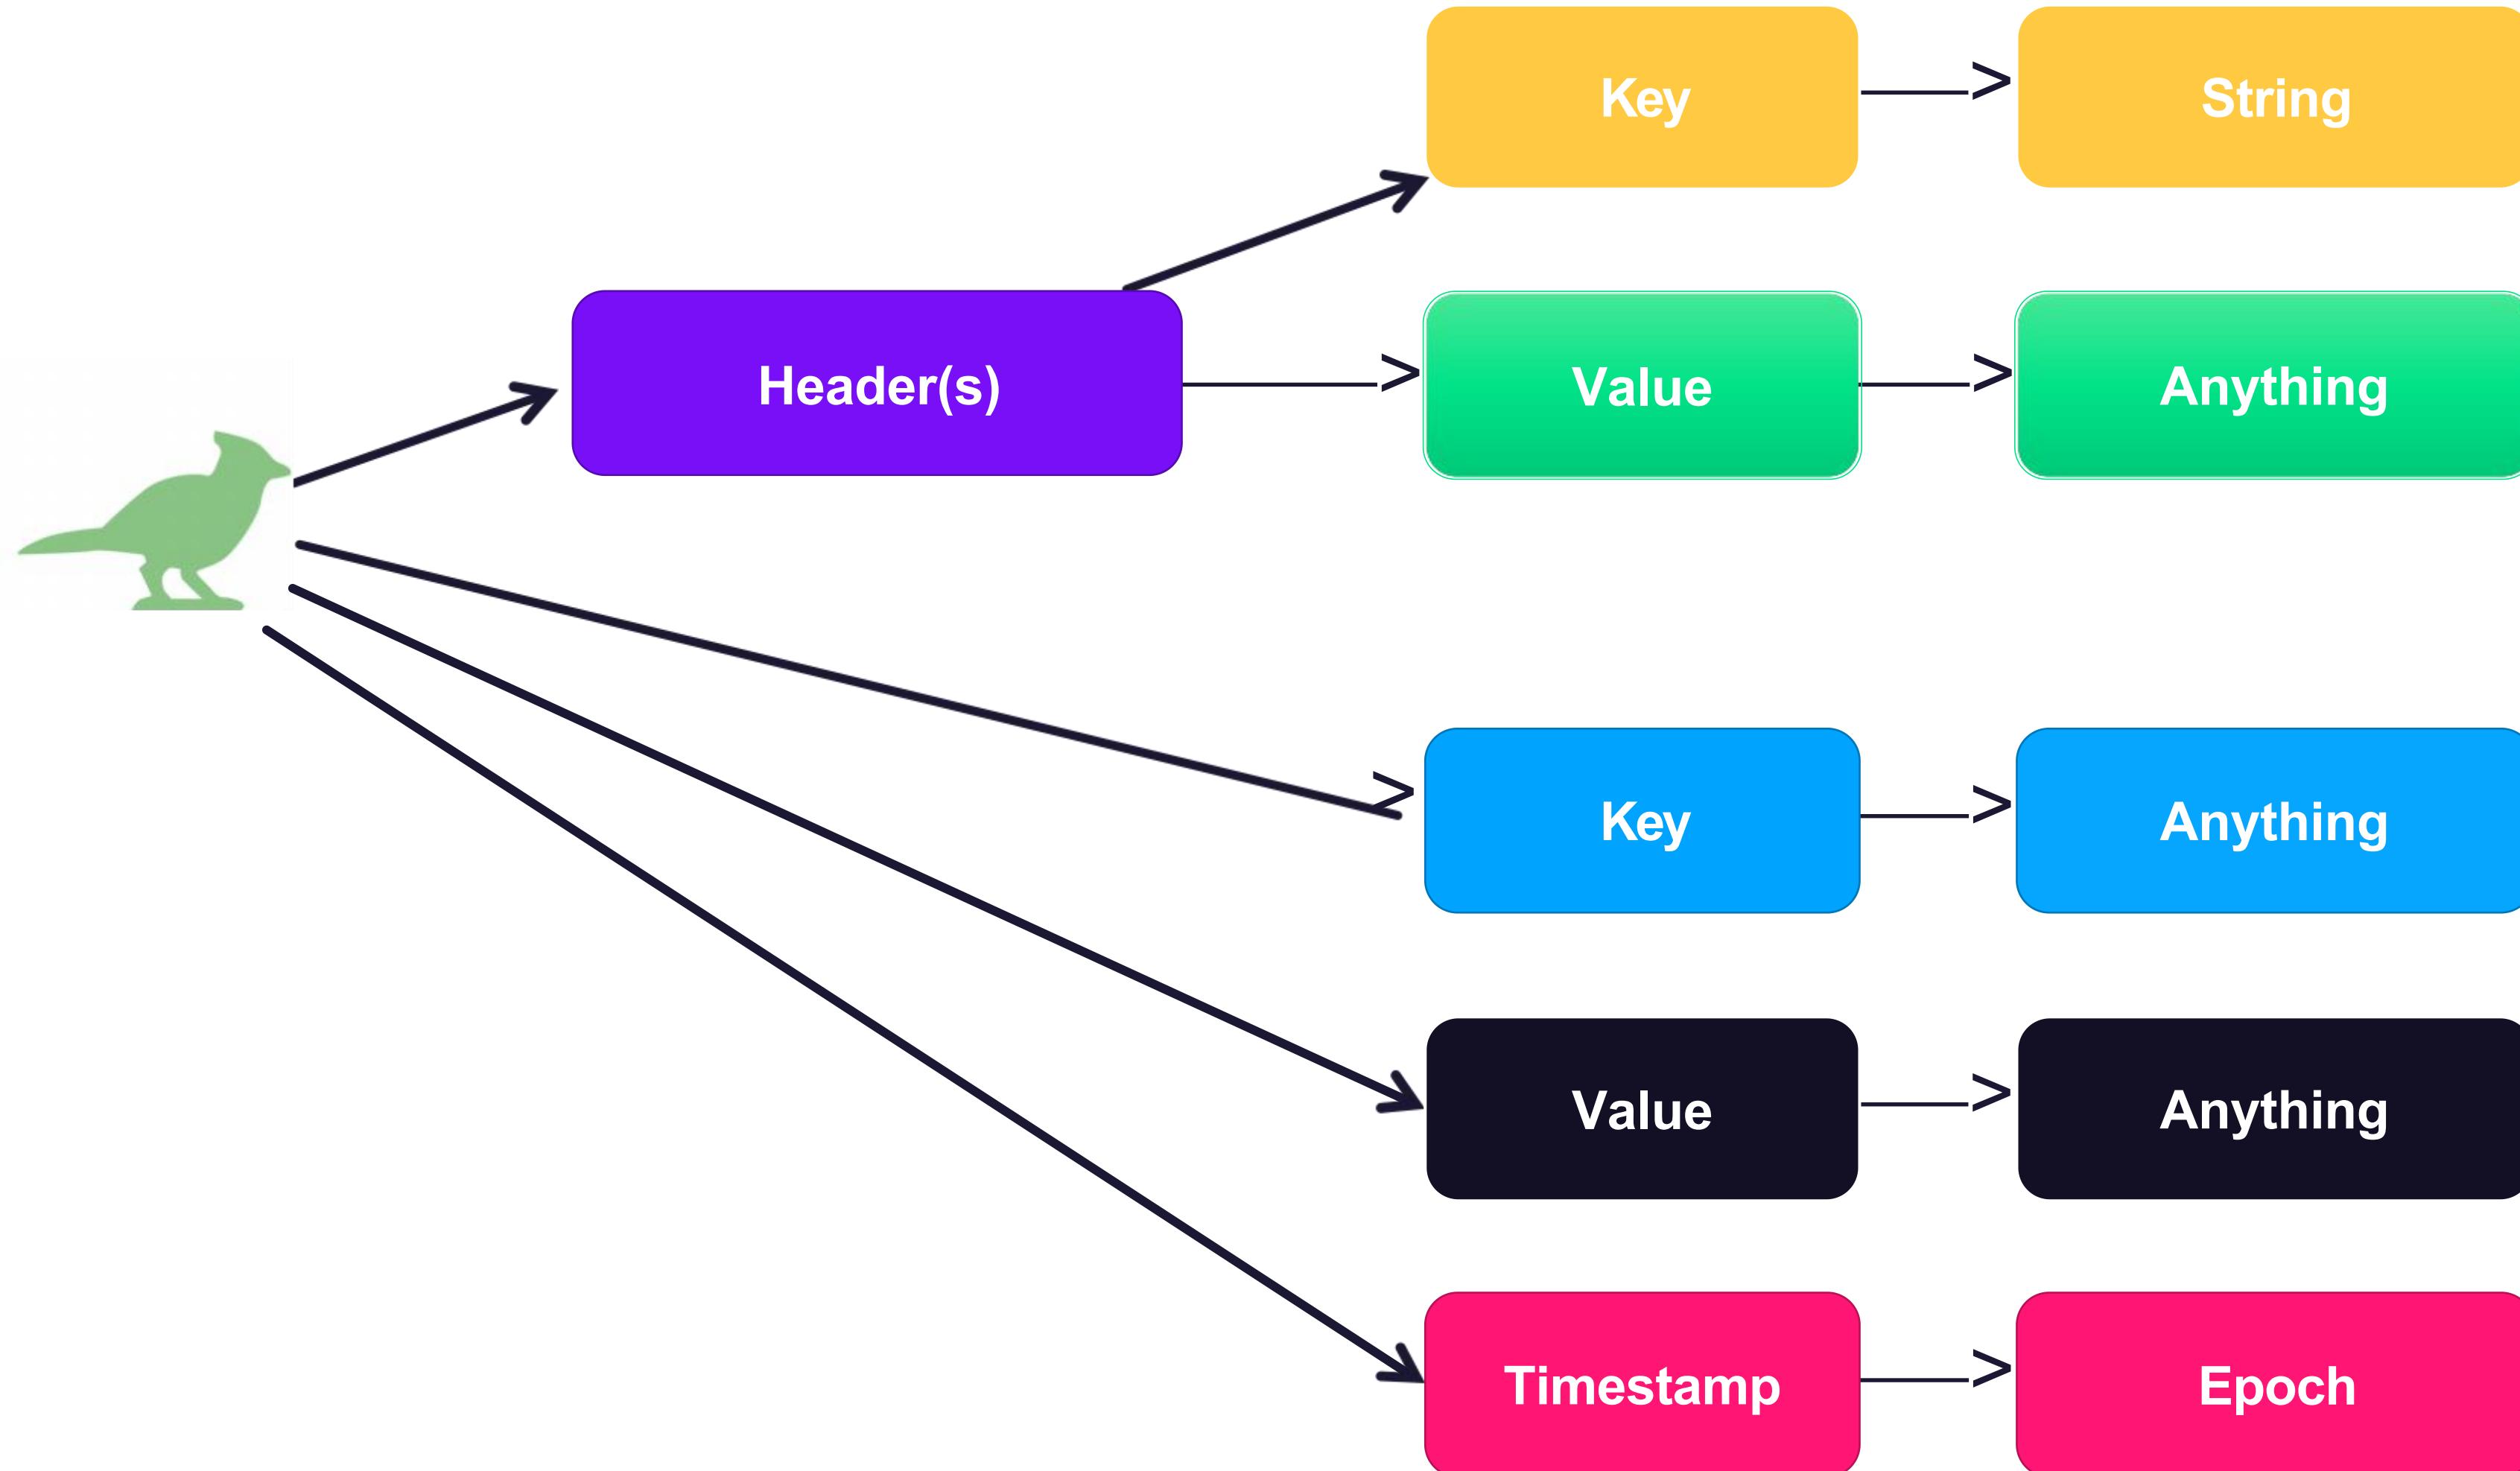


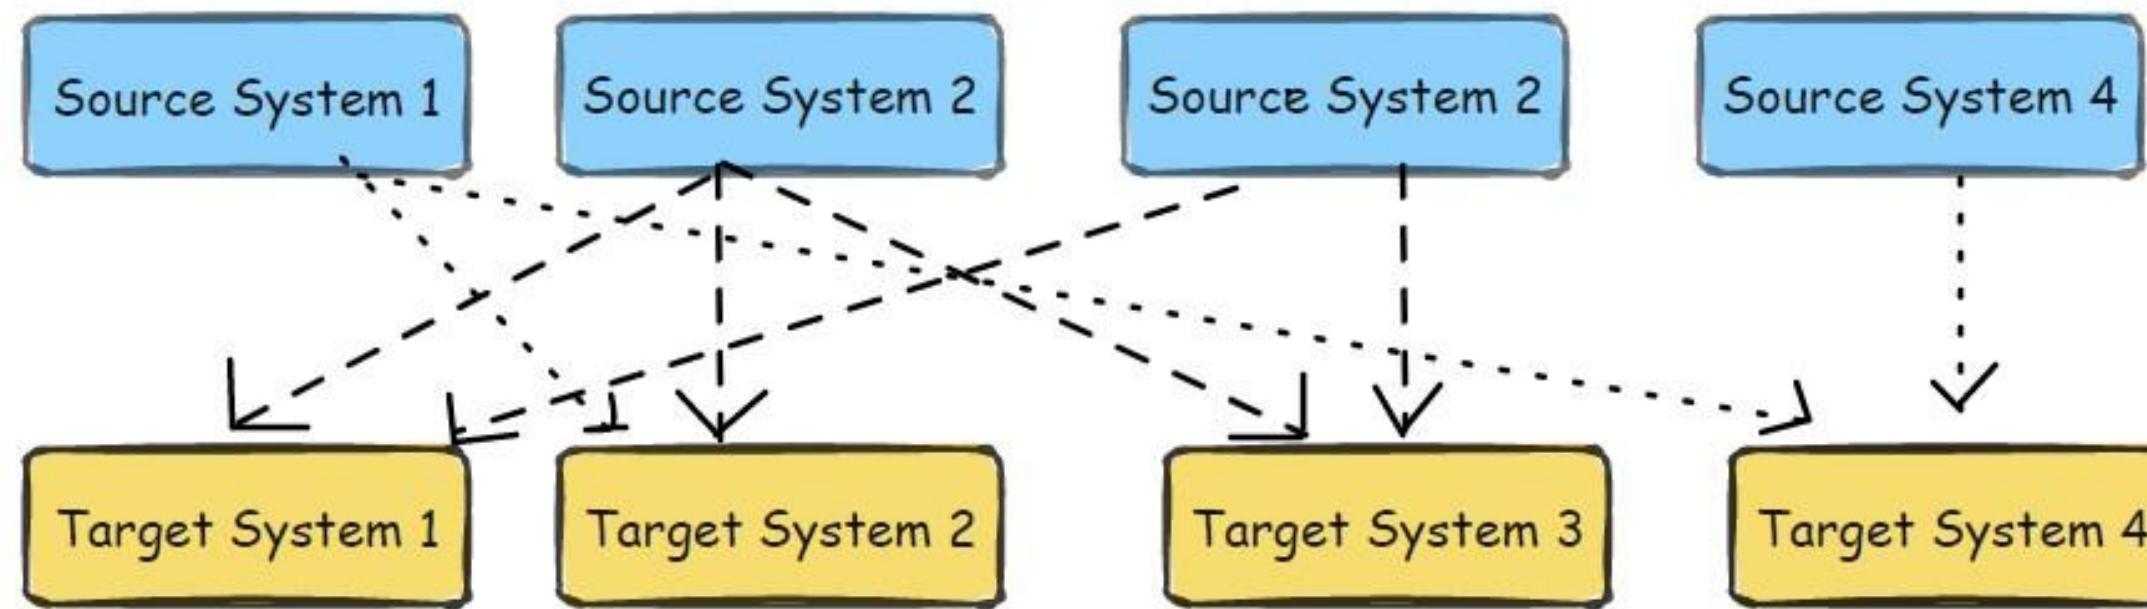
- Service Discovery
- Service Address Retrieval
- Management of Elections











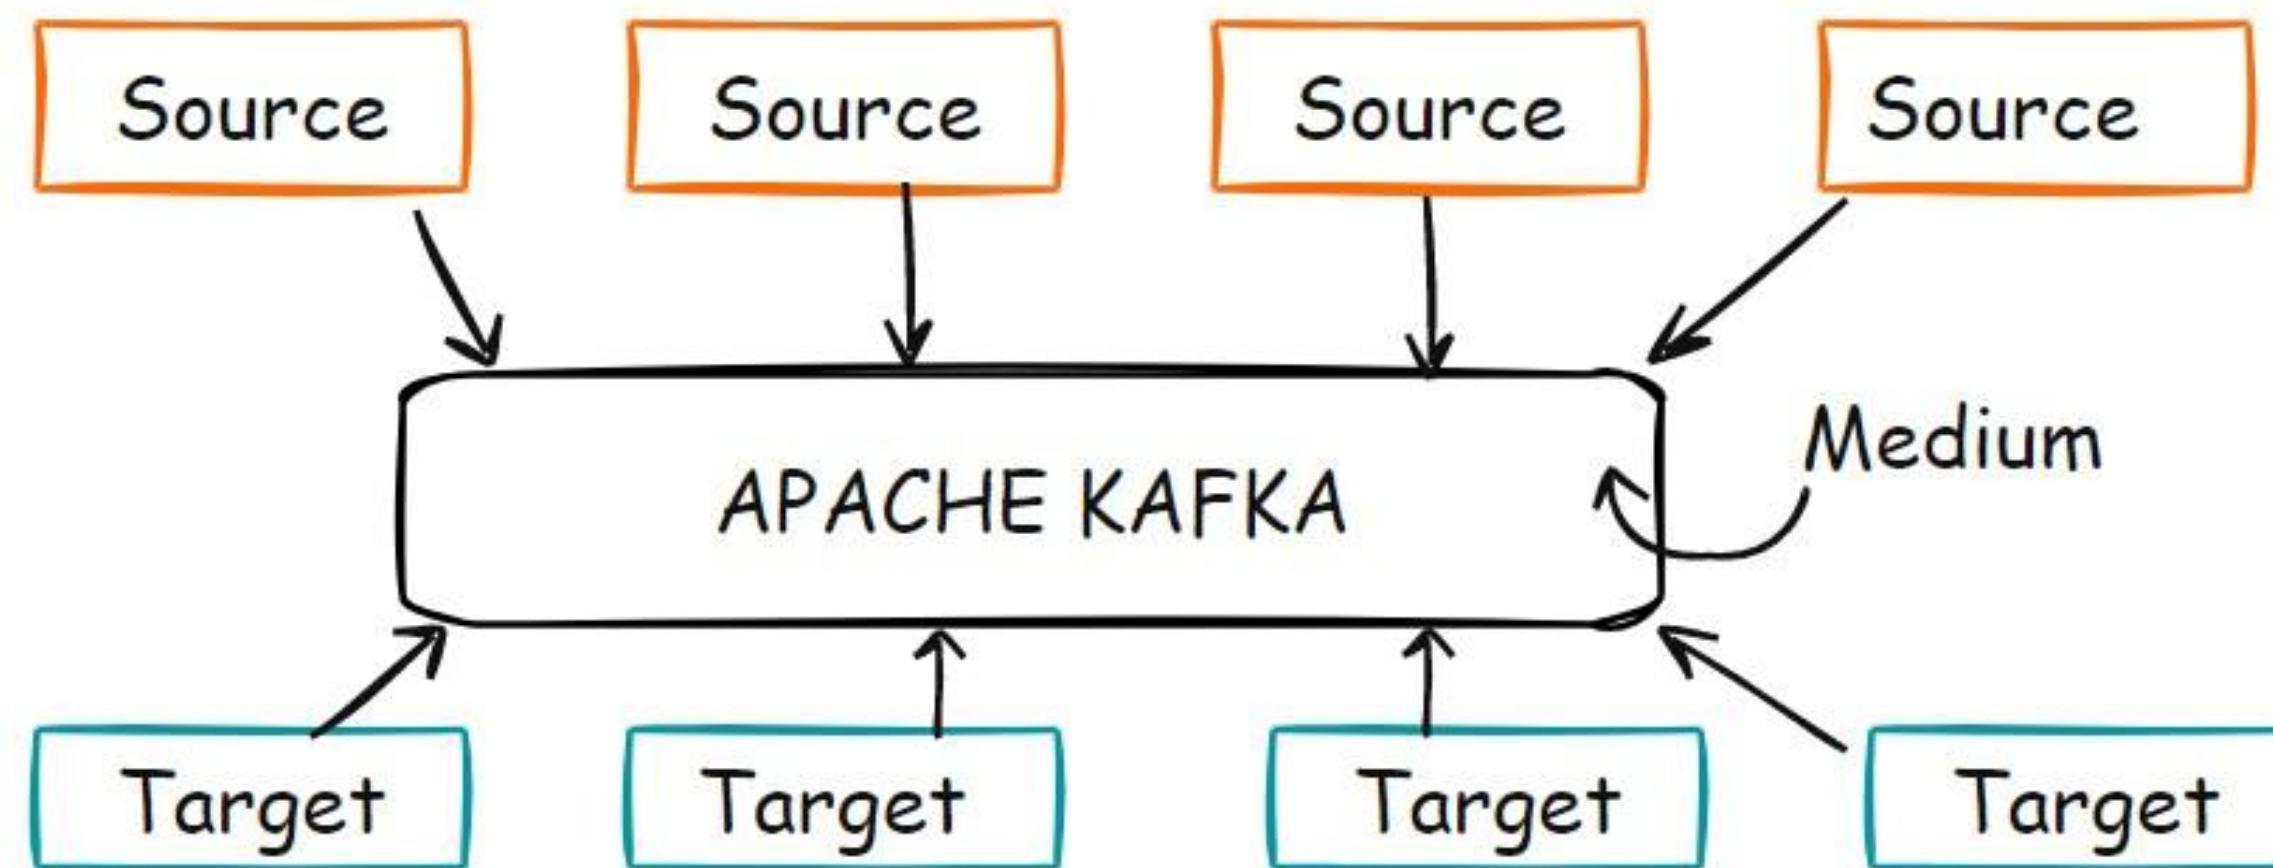
After a while it becomes complicated

- 4 source system X 6 targets system
- 4 X 6 integrations

Challenges we will face

- Protocol – TCP/HTTP/JDBC/REST [How data is transport]
- Data Format – Data is passed in which [Binary , CSV, Array]
- Schema – How the data is shaped
- One more source is introduced , it will increase the network load
- Source & Target system have a very tight coupling

- Apache Kafka – Decoupling of data stream
- pub-sub model :- pub- Publisher , sub-Subscribes
- Pattern where a publisher sends a piece of data – not directly to the receiver, in fact it will send it through a medium



Why Kafka ?

- Created at LinkedIn , -- open source project– Distributed, Resilient , Fault Tolerant Confluent
- Horizontal Scaling
- 100's of brokers , can scale to millions of messages per second
- High Performance
- Latency of less than 10 ms real time

- Apache Kafka – Use Cases

- Messaging System
- Activity Tracking
- Stream Processing
- Gather metric from different location
- Gather application logs
- Integrated with big data modules

3 big differences which set Kafka aside

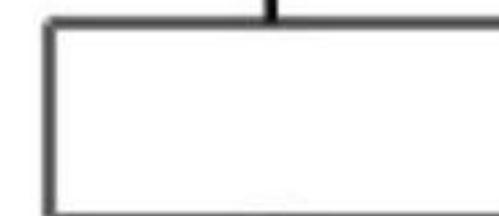
1. Distributed System – rather than running dozens of individual messaging brokers; hard wired to different apps. We have a central platform that can scale
2. Kafka is a true storage system – it is built to store the data as long as you want (default– 1 week) , Real Delivery Guarantee— [replicated | persisted]
3. Kafka supports Stream Processing → Other Messaging System → just hand-out the messages , With Stream Processing in Kafka – allows us to compute derived stream & generate more meaningful data

Producer



Consumer

Producer



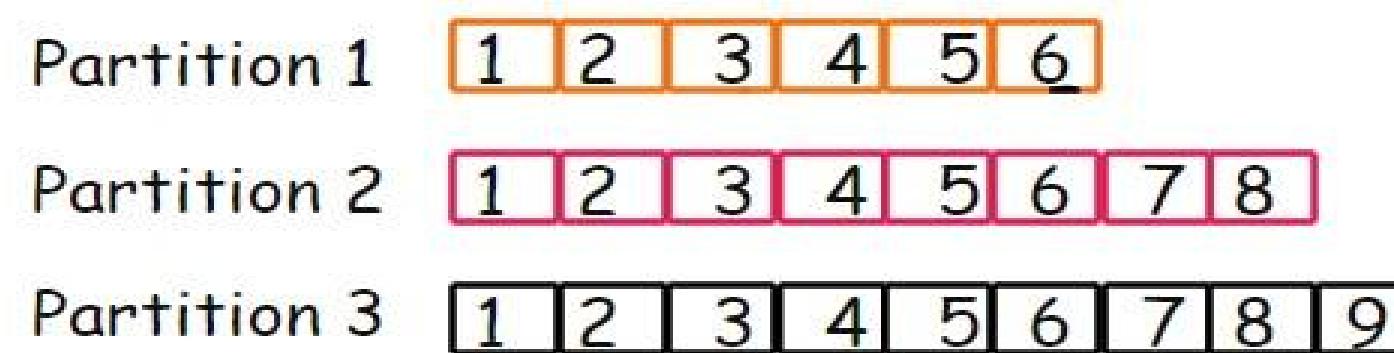
NRT Near Real
Time Processing



Consumer

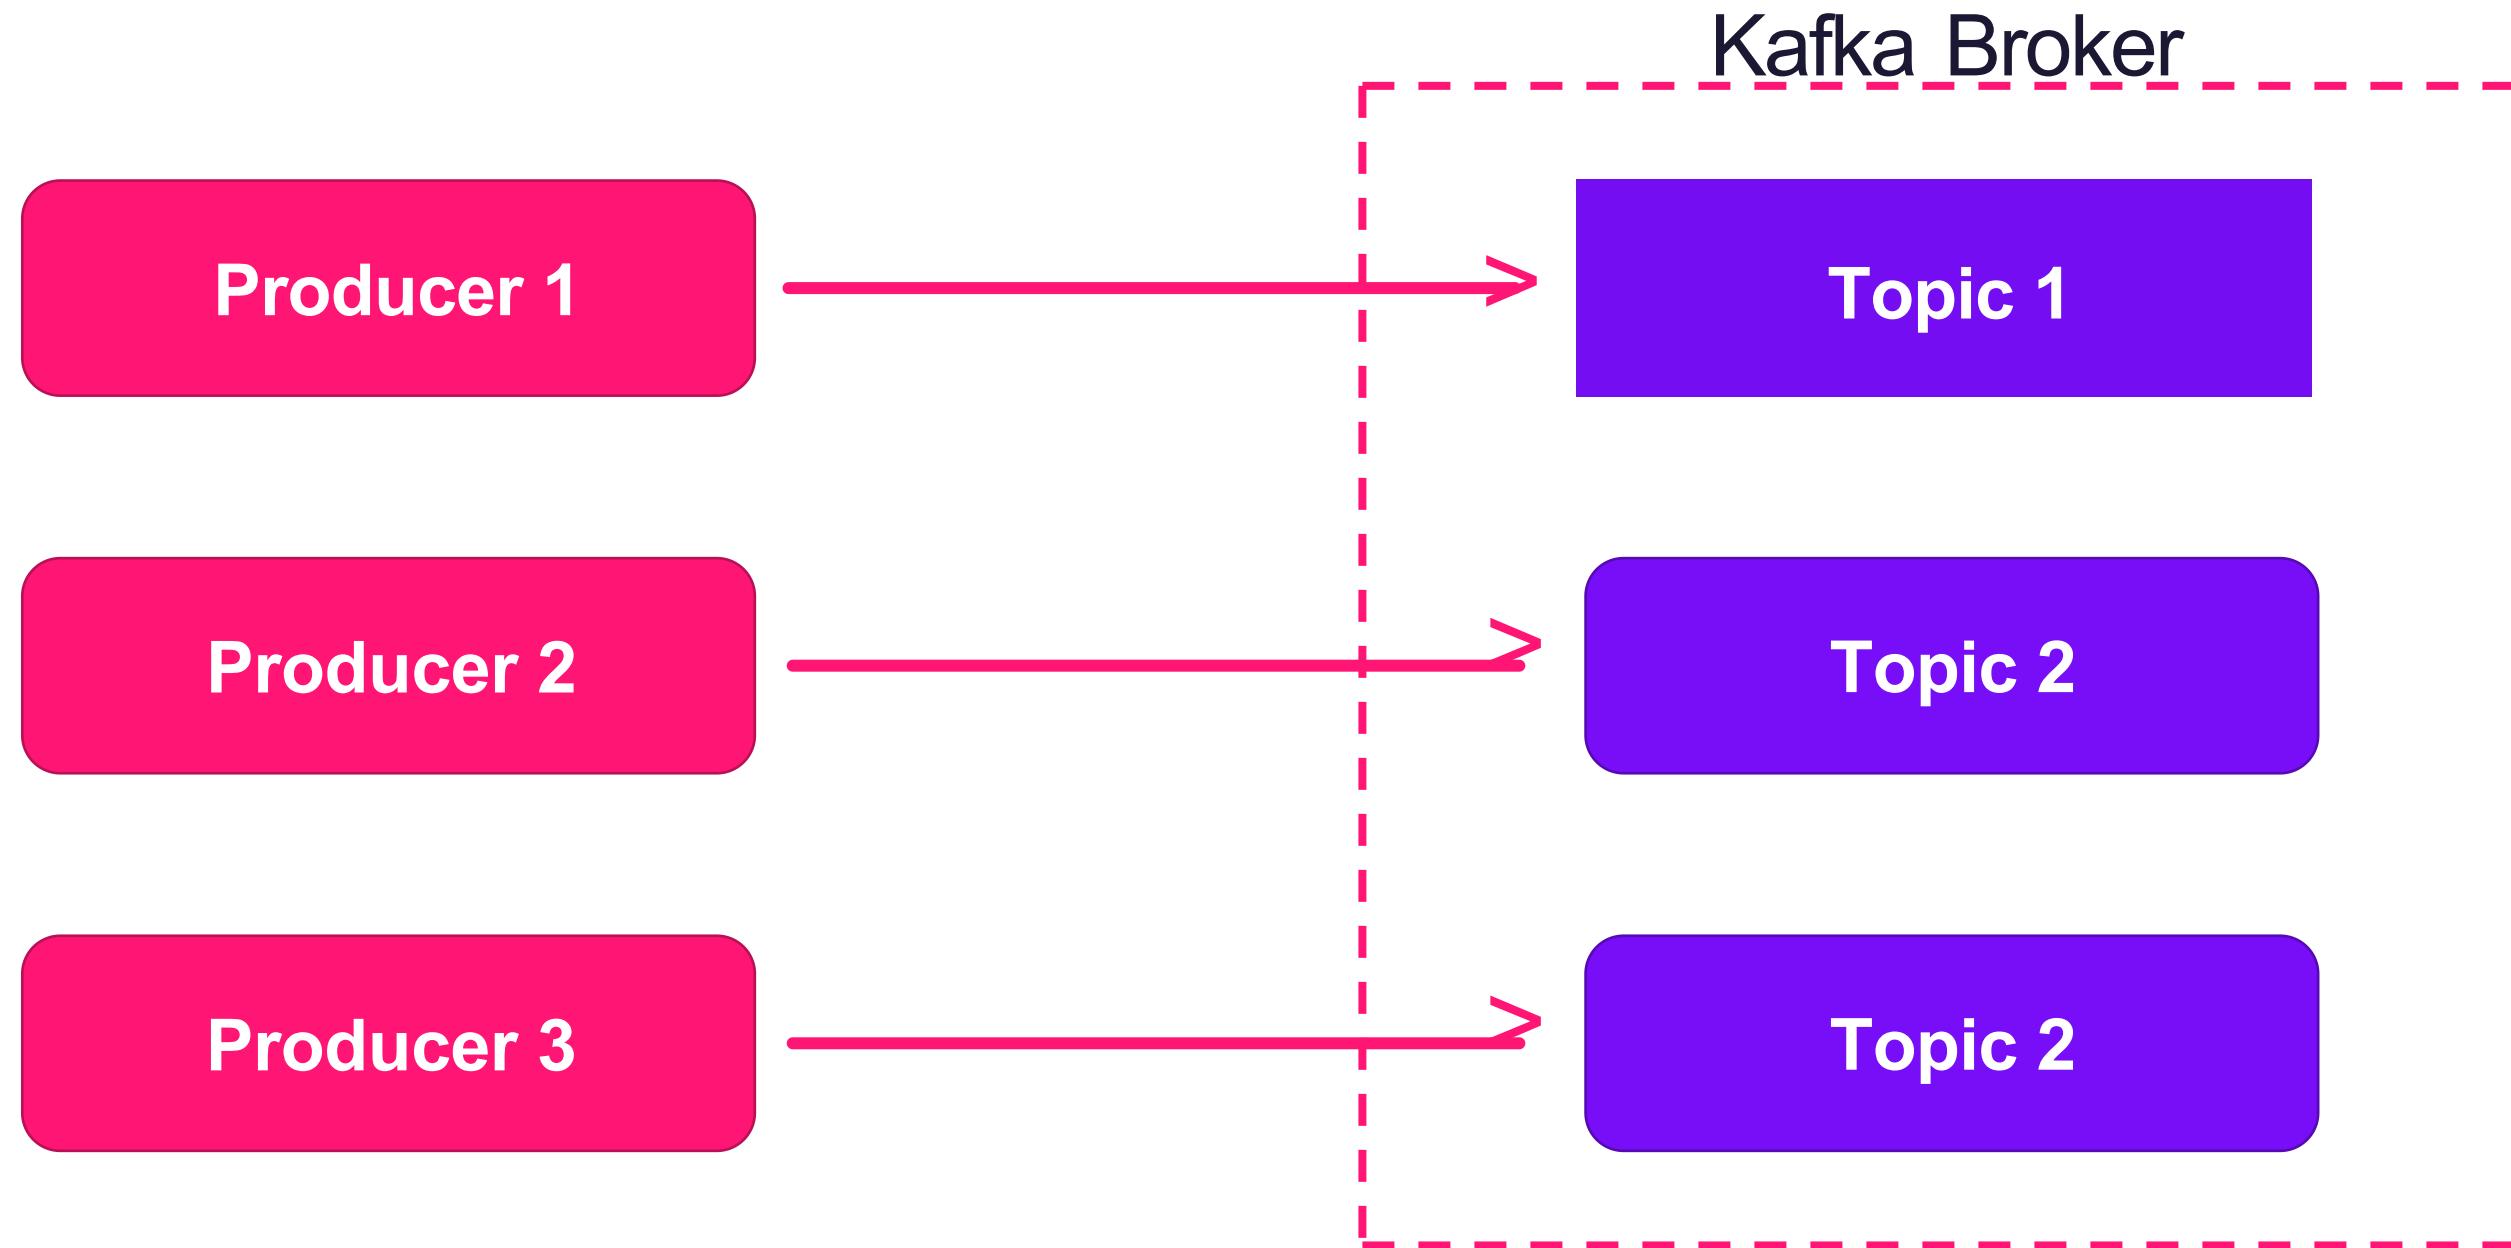
Topics , Partitions , Offset

- Topics- A particular stream of data similar to a table in a database (without all constraints)
- You can have as many as topics you want
- A topic is identified by its name
- Topics are split into Partitions (Kafka Topic)



- Each partition is ordered
- Each message in a partition gets an incremental id , called as offset

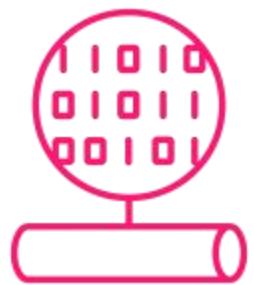
Topics



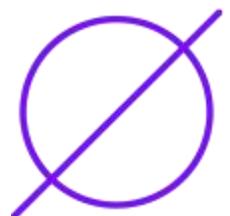
Retention Can Be Set To:



Days: Meaning only keeps messages of up to N days and then delete

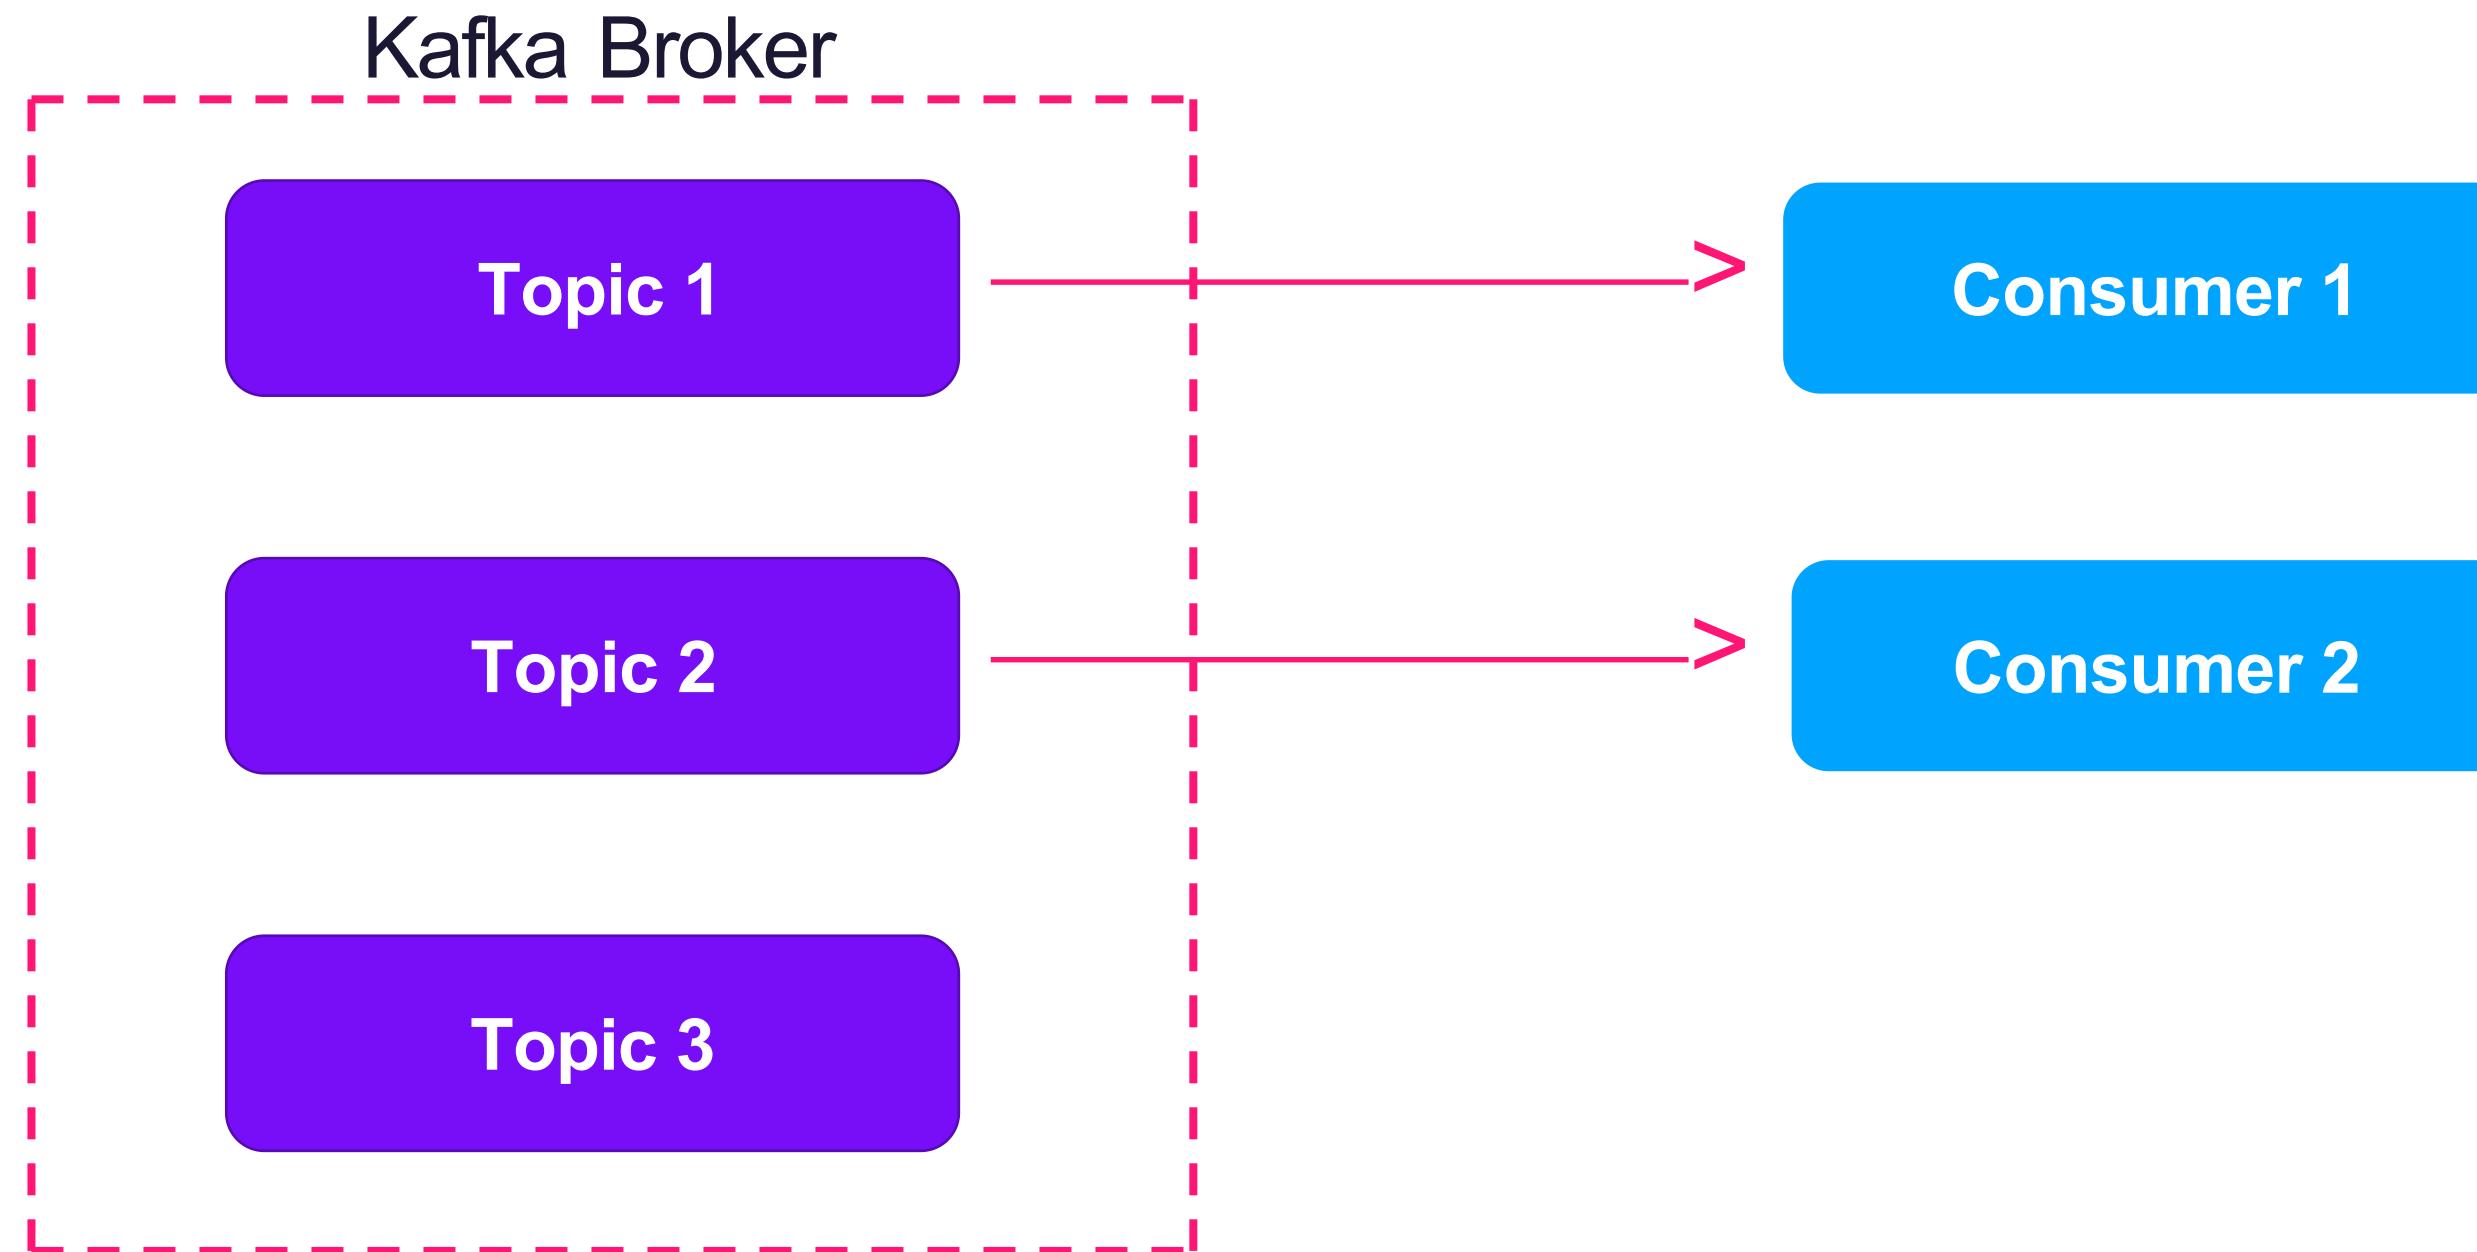


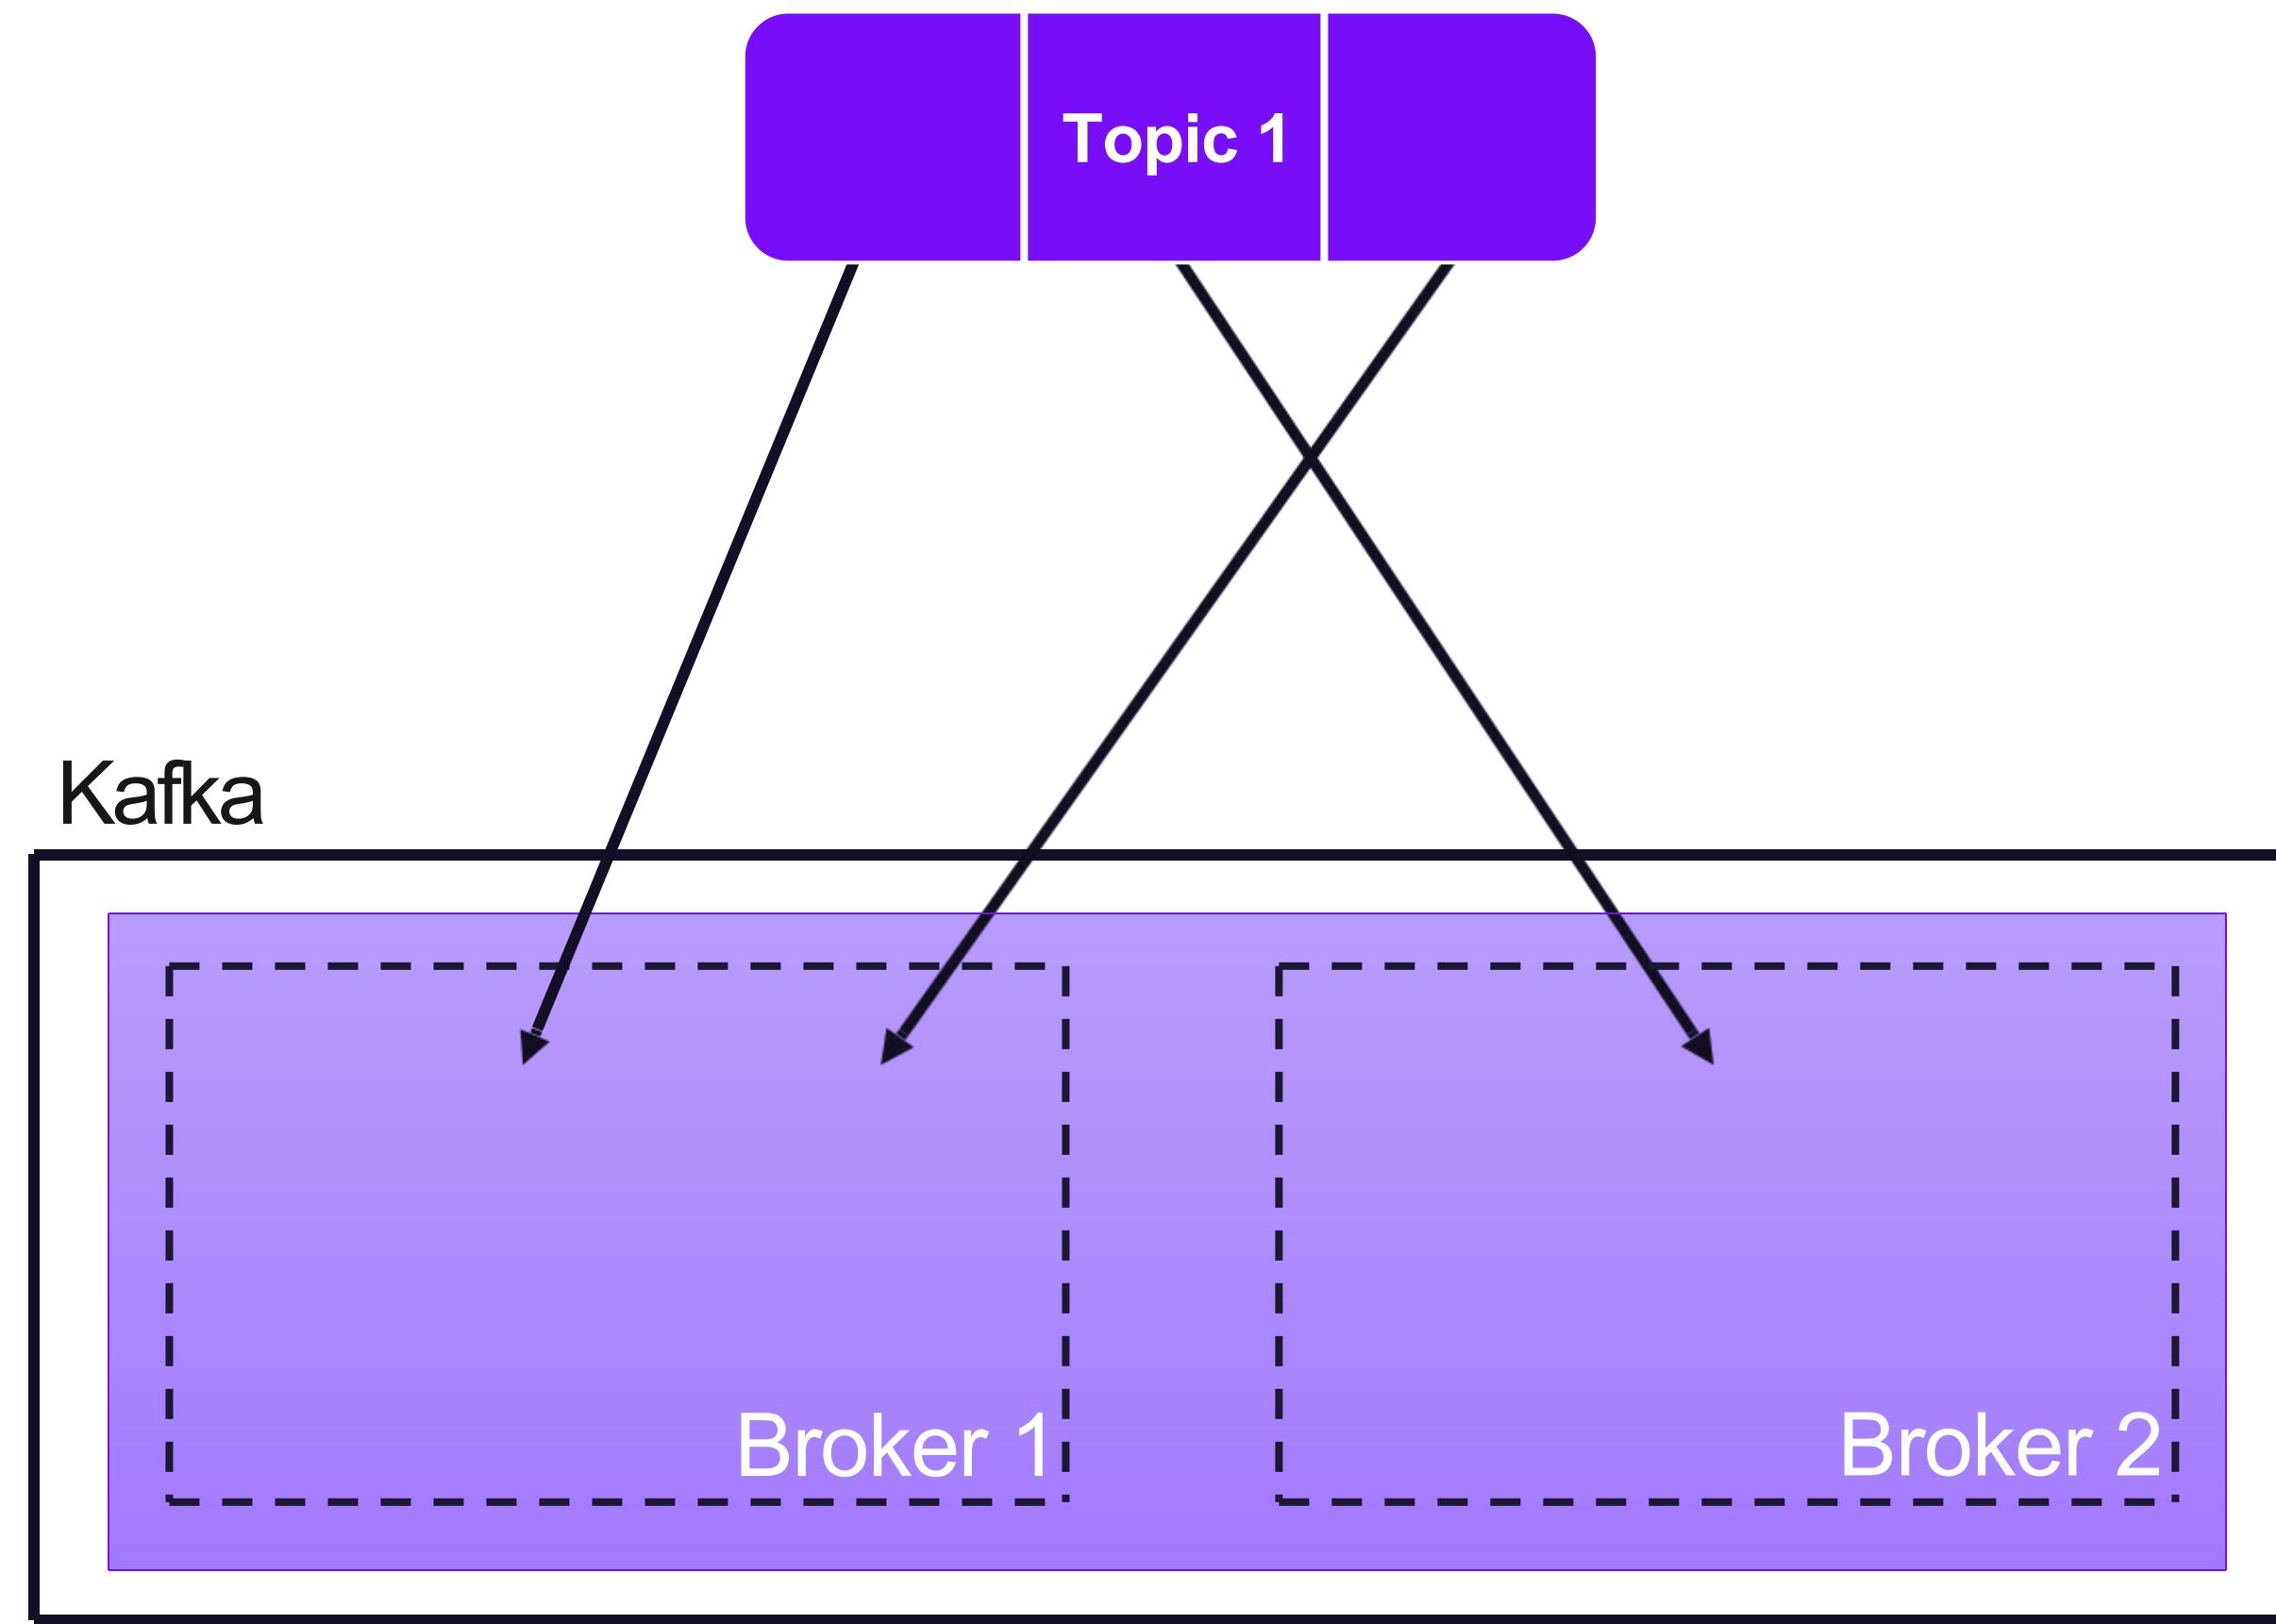
Bytes: Meaning only keep messages until the topic has more than X bytes and then delete



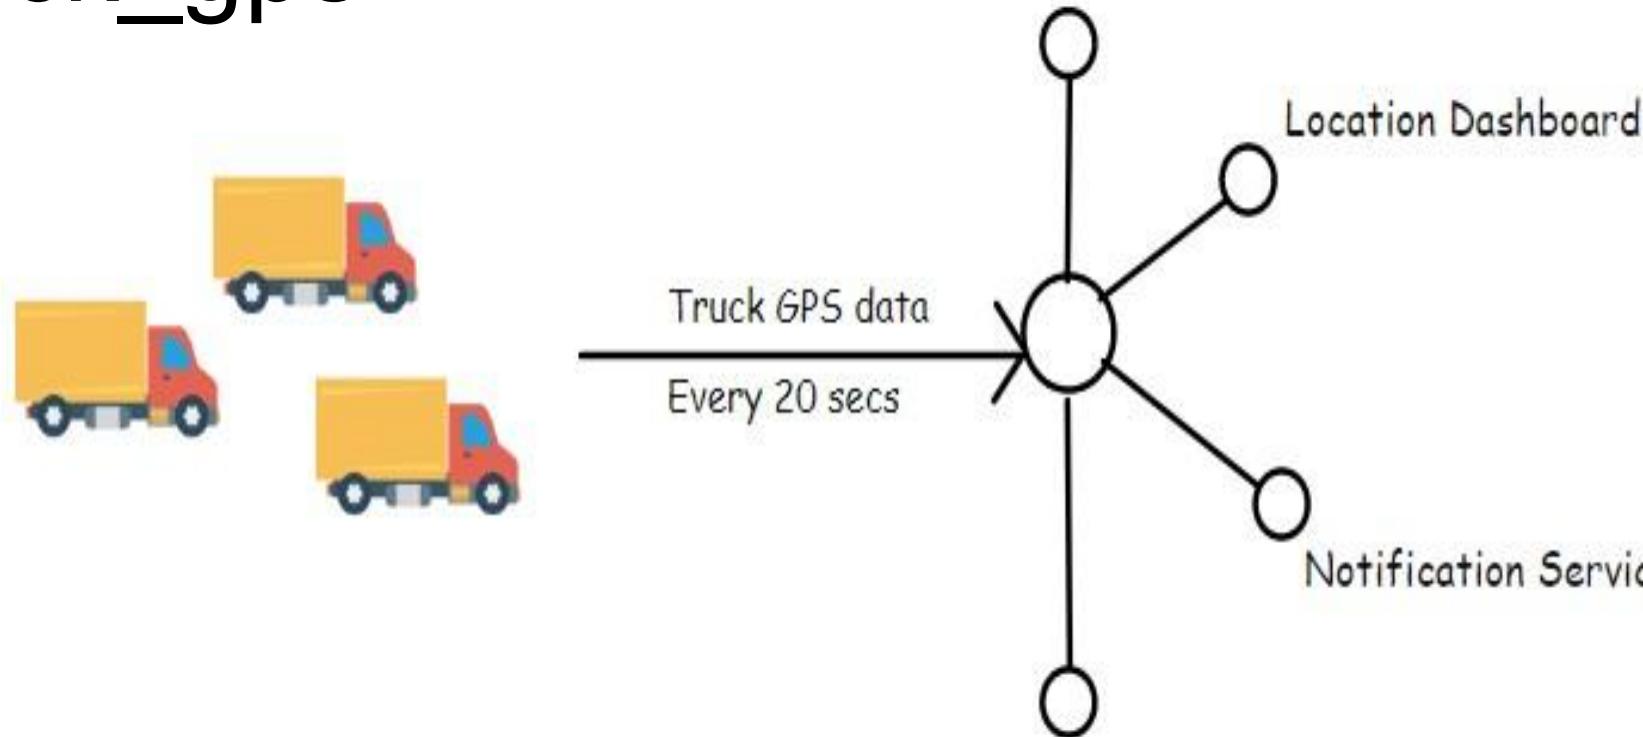
None: Meaning store all messages from the beginning of time

Topics





Topic :- Truck_gps



- Each truck sends the GPS position to Kafka
- Data is sent to the topic- truck_gps
- Each truck sends a message after 20 seconds each message contains truck_id & the truck position (latitude & longitude)

Offset



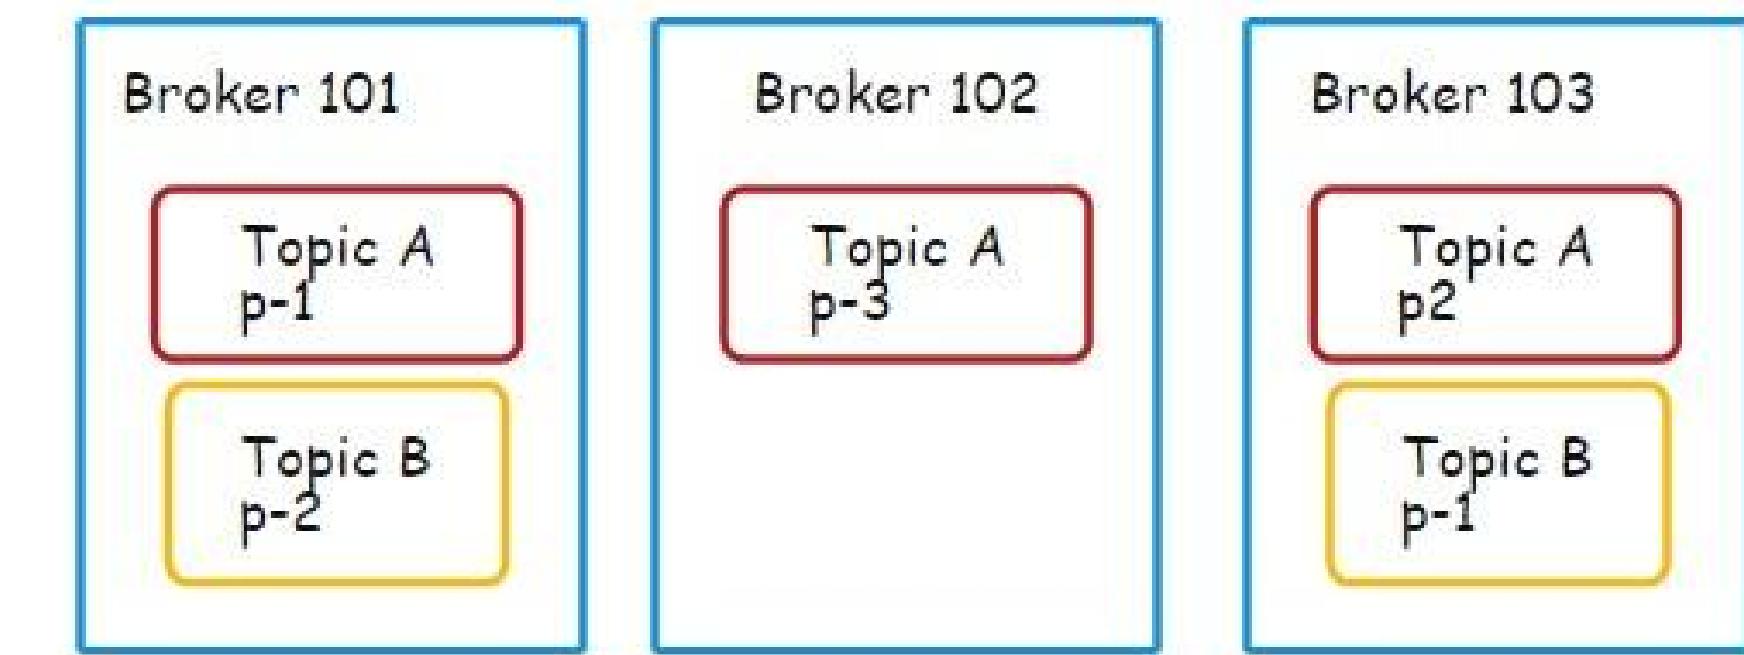
- Offsets have meaning only for a specific partition
- Order is guaranteed only within a partition (not across the partition)
- Data is kept for 1 week (configurable)
- Data once written to the partition it cannot be changed (immutable)
- Data is assigned to the partition randomly unless the key is provided

Brokers & Topics

- A Kafka cluster is composed of multiple brokers (servers)
- Each broker is identified with its id(integer)
- Each broker contains certain topic partitions
- We connect to any broker— Bootstrap Server
- No of brokers-3

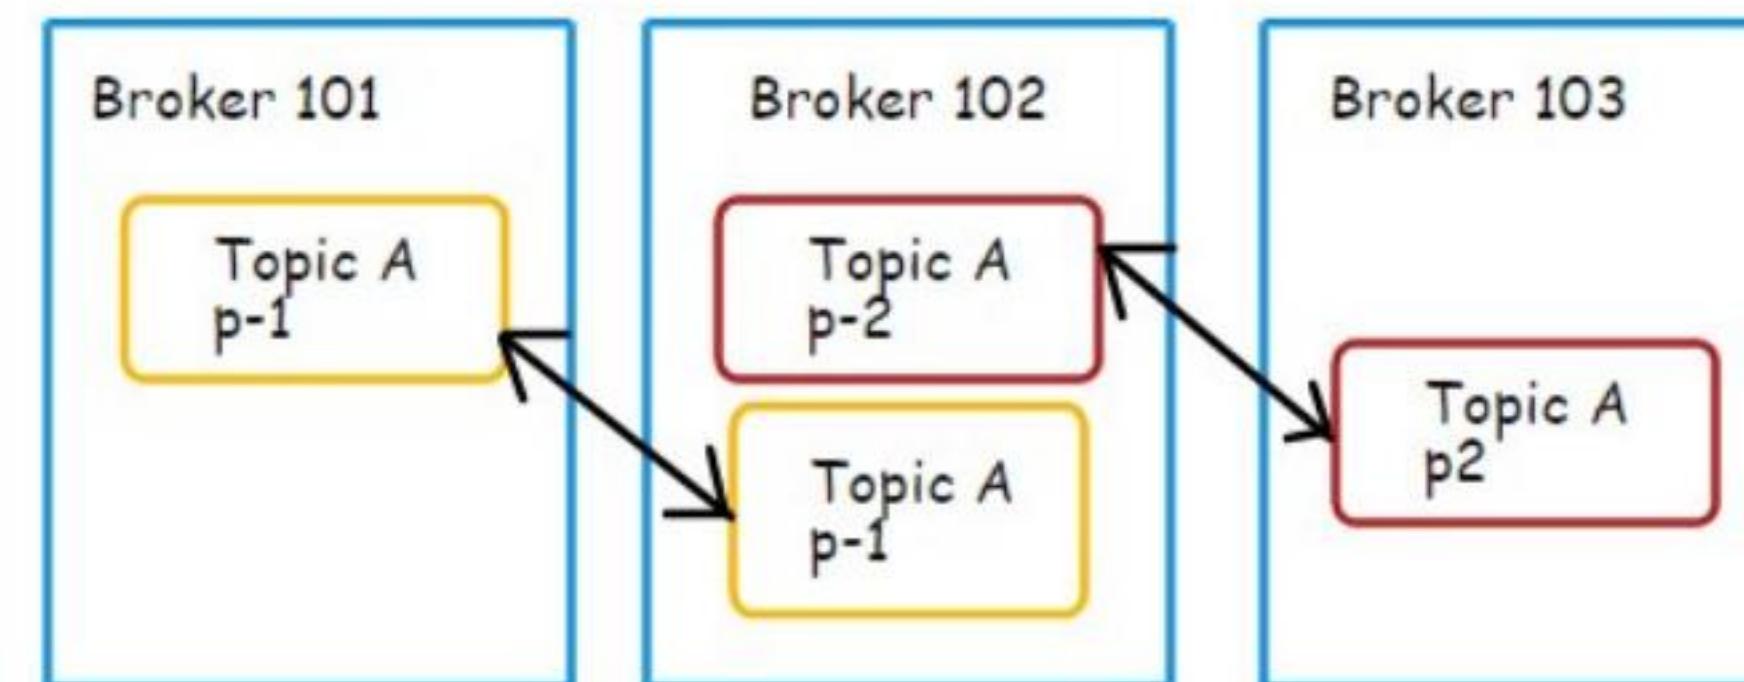
Topic A- 3 partitions

Topic B- 2 partitions



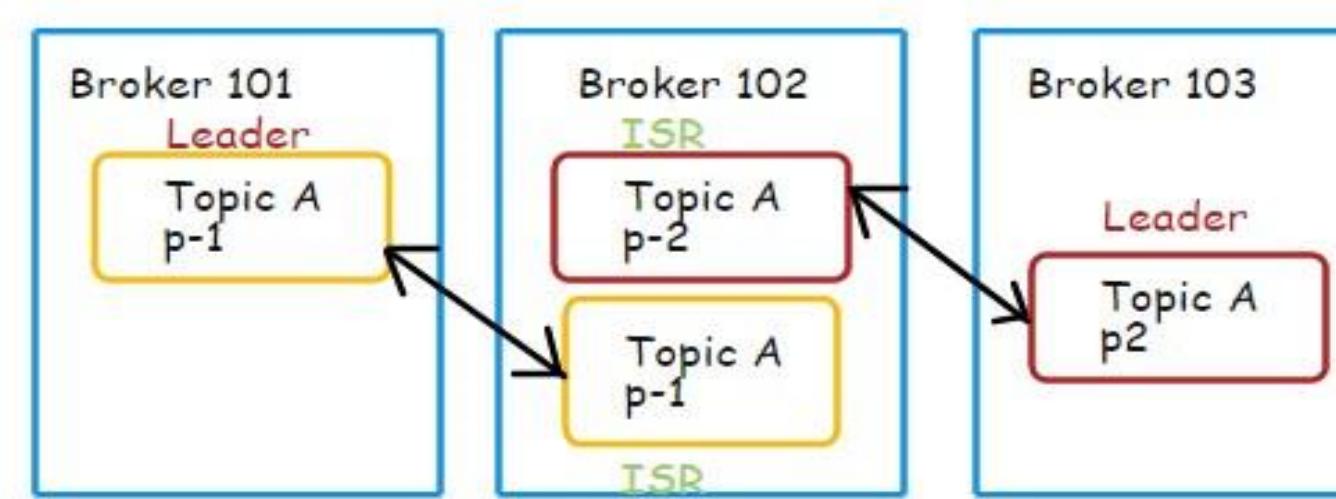
Topic Replication Factor

- A topic should have Replication Factor greater than 1 (usually between 2&3)
- Replication Factor make sure that if a broker is down ; another broker can serve the data
- Rf-2 [1 original + 1 copy]
- Topic A – 2 partition & RF- 2



Concept of Leader for a Partition

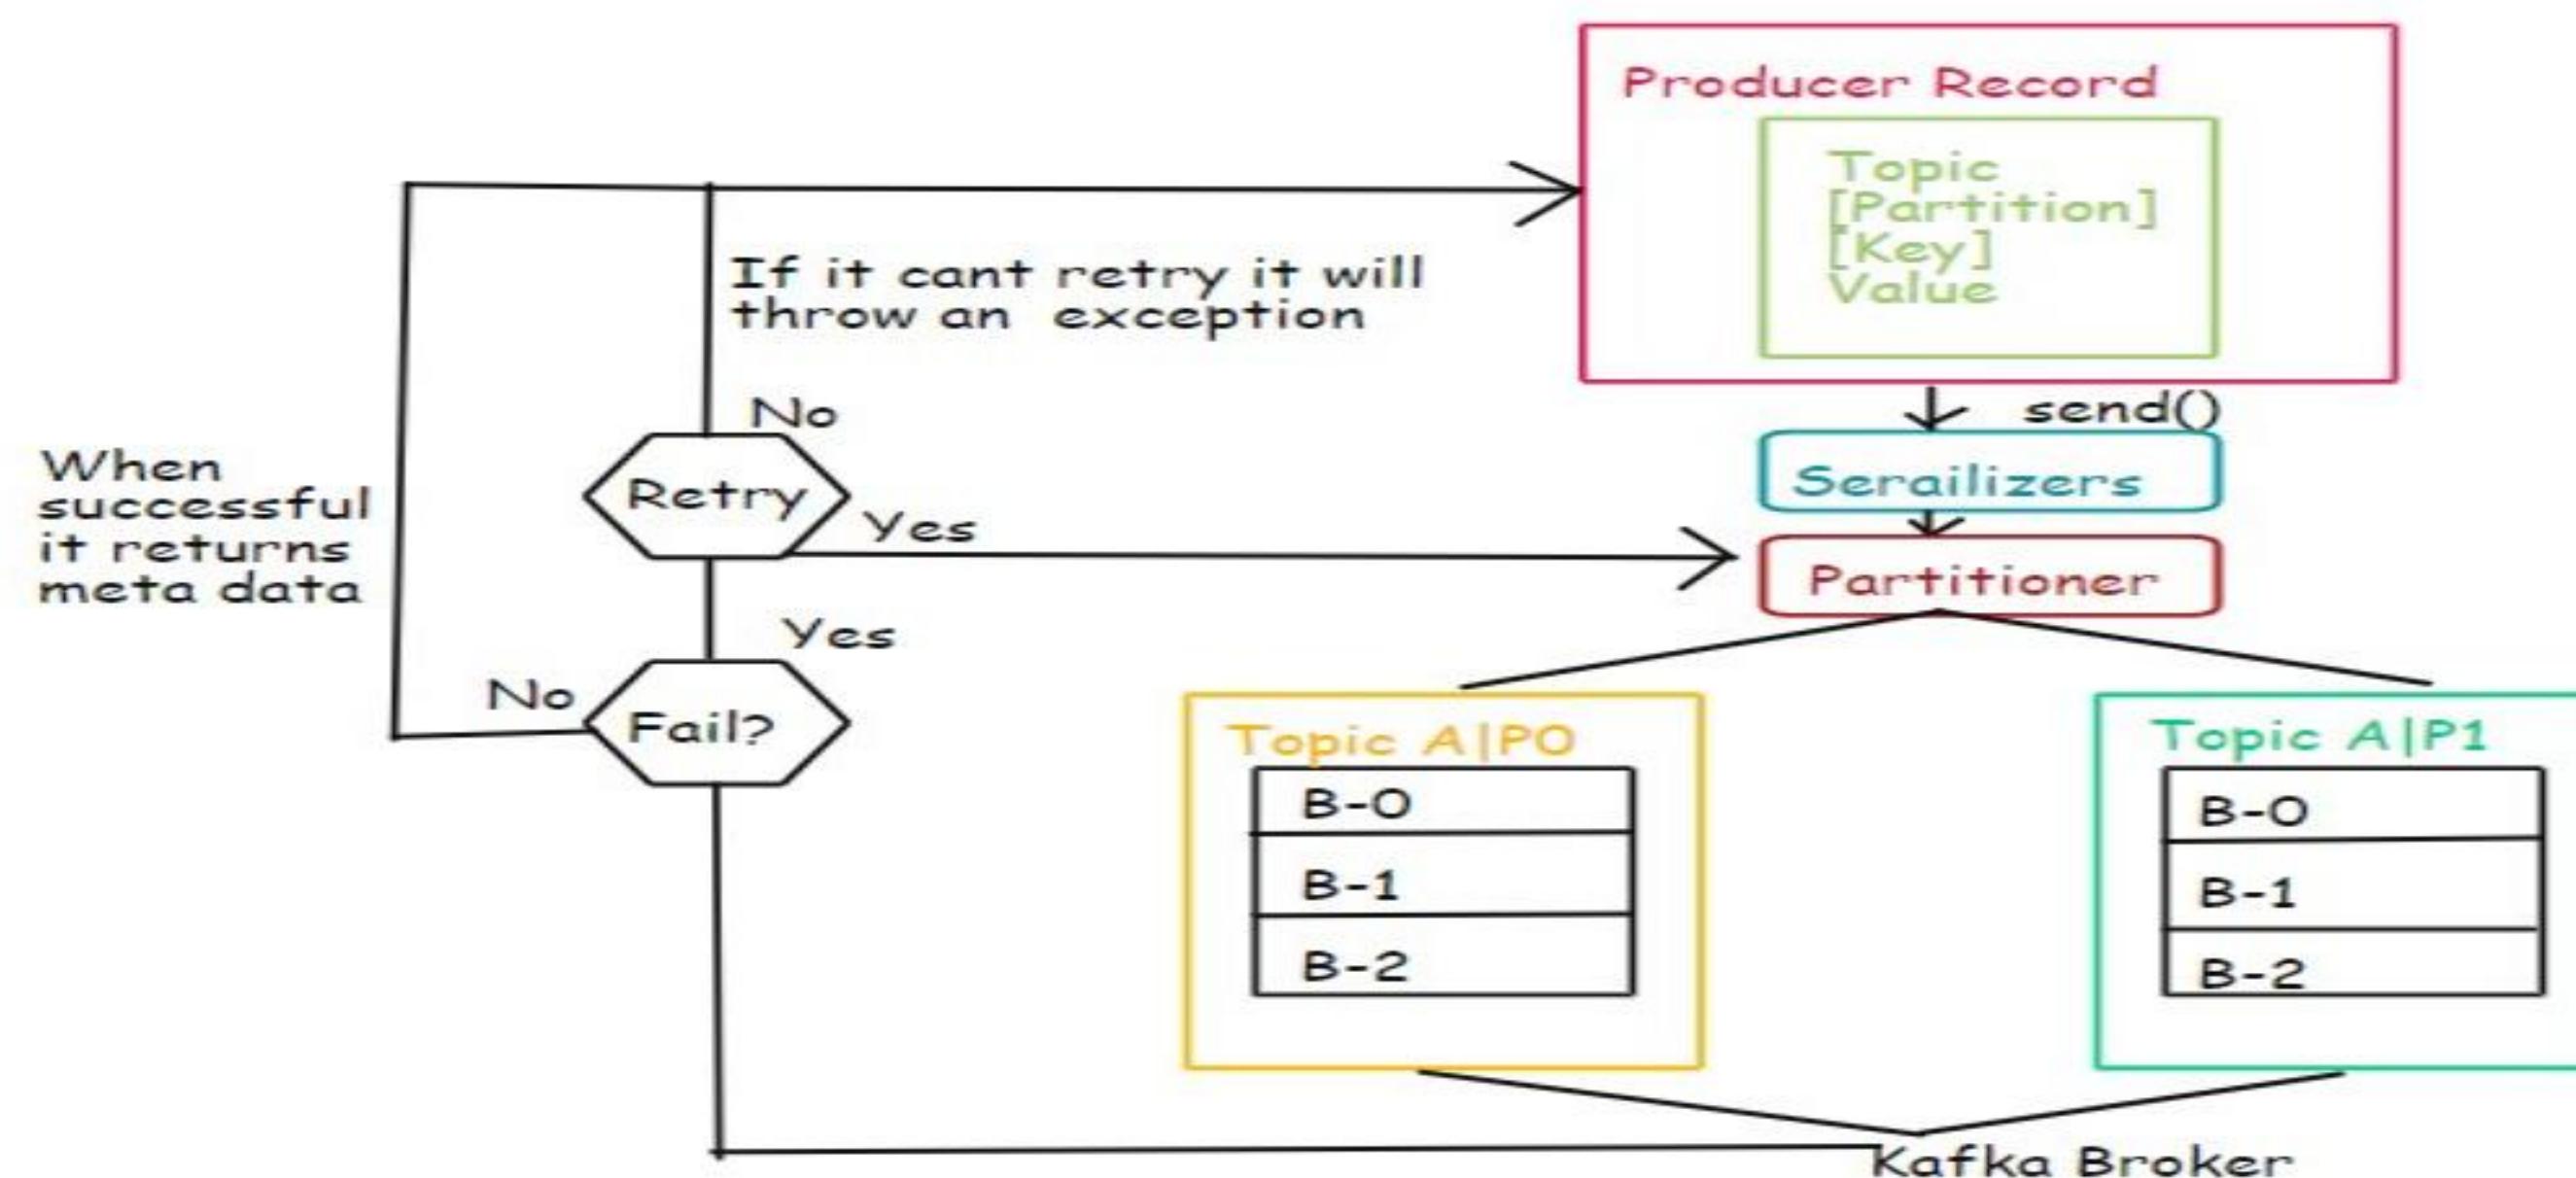
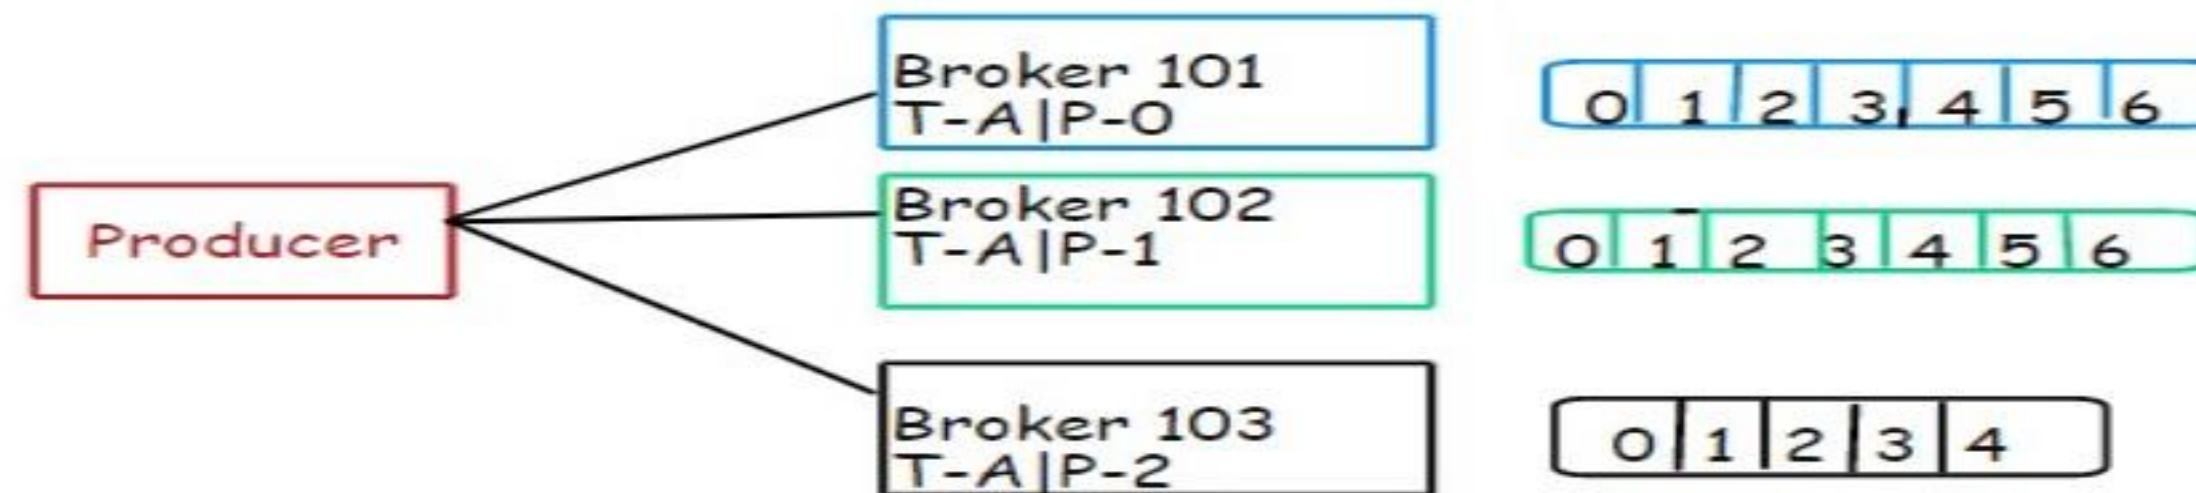
- [At any point in time; only one BROKER can be a LEADER for a PARTITION]
- [ONLY that LEADER can RECEIVE & SERVE DATA FOR that Partition]



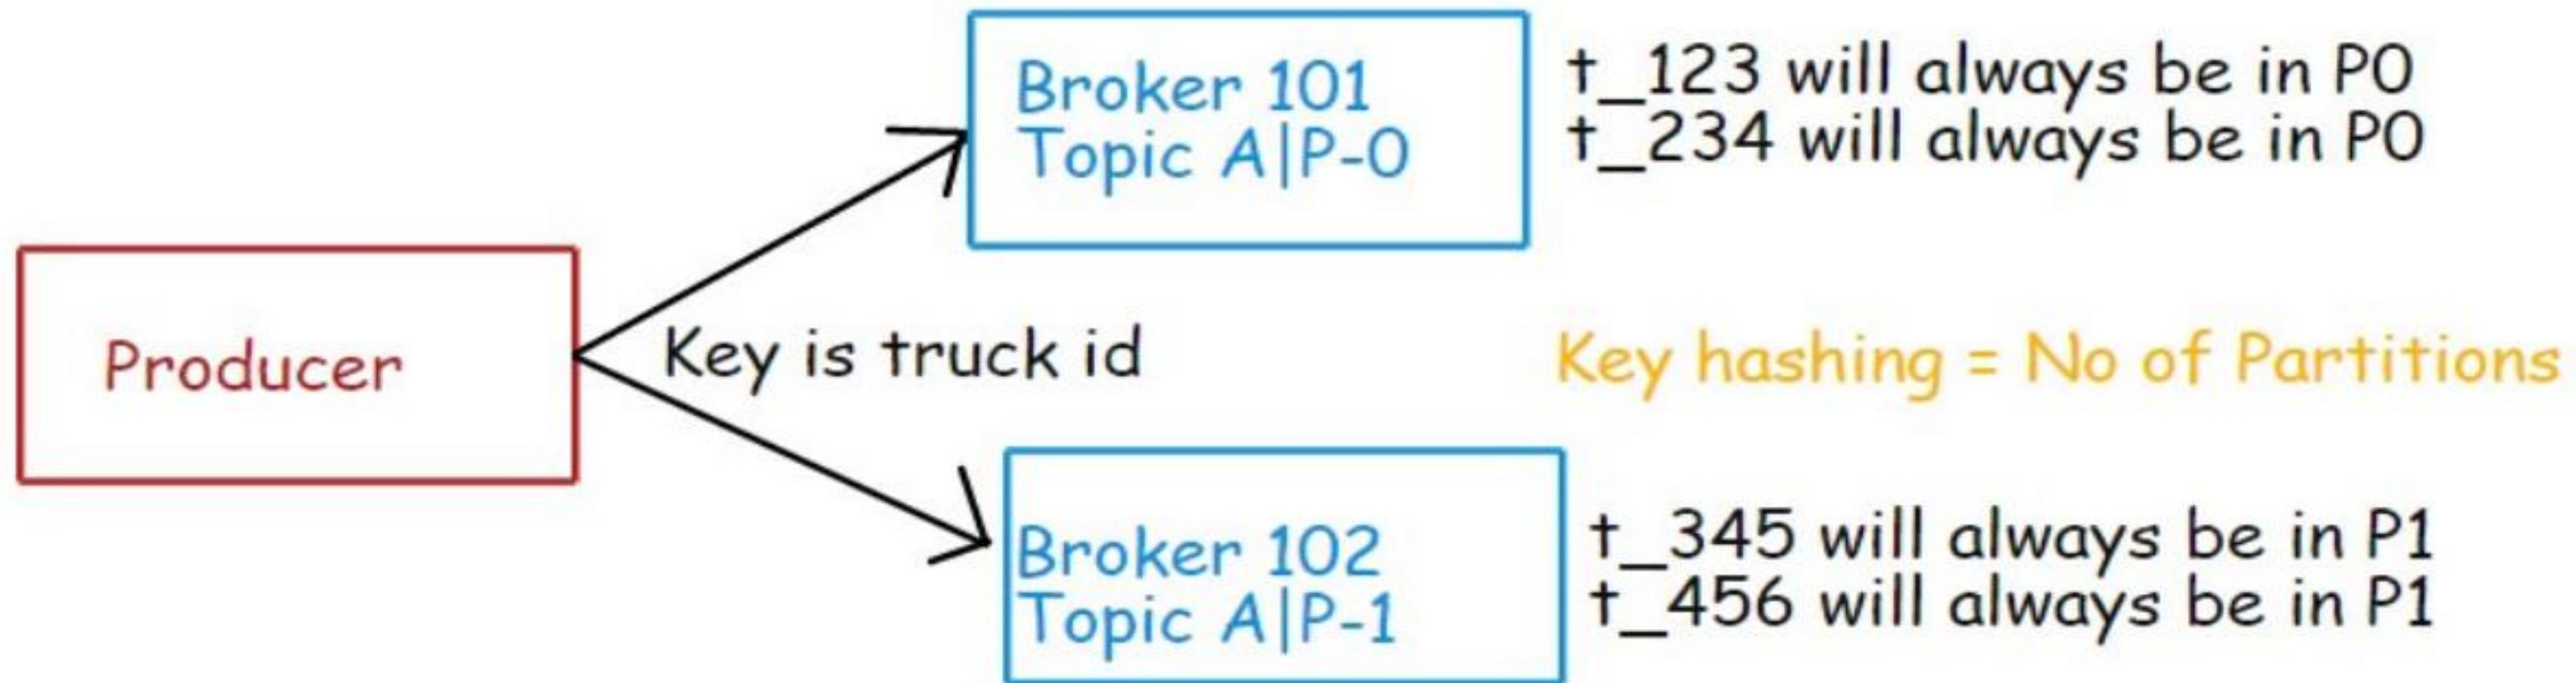
- Other brokers will synchronize the data
- Each partition has ONE LEADER & Multiple ISR[In Sync Replicas]

Producers

- Producer writes the data to the topics
- Producers automatically knows to which broker & partition to write to.
- In case a Broker fails; Producer will automatically recover

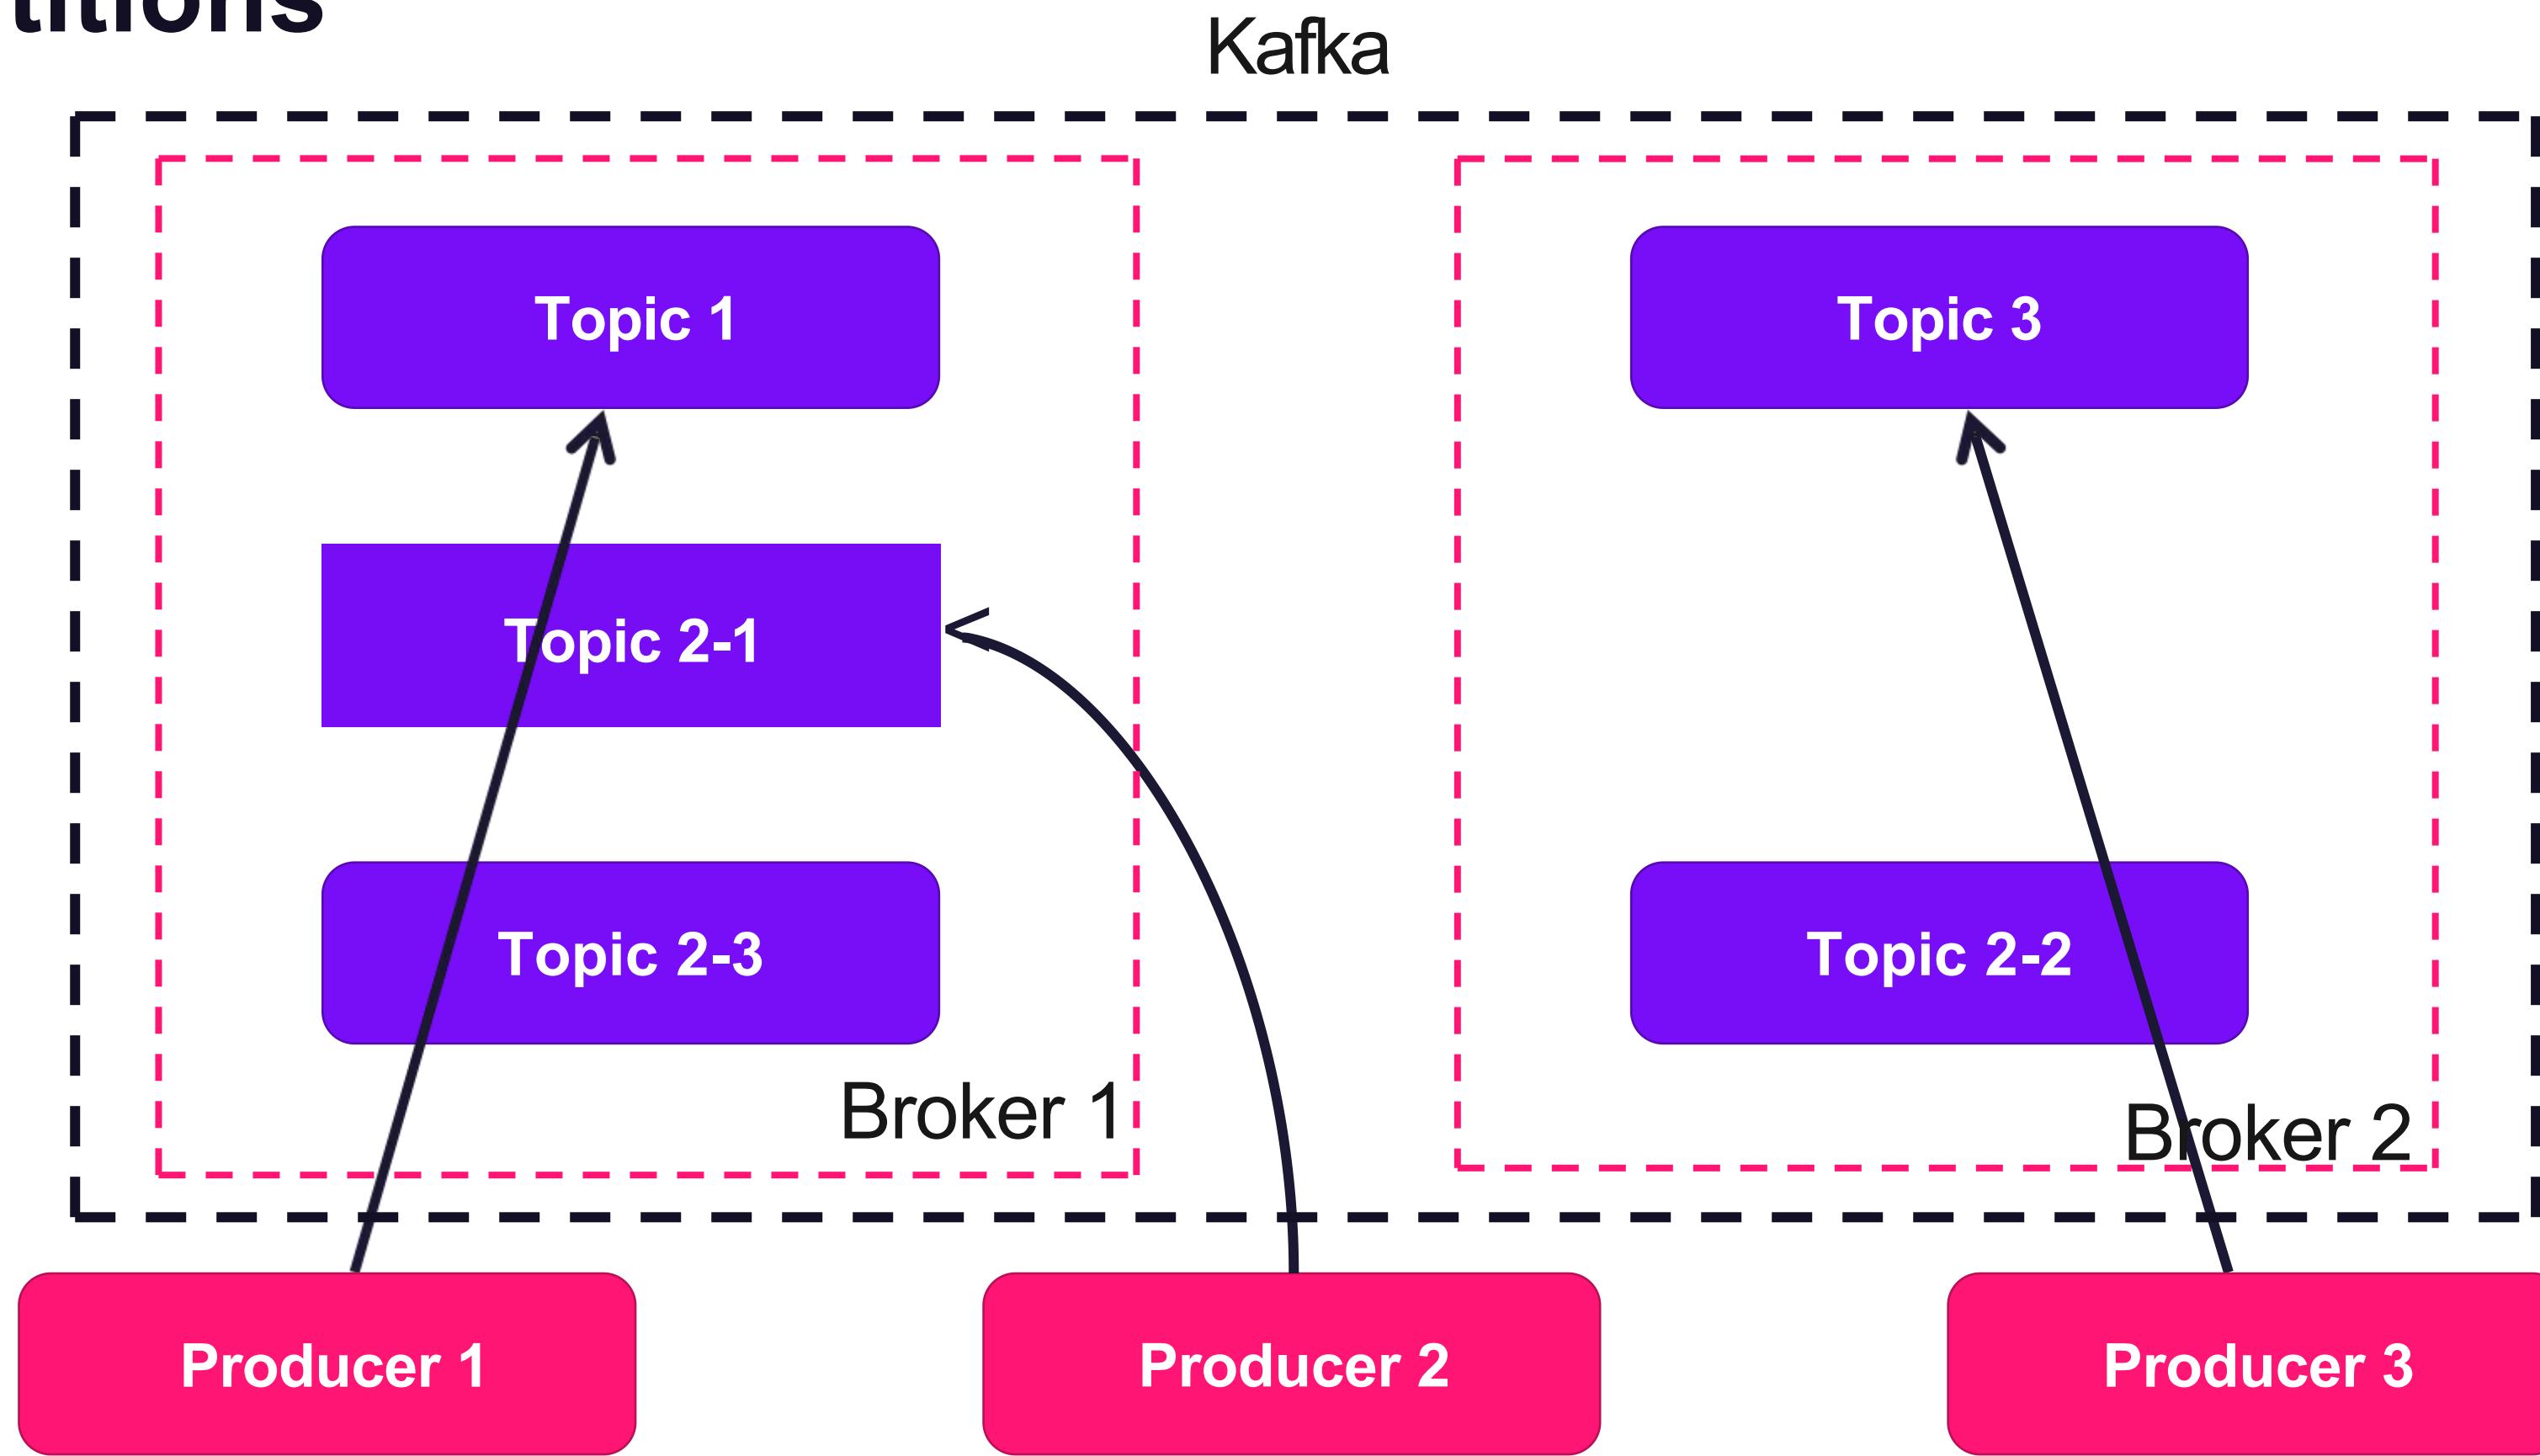


- Producers can choose to receive the acknowledgment of data writes
 - Acks=0 | Producer will not wait for acknowledgment
(possible data loss)
 - Acks=1 | Producer will wait for ‘Leaders Acknowledgment’
(limited data loss)
 - Acks=all | Leader + Replicas Acknowledgment **(no data loss)**
-
- Producer– Message Key’s
 - Producer can choose to send a key with the message (string , number, etc)
 - If key=null; data is sent in Round Robin (101-102-103)
 - If key is sent ; then all the messages for that particular key will always land up in same partition
 - ❖ A key is normally sent with the message when we need ordering for a specific field ex:truck_id

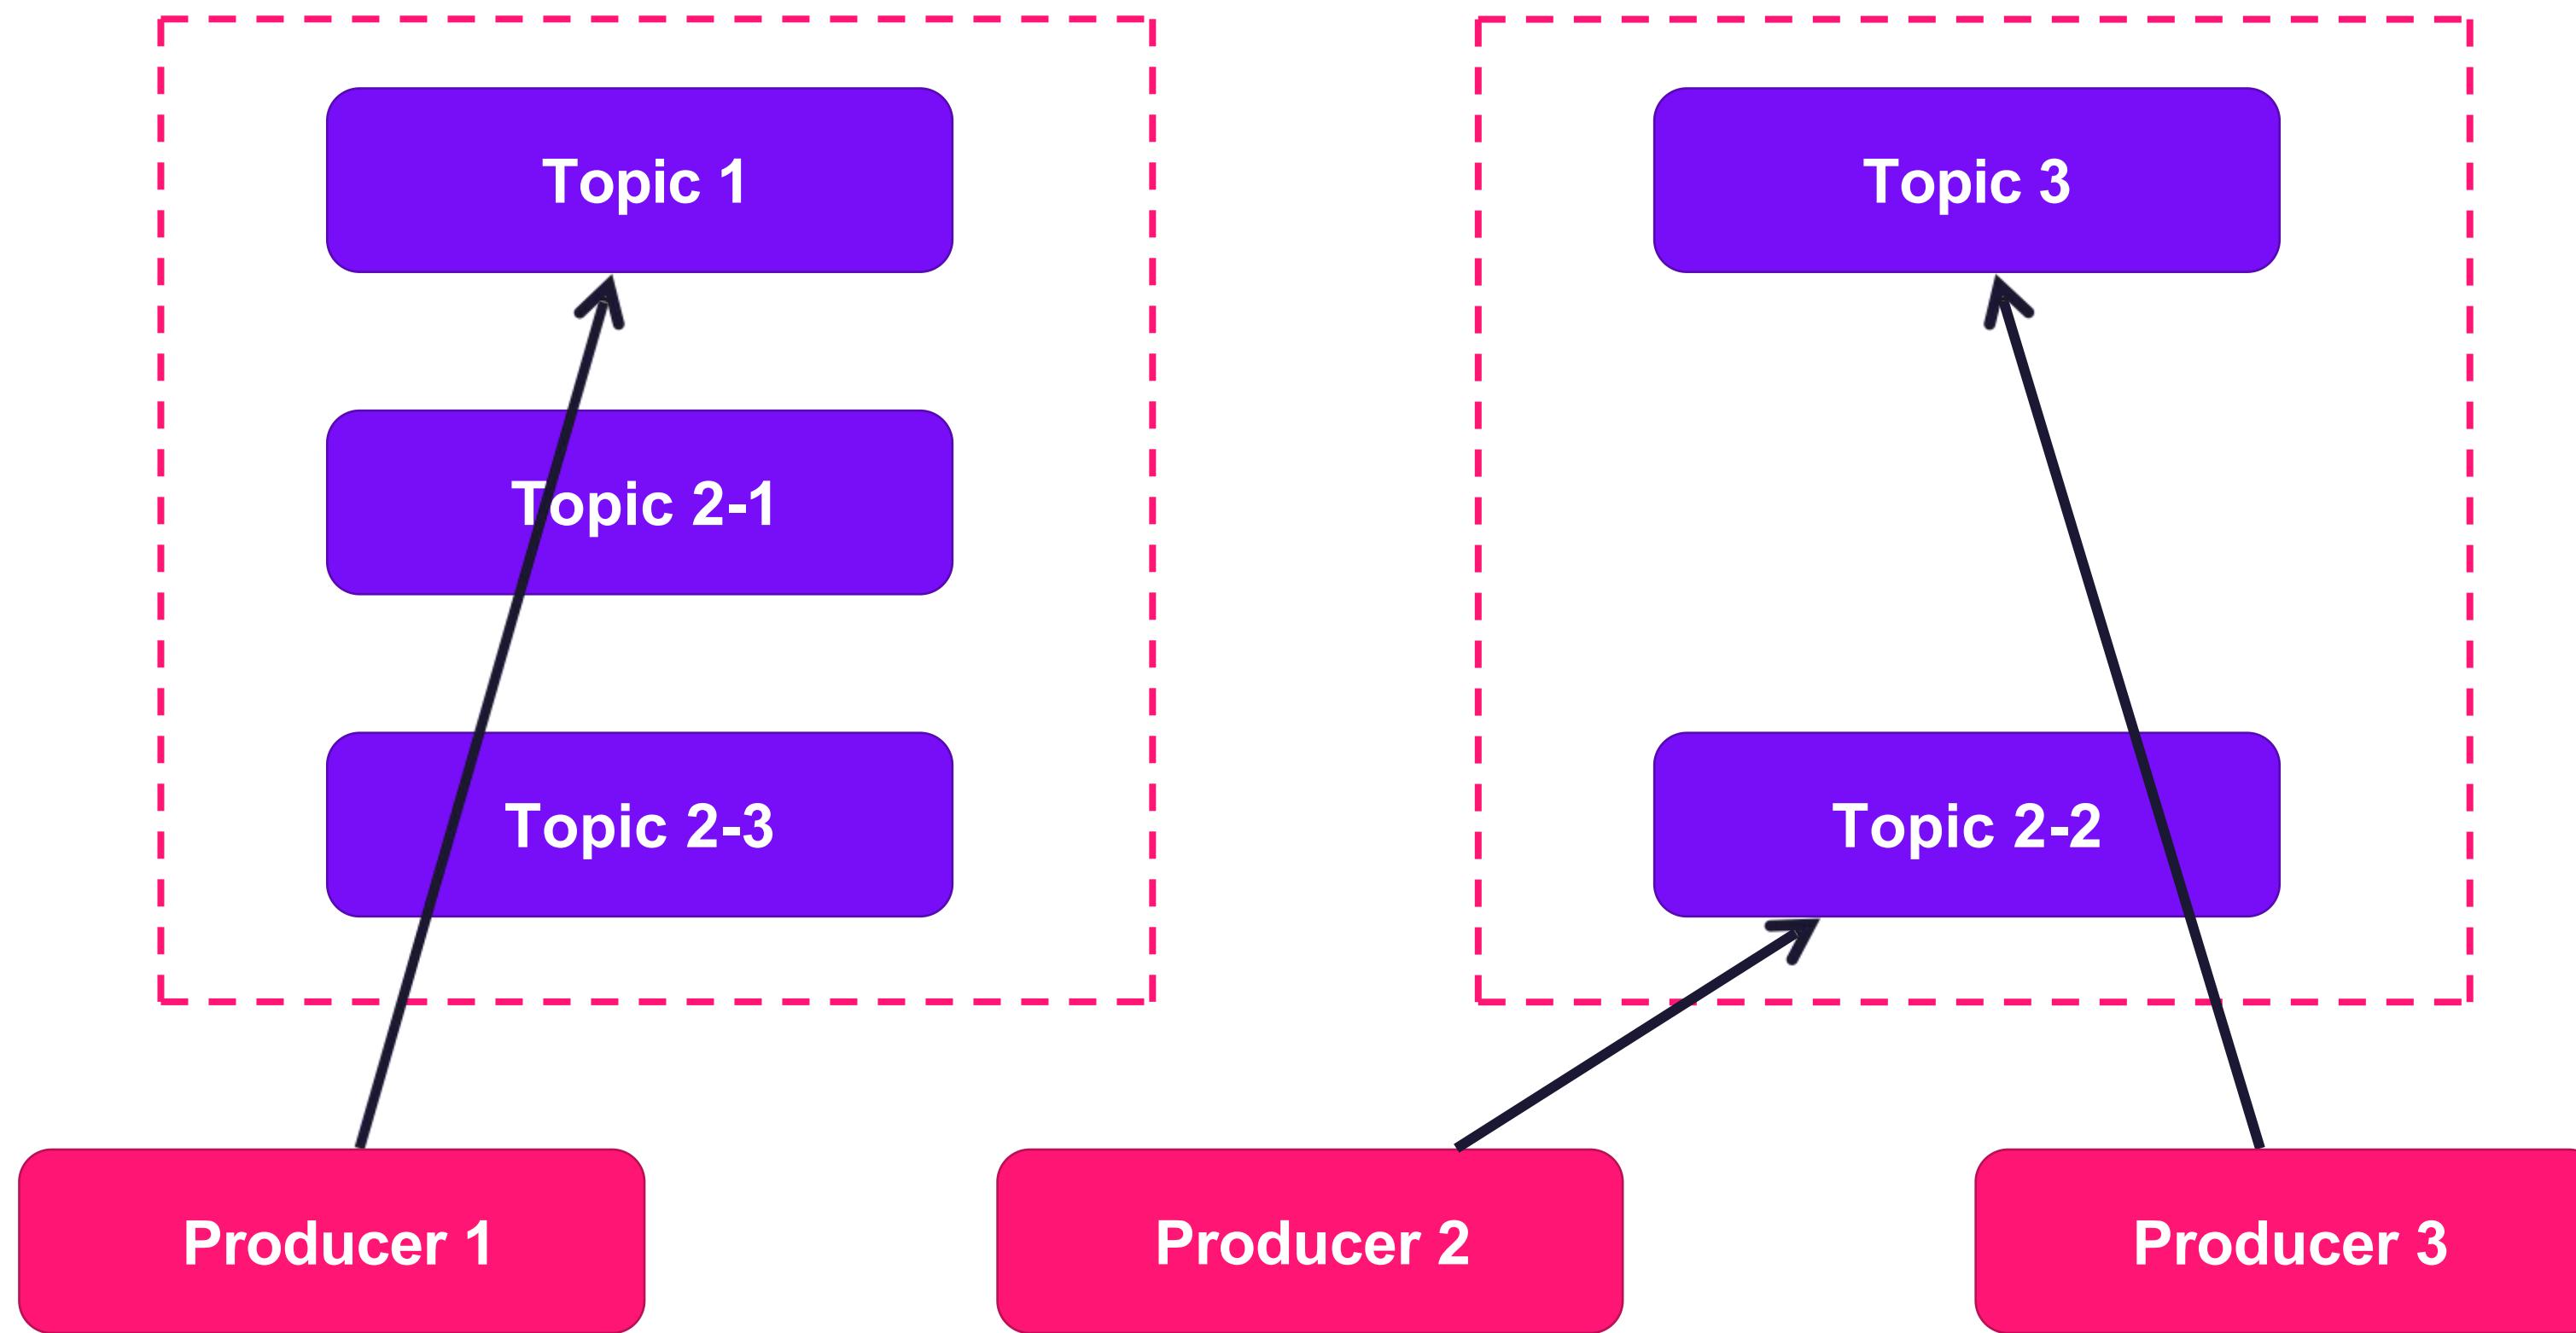


The partitioner uses the key
to distribute the messages
to the correct partition

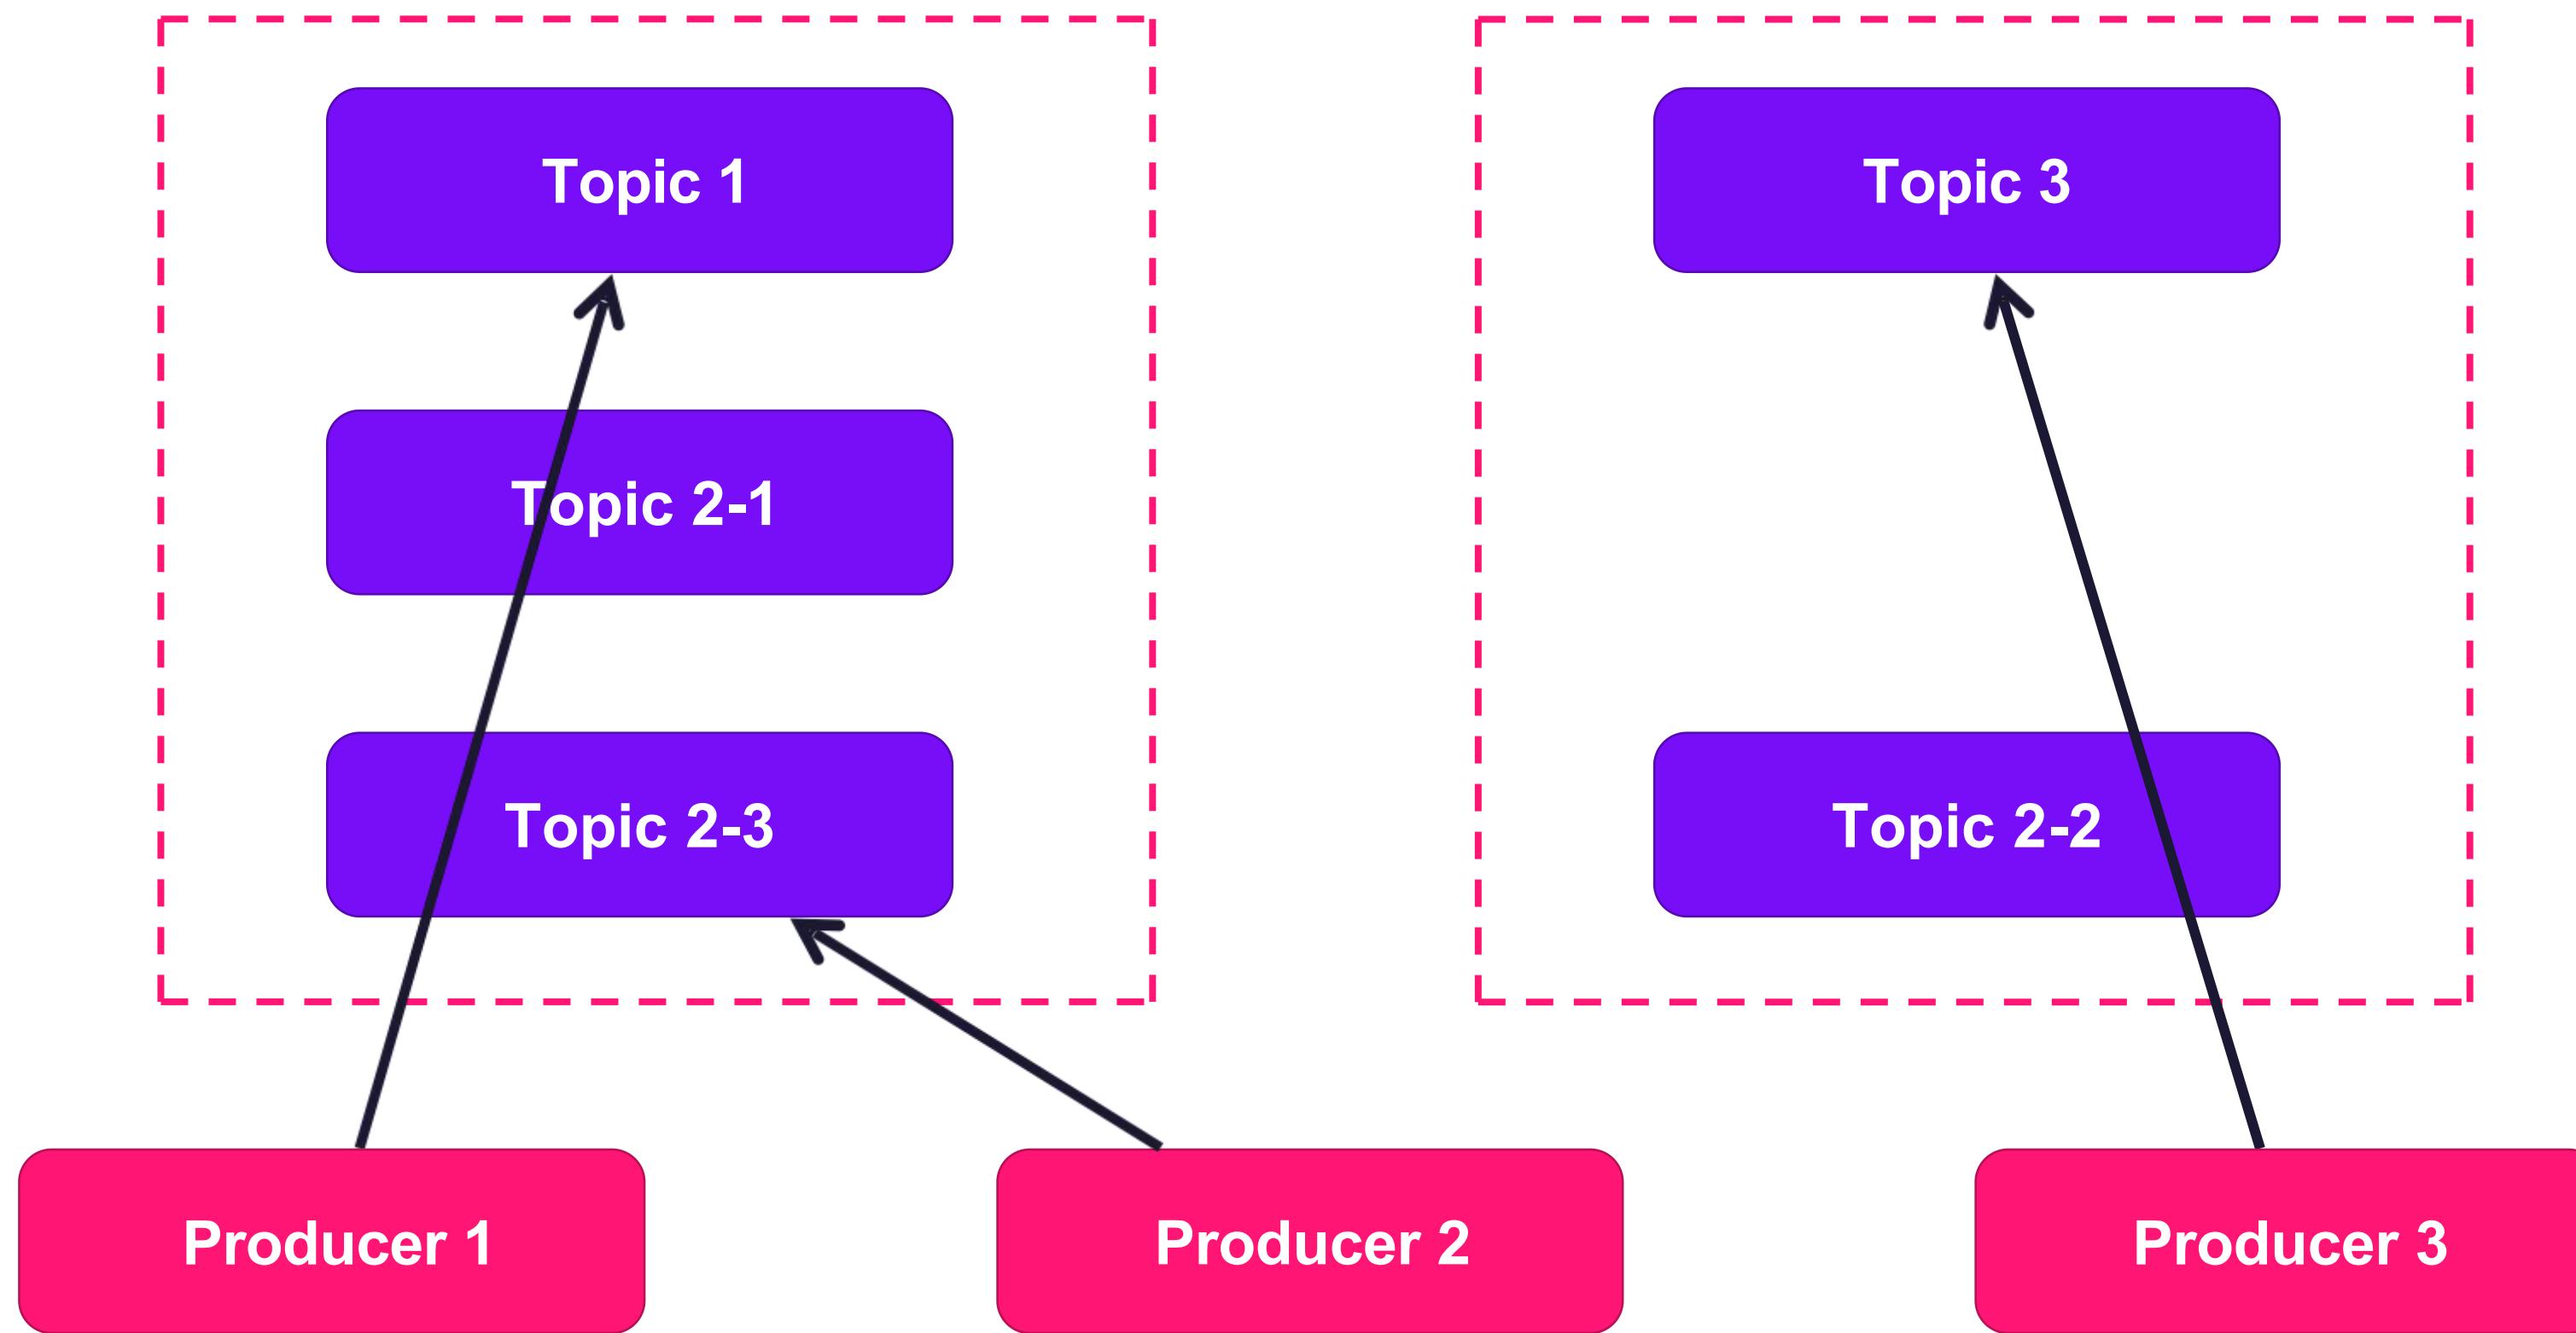
Partitions



Partitions

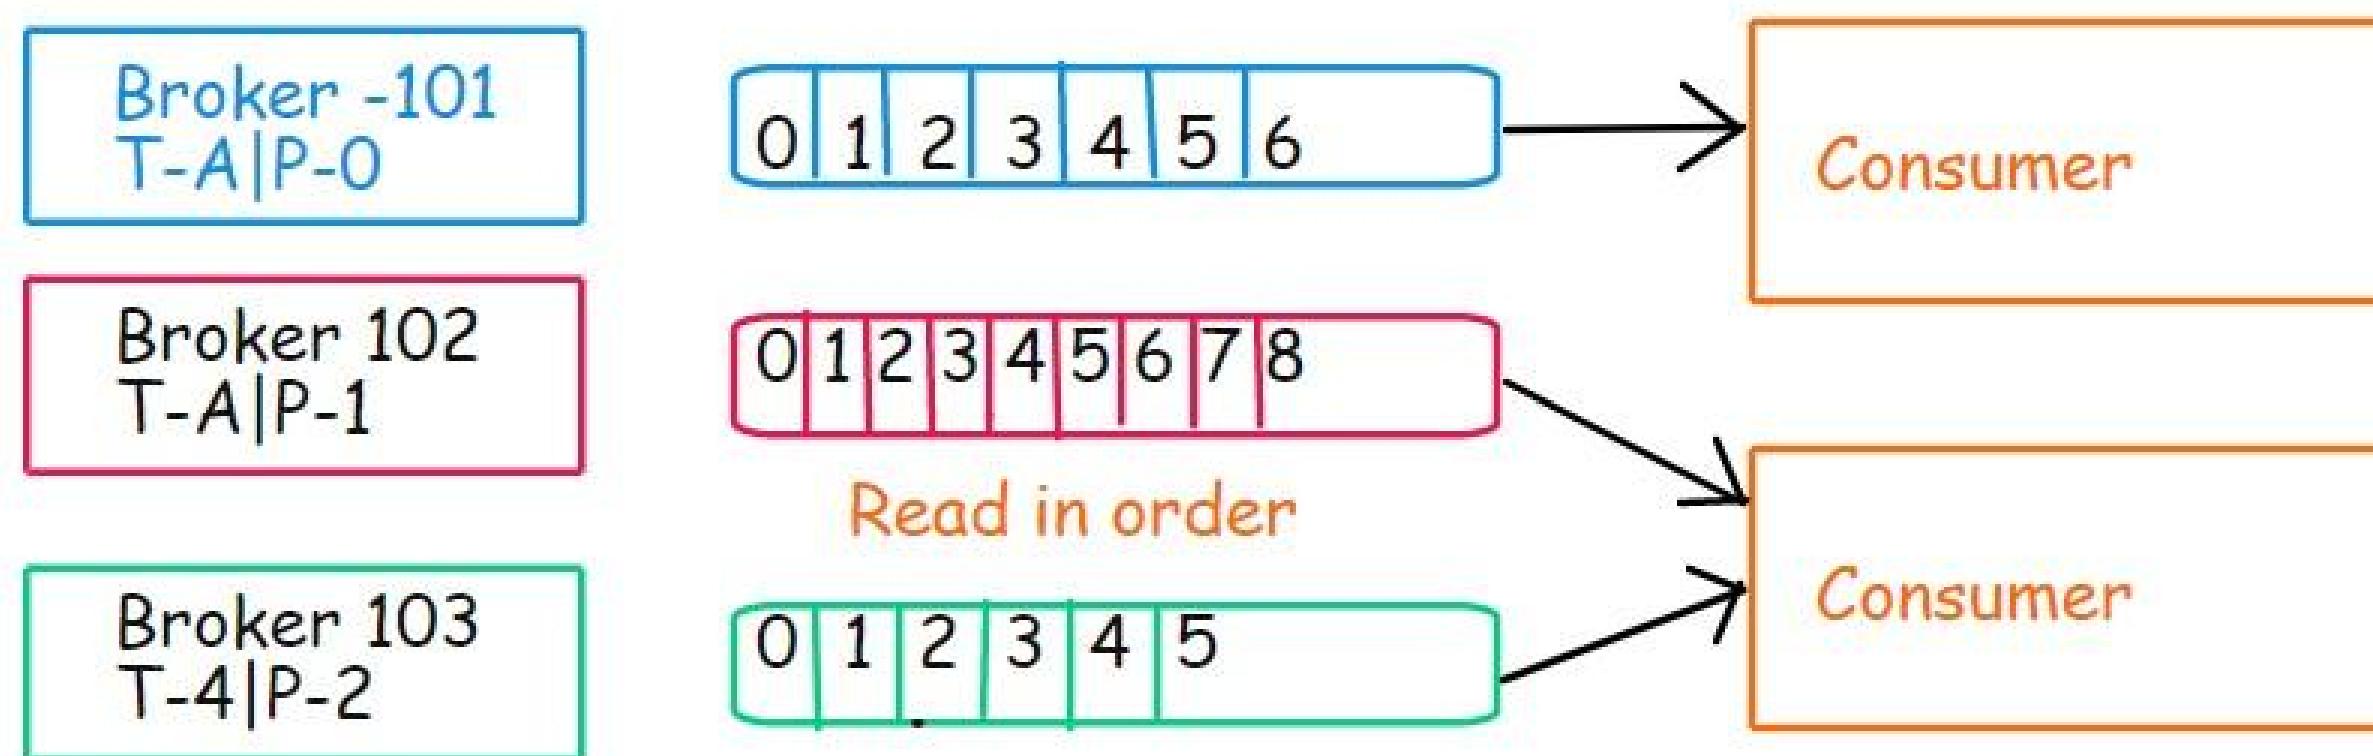


Partitions



Consumer

- Consumers read the data from a topic
- Consumer knows which broker to read from
- In case of broker failure ; consumer knows how to recover
- Data is read from each partition in an order

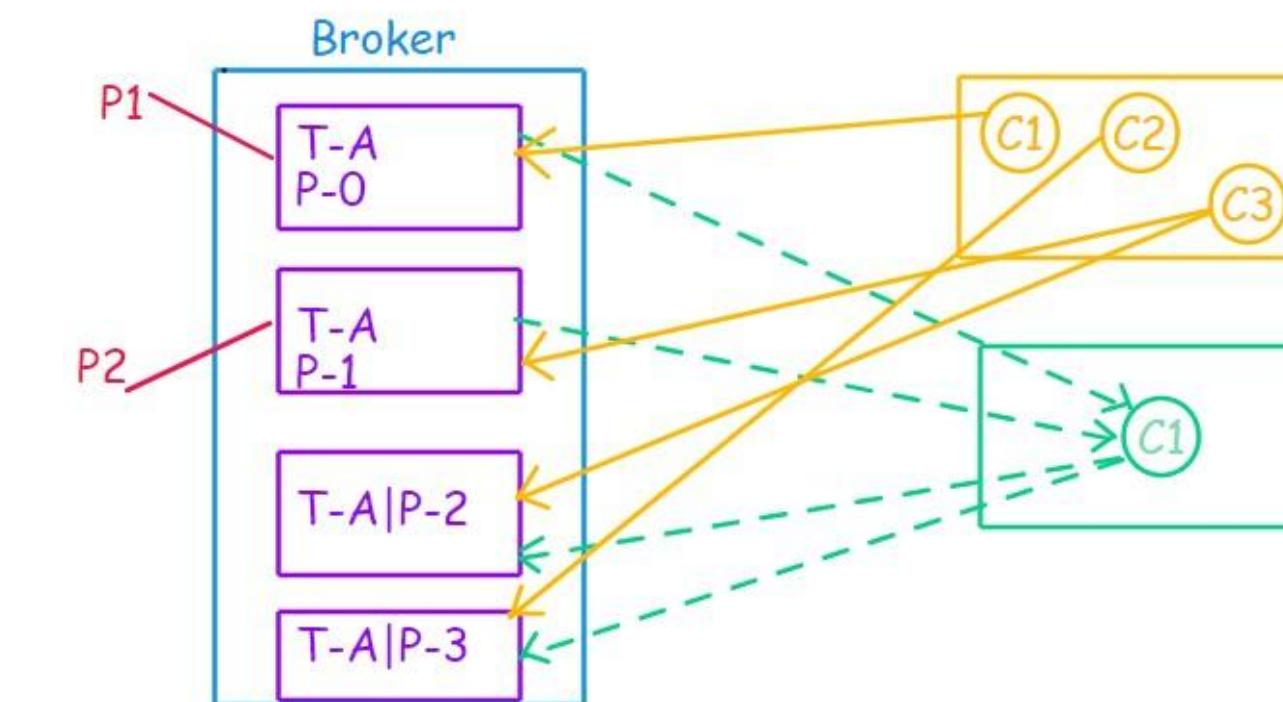
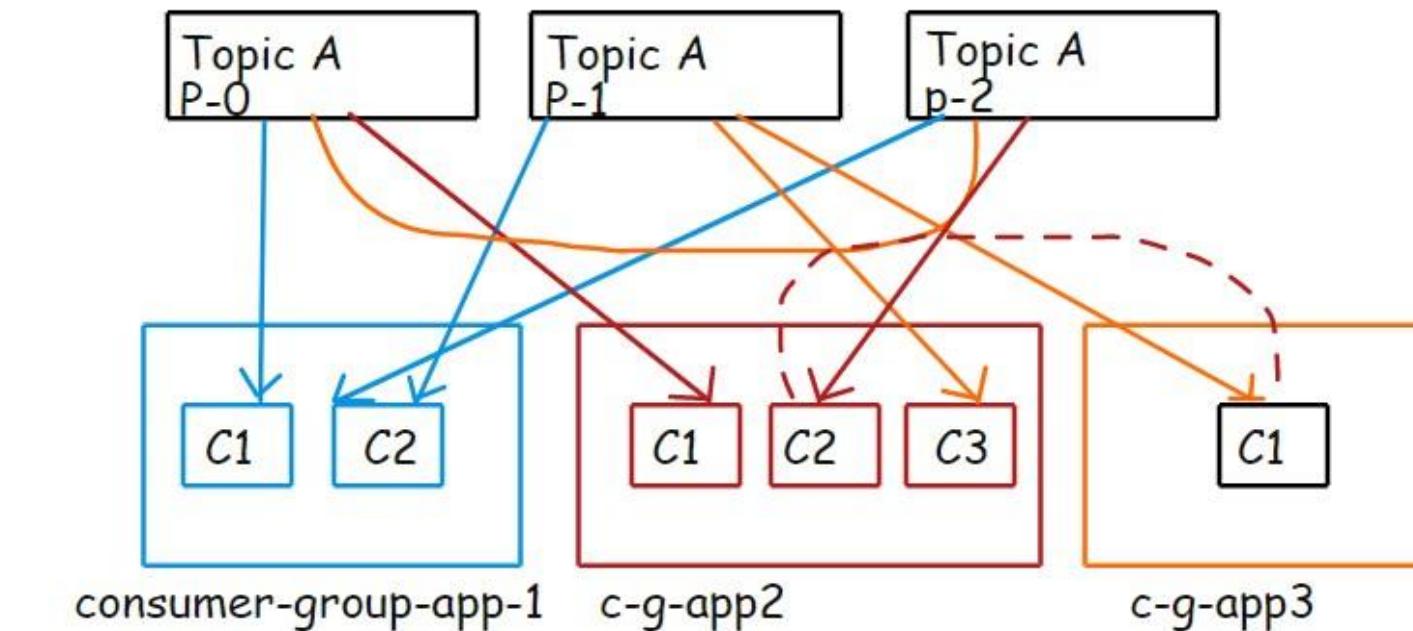


Consumer Group

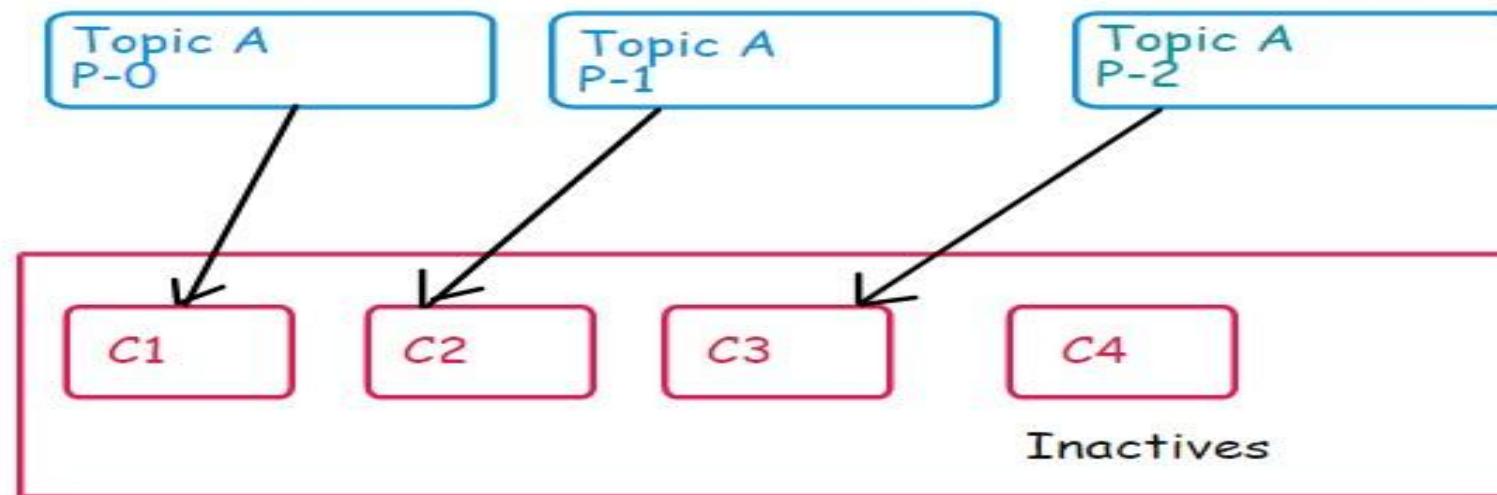
- Consumer reads the data in Consumer group
- Each customers within a group reads from a exclusive partition
- If you have more consumers than the partitions; some consumers will be inactive
- Consumer belonging to the same group shares a group_id
- Consumer in the group divides the topics partition fairly amongst themselves

↓

- Each partition is only consumed by a single consumer in that group



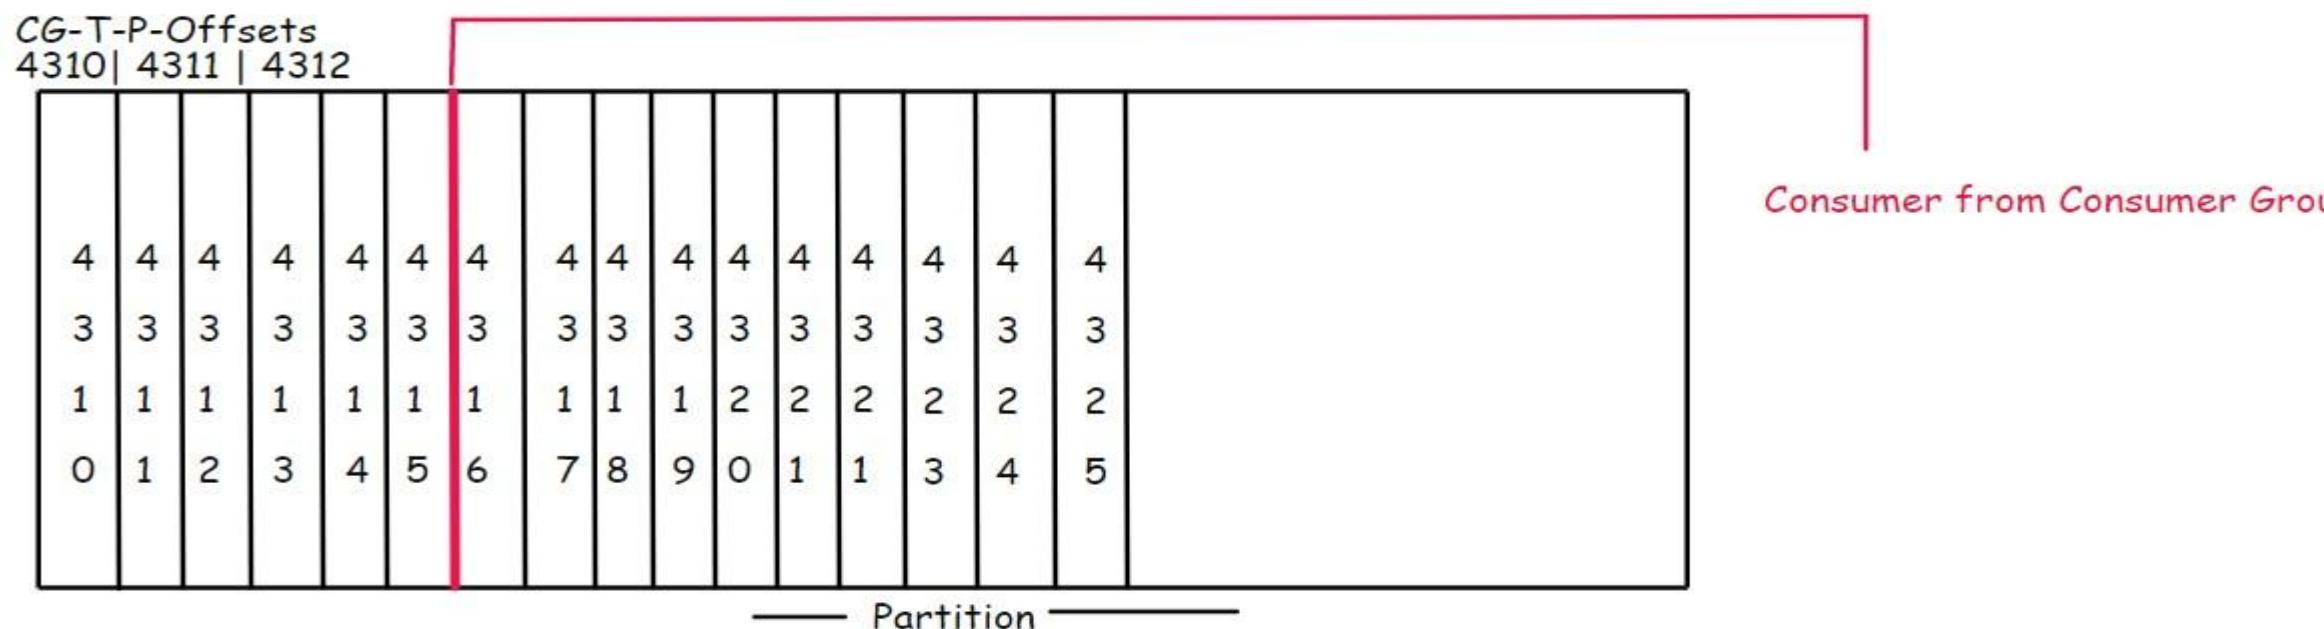
- What will happen if number of consumers are more than number of partitions?



- **Rebalancing** :- When a consumer group scales up or scales down Consumers in the group split the partition amongst themselves
- Rebalancing is triggered by a shift in ownership between the partition & consumer
- CG → (C1 , C2 ,C3) → C2 (suffers a failure C1 & C3 will briefly pause consumption of messages from their partitions all the partitions will be up for reassignment

Consumer Offsets

- Kafka stores the offset at which consumer group has been reading
- The offset committed → they are stored in a Kafka topic `_consumer_offsets`
- When a Consumer group have processed the data received from Kafka : it commits the offsets

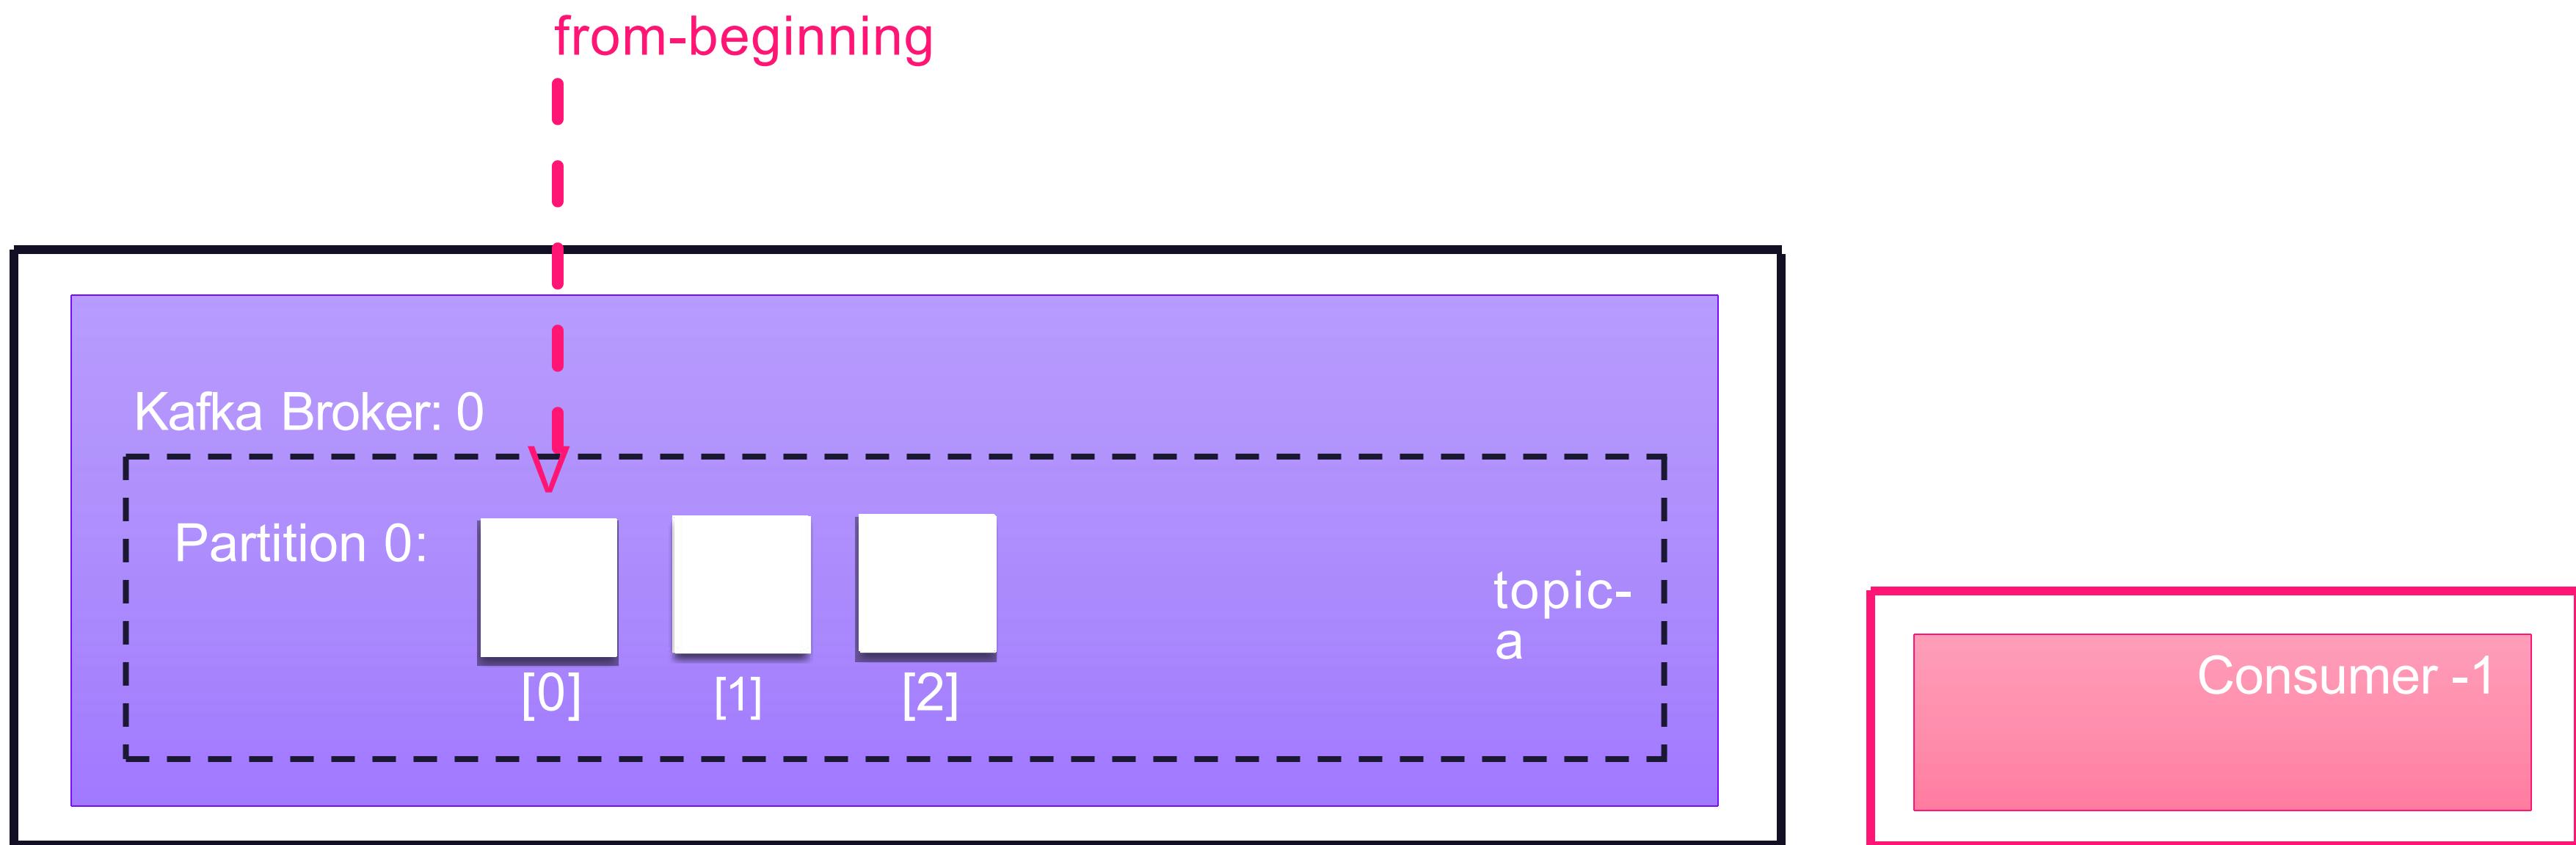


- If a consumer is down; it will be able to read from where it left

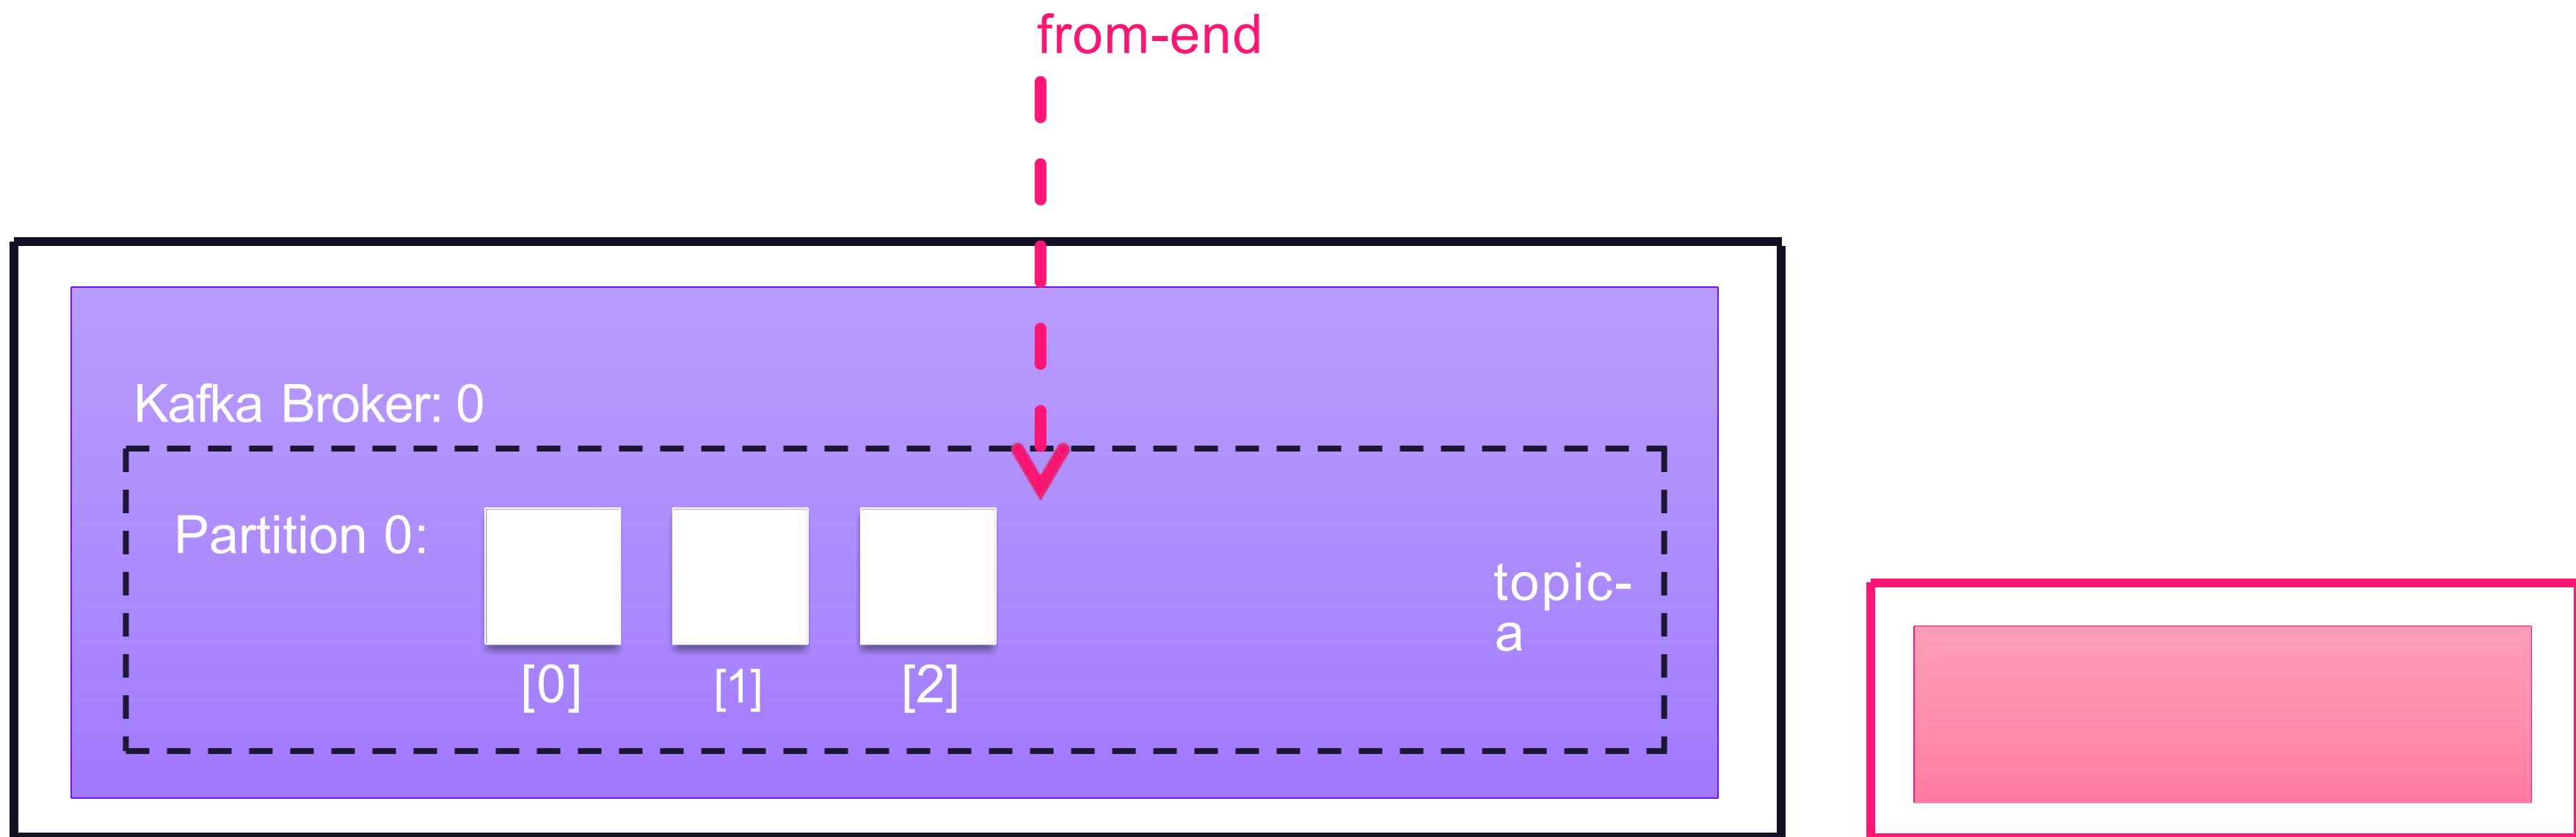
Delivery Semantics for Consumer

- Consumer can choose when to commit offset
- There are 3 delivery semantics
 - At most once
 - ❖ Offsets is committed as soon as the message is received
 - ❖ If processing goes wrong; message will be lost
 - At least once (**usually preferred**)
 - ❖ Offset is committed after the message is processed
 - ❖ If the processing goes wrong; message will be read again
 - ❖ **Idempotent** (if process is done multiple time it doesn't affect the state of application)

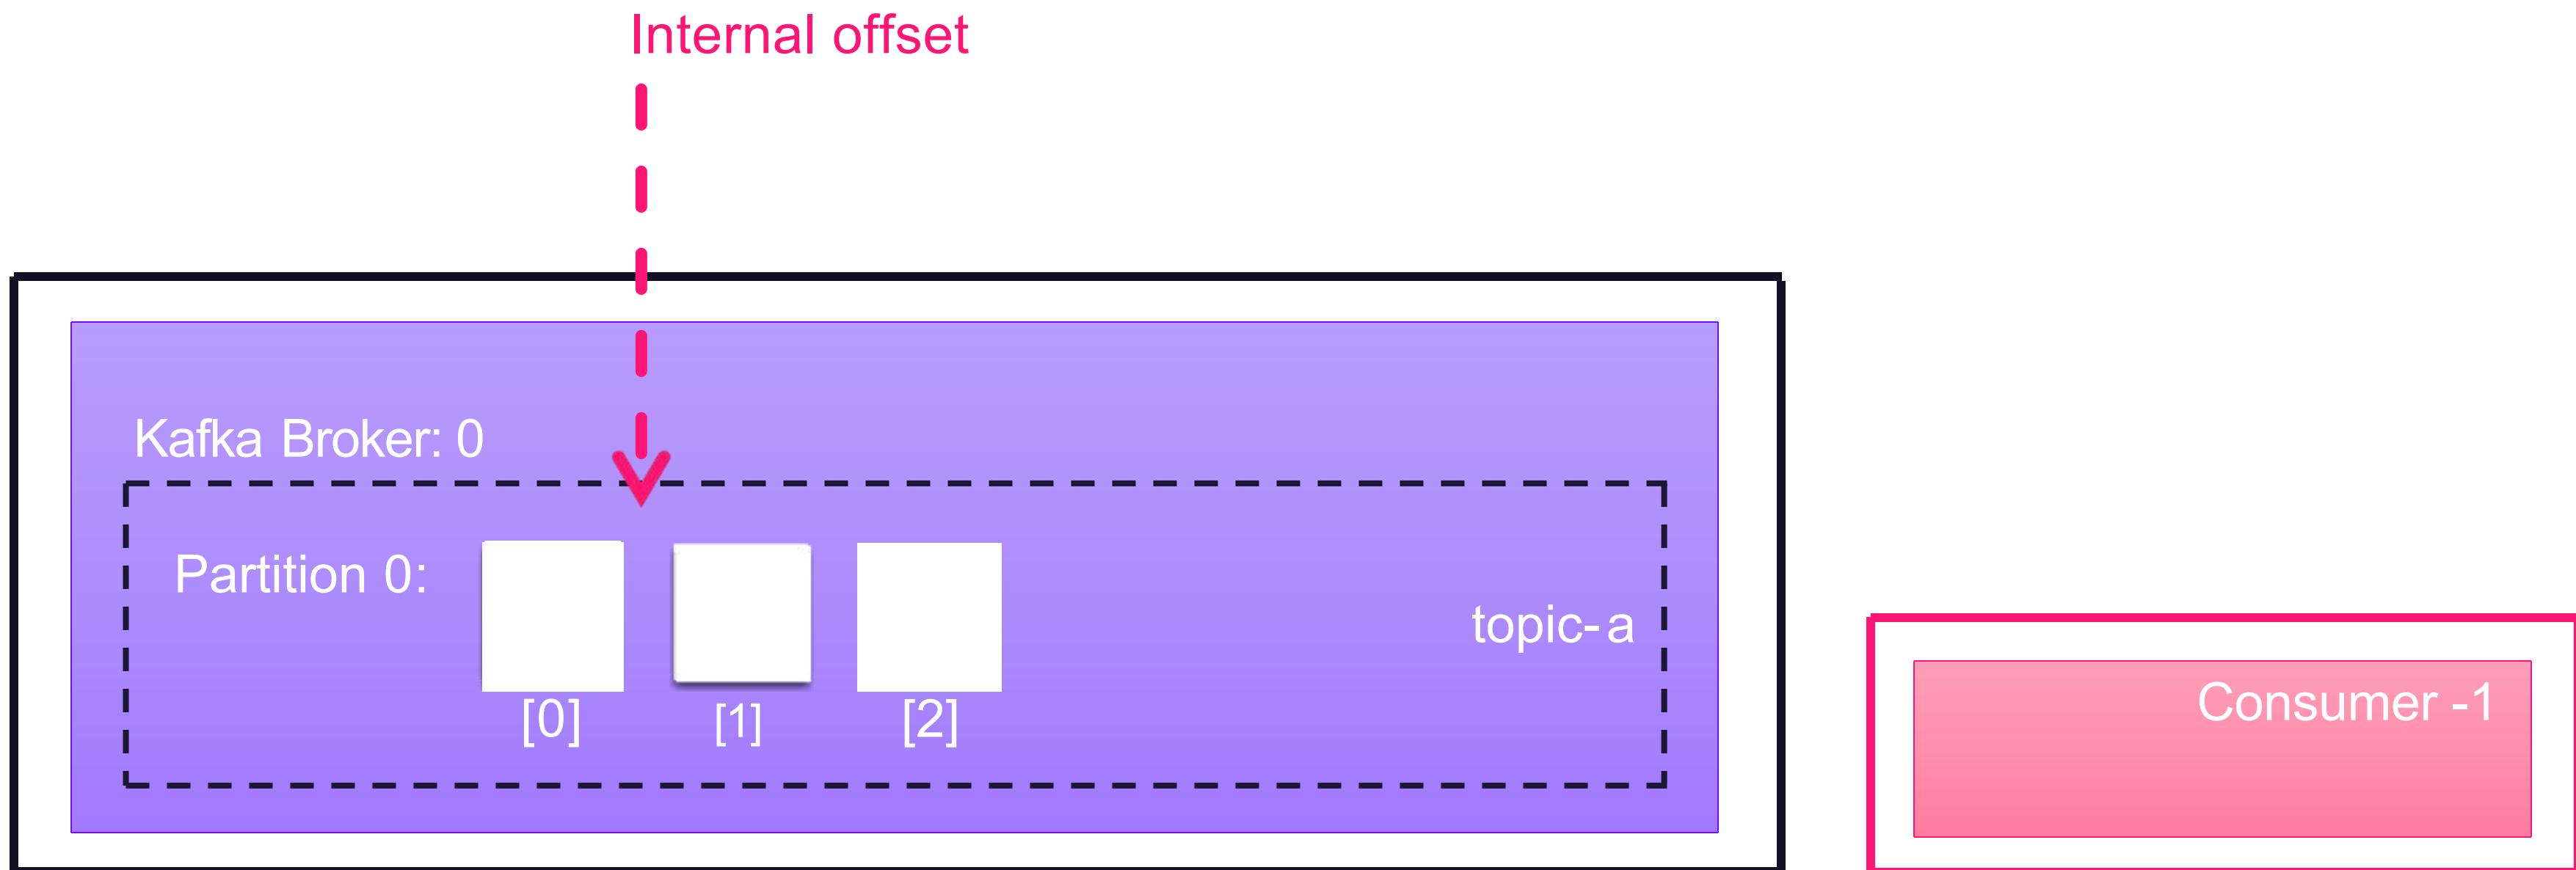
Consumer Offset: From Beginning



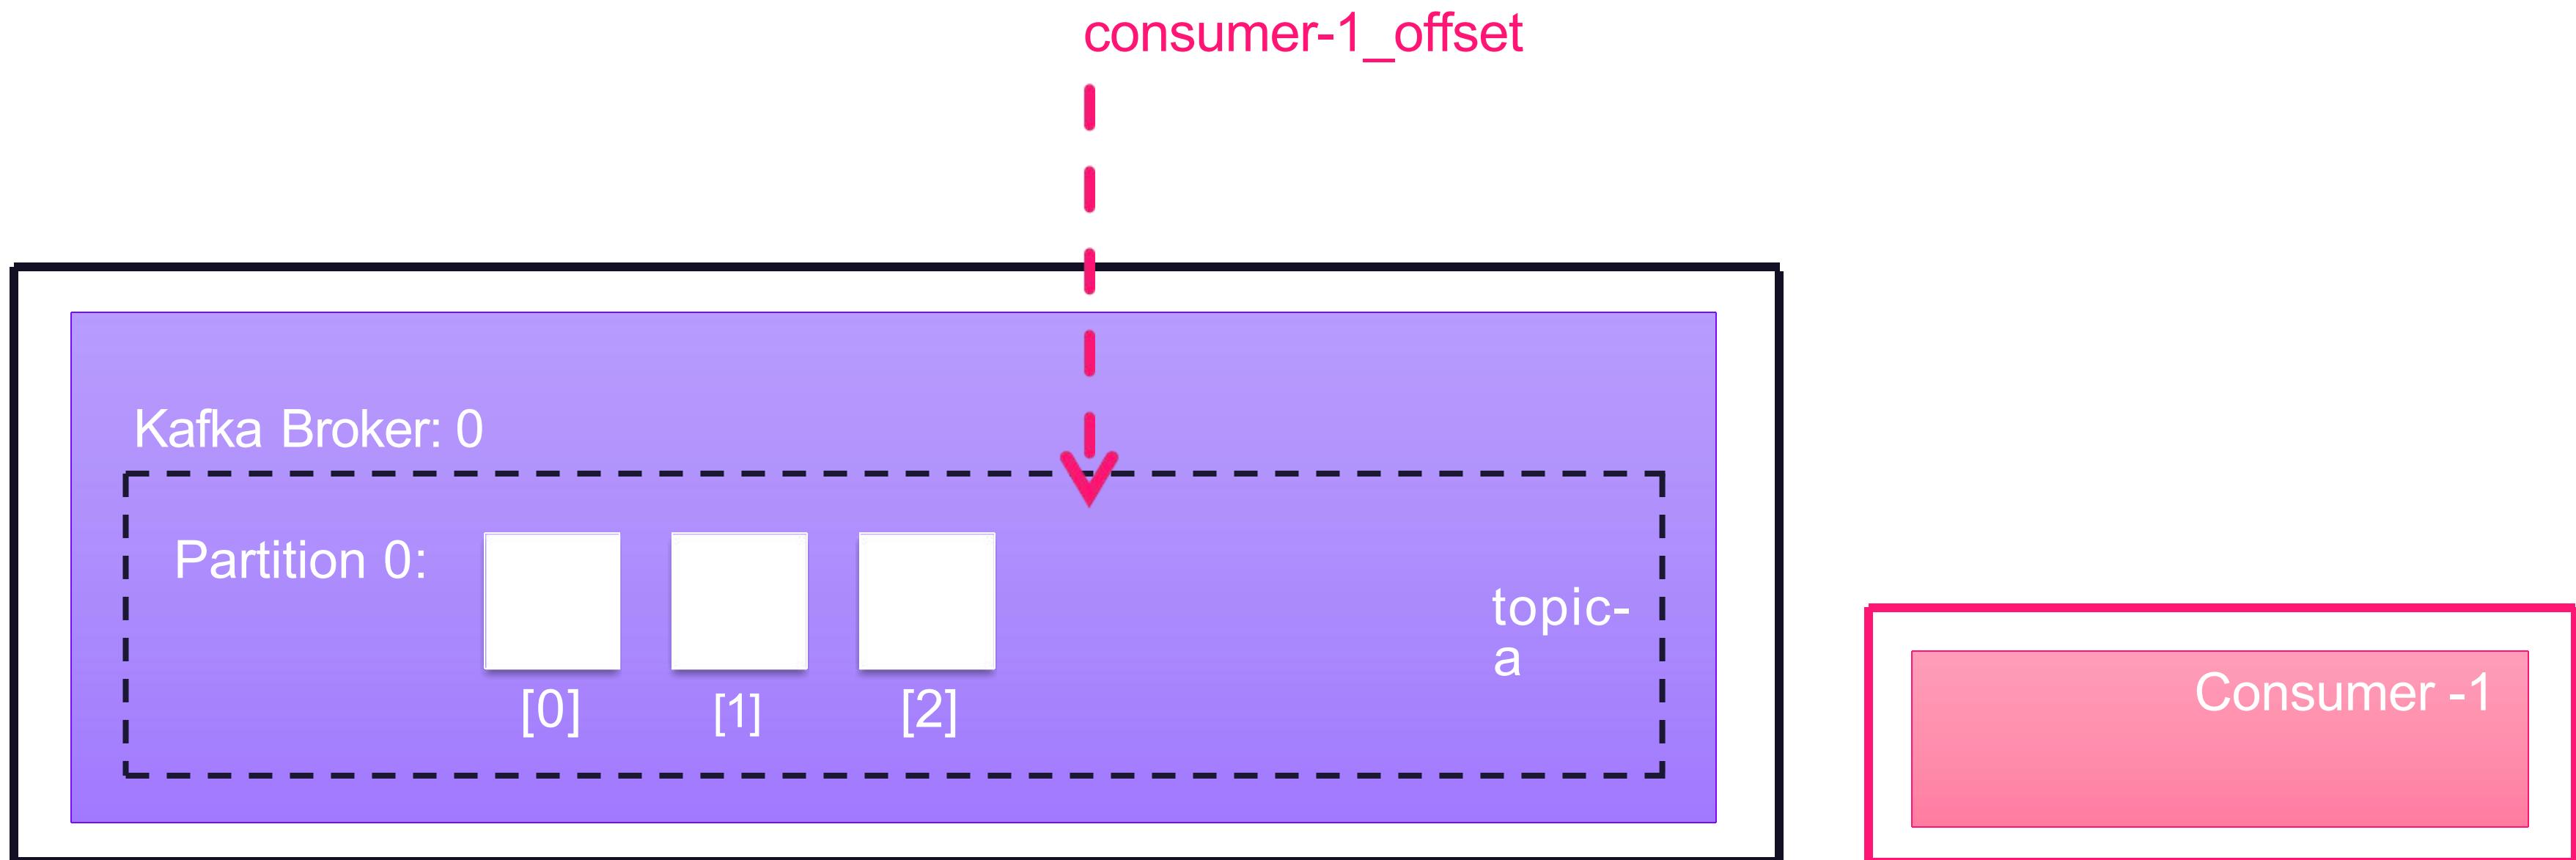
Consumer Offset: From End



Consumer Offset: Already Have Been Here

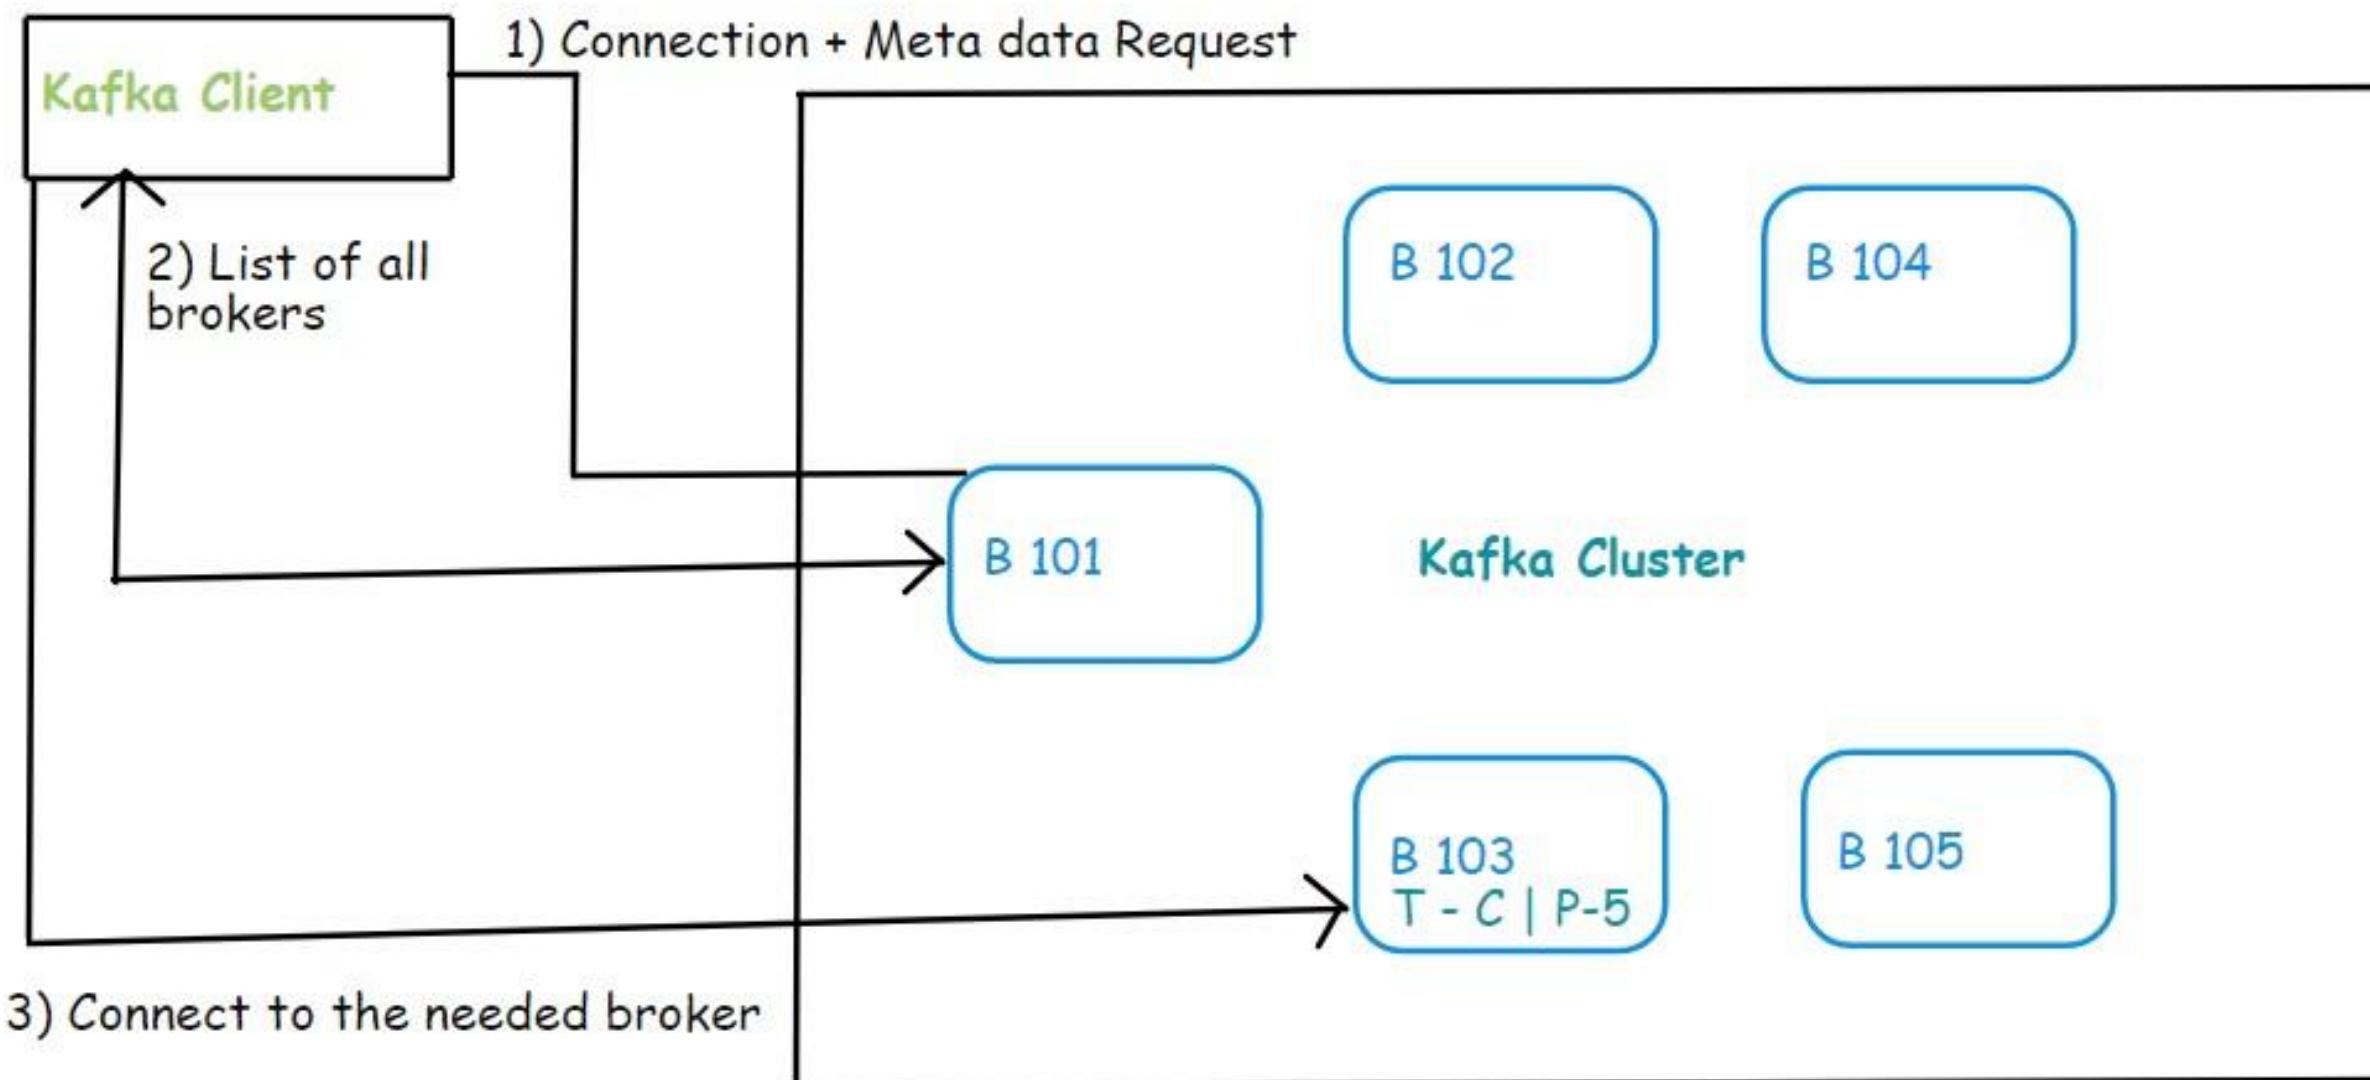


Consumer Offset



Kafka Broker Discovery

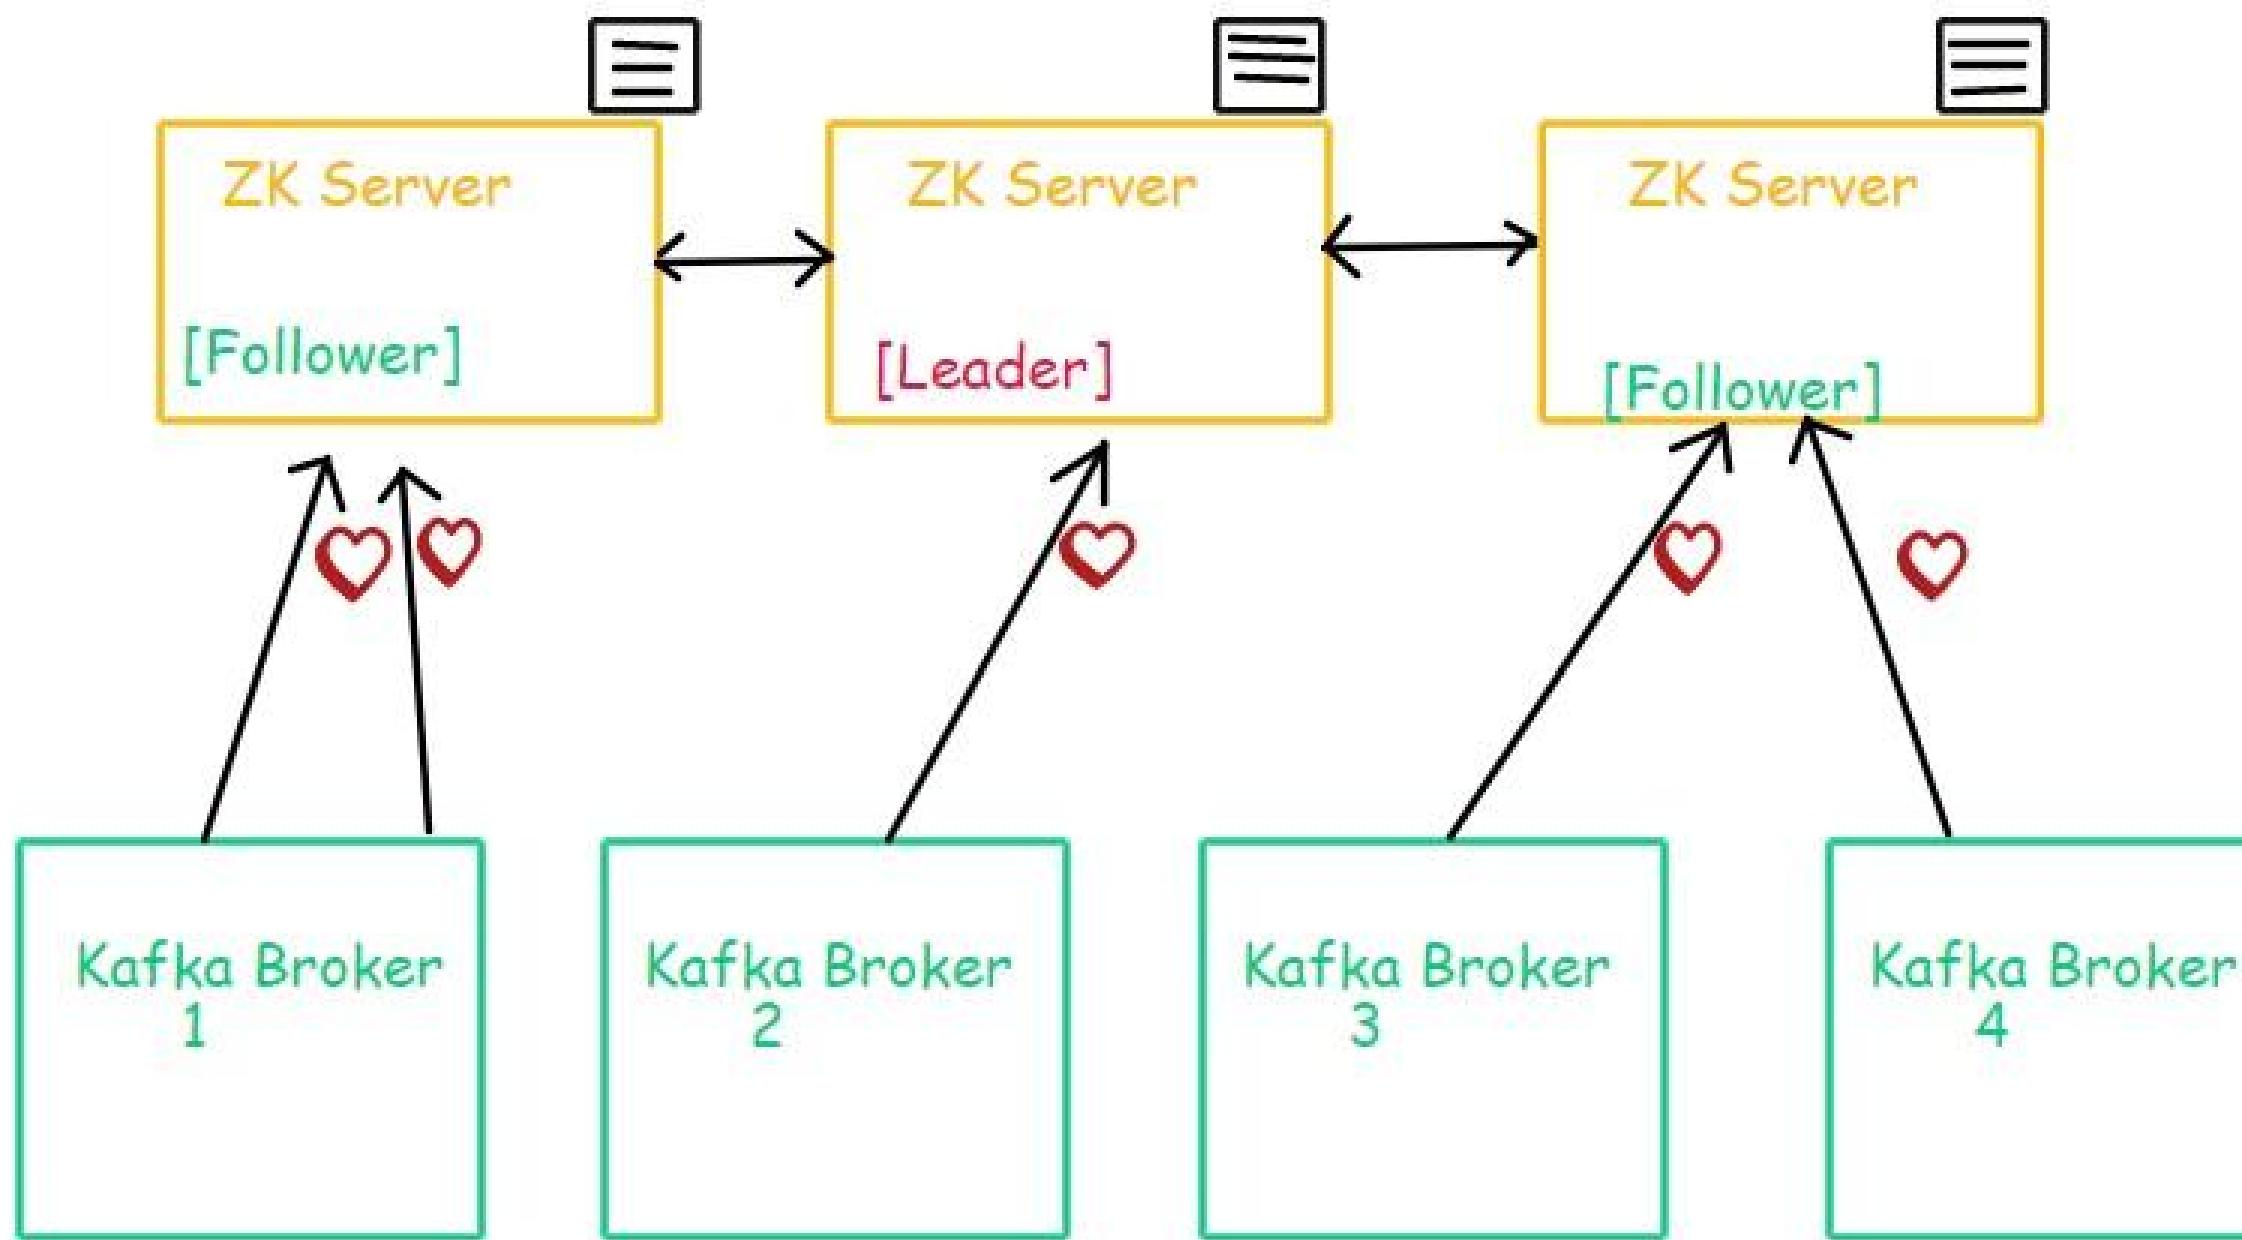
- Every Kafka Broker is called as “Bootstrap Server”
- You need to connect to any one broker & you will be connected to entire cluster



Zookeeper

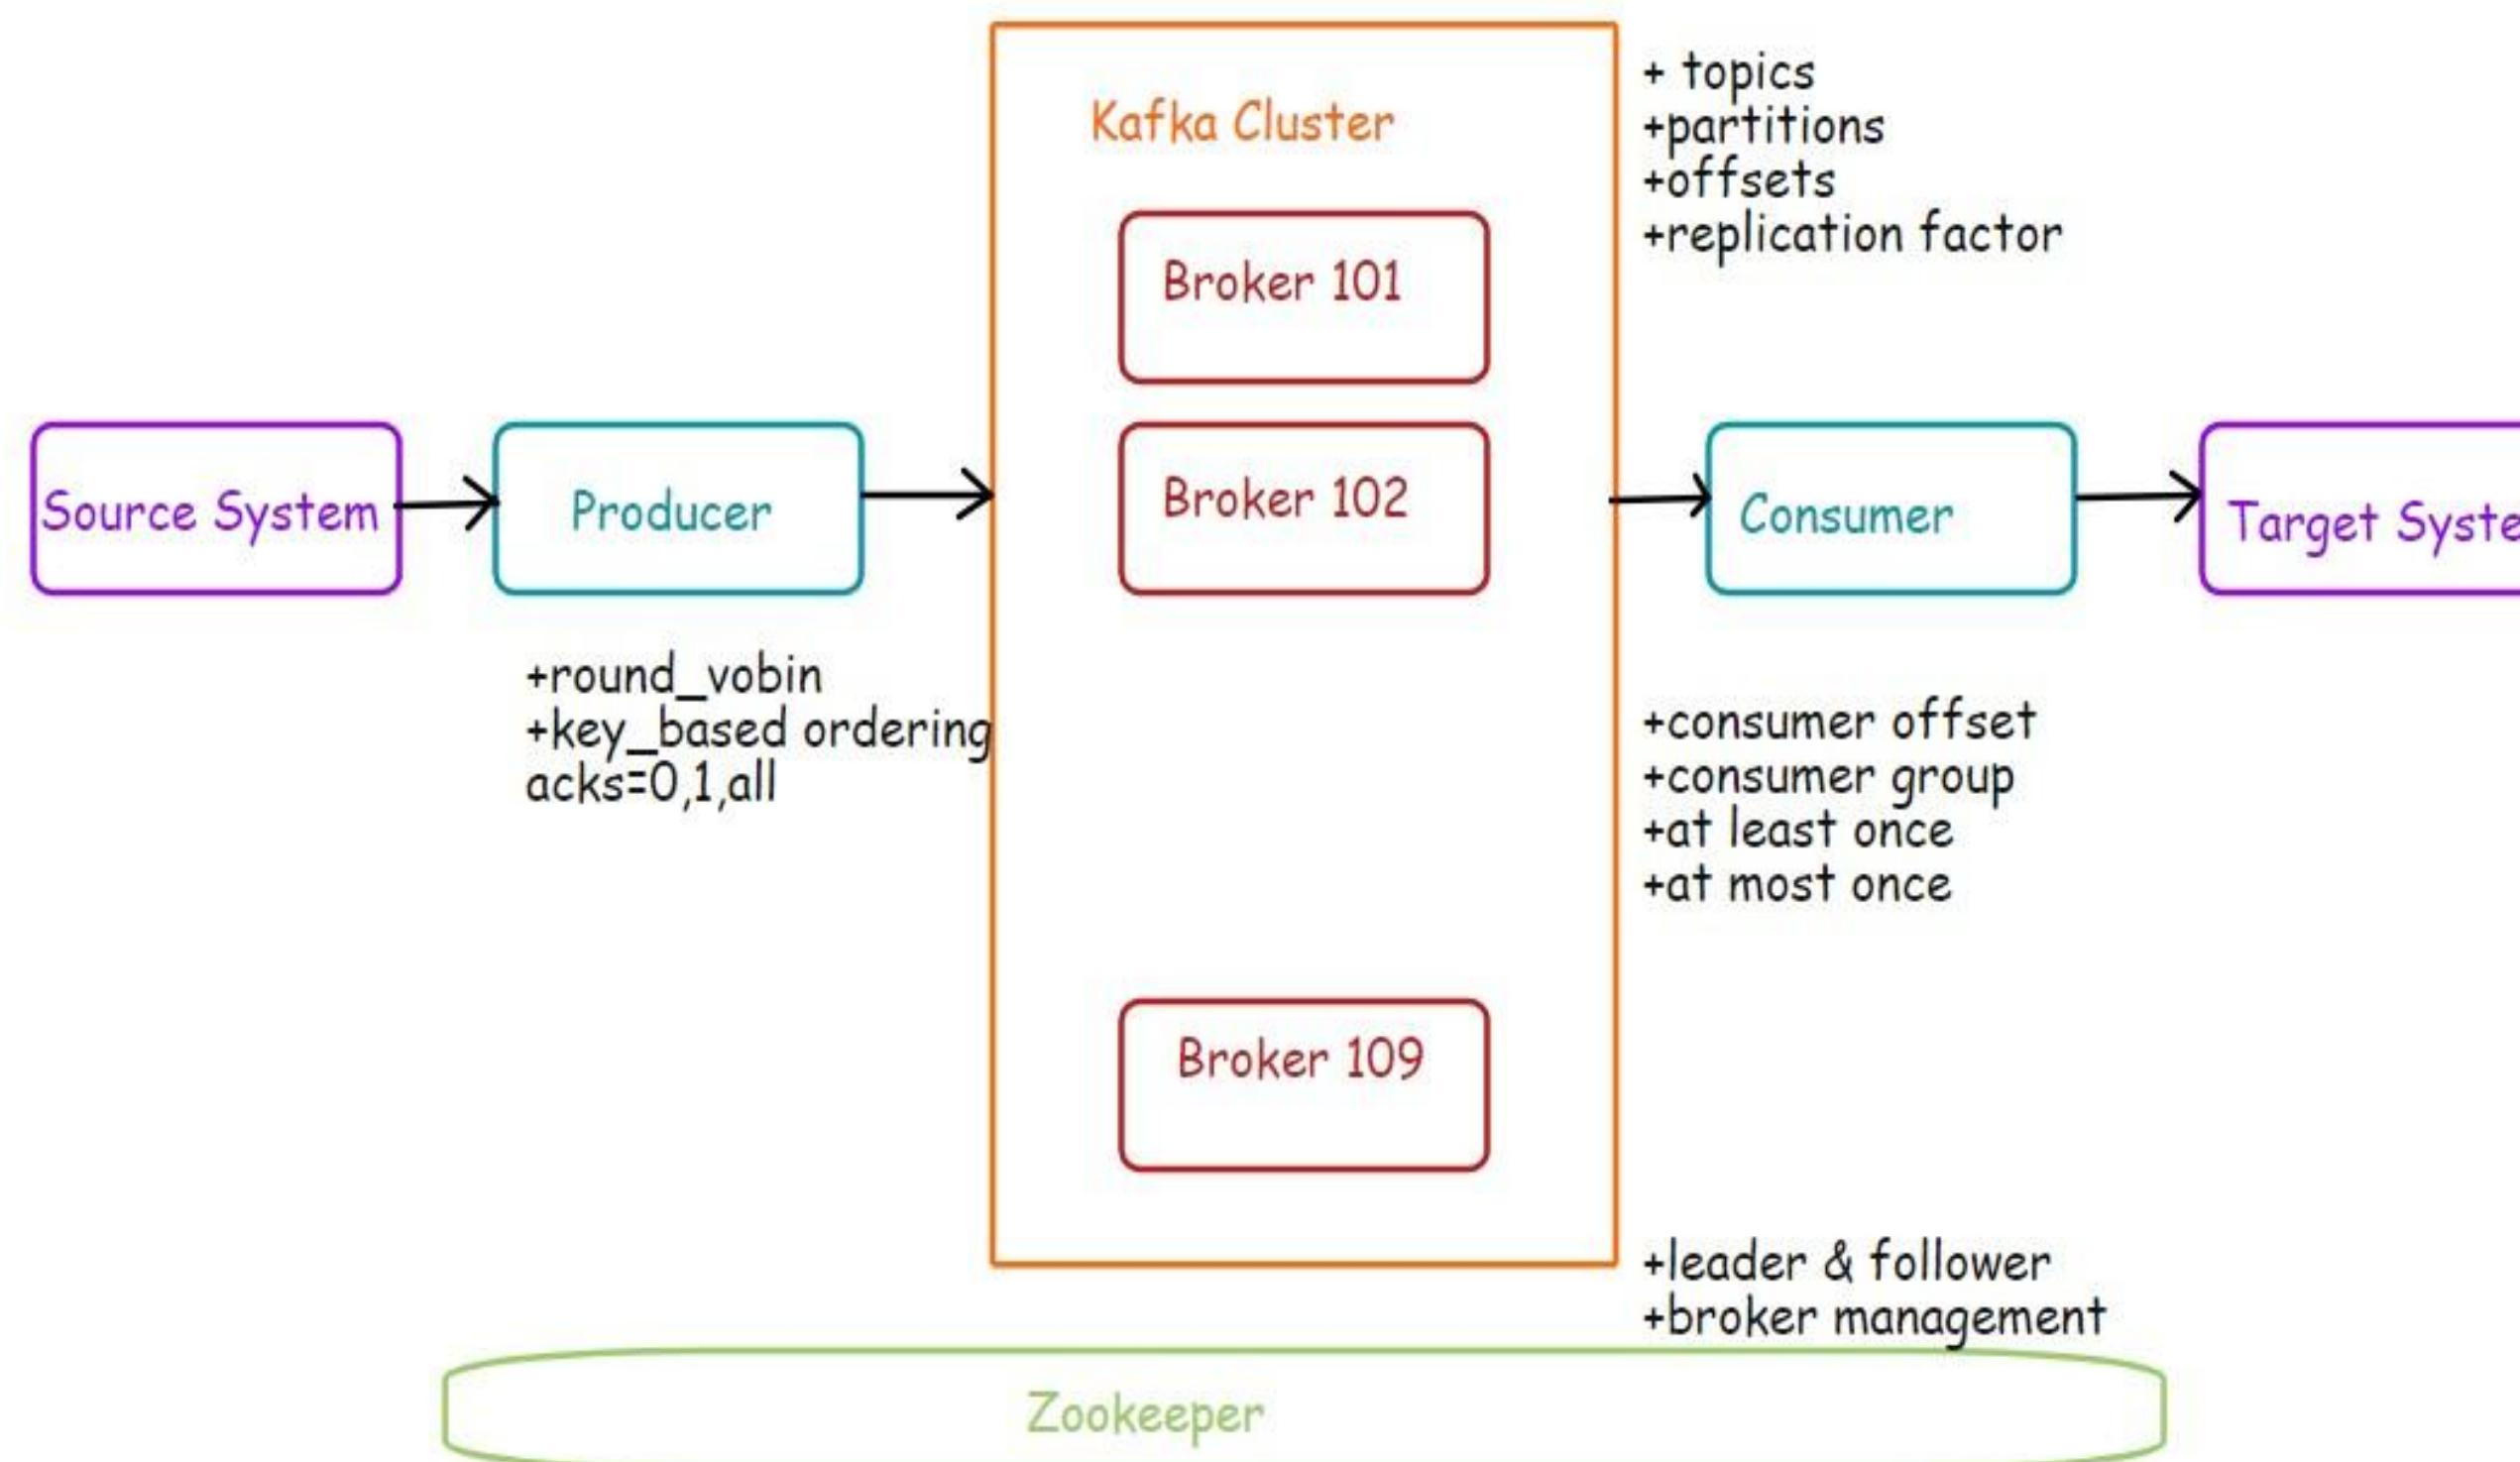
- Zookeeper manages the brokers (keeps a list of them)
- Zookeeper helps in performing leader elections for partitions
- Zookeeper sends notification to Kafka in case of
 - ❖ New topic
 - ❖ Broker dies
 - ❖ Broker comes up
 - ❖ Deleted Topic
 - ❖ Etc.
- Zookeeper operates with odd number of servers

- Zookeeper has a leader (handle write) & rest of the servers are followers(handle read)



Kafka Guarantees

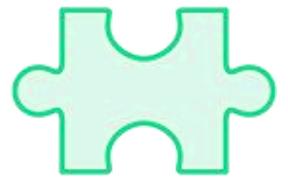
- Messages are appended to a topic-partition in order they are sent
- Consumer reads message in order stored in topic-partition
- With $RF \rightarrow N$; producers and consumers can tolerate $N-1$ brokers being going down
- RF-3 is a good idea
- As long as number of partitions remains constant; same key will go to same partition



In the Logging Area, for Example, We Have Plugin For:



syslog



fluentd



Logstash



Filebeat

DBs



JDBC connector to cover RDBMS Dos



Debezium also offers a connector that include no-SQL Dos



Presto connector that allows us to query Kafka using Presto queries

We Learned



Topics can be partitioned



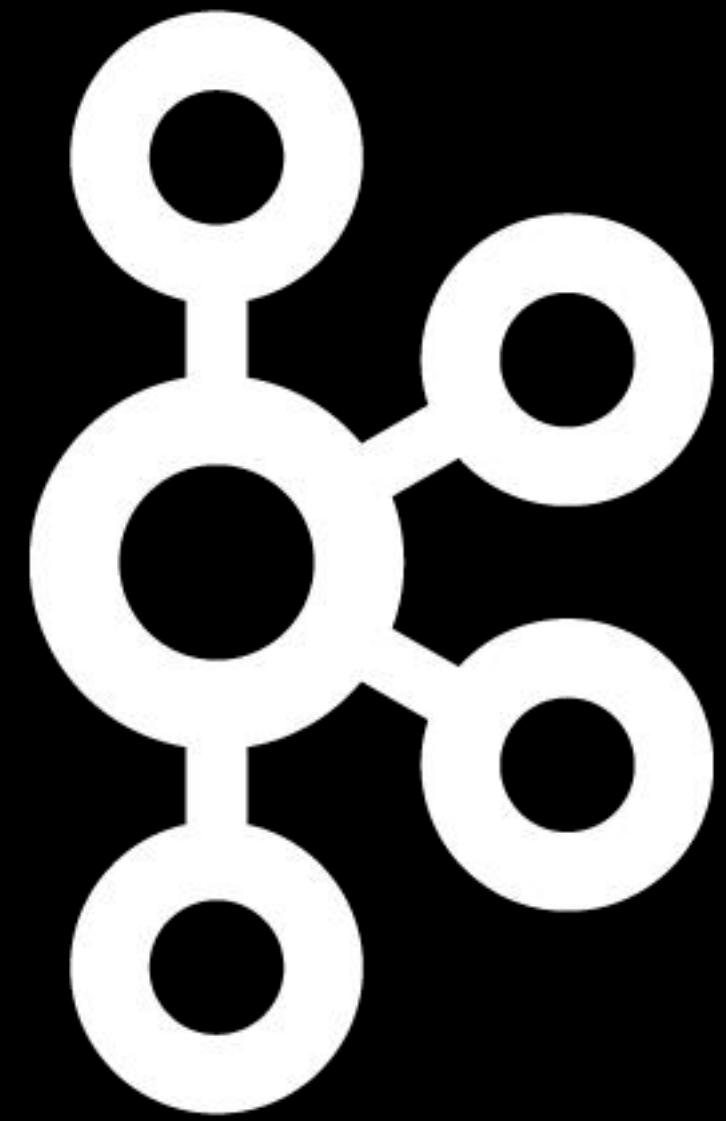
The partitioner uses the message key to redirect messages to partitions



Consumers read from partitions and have an offset



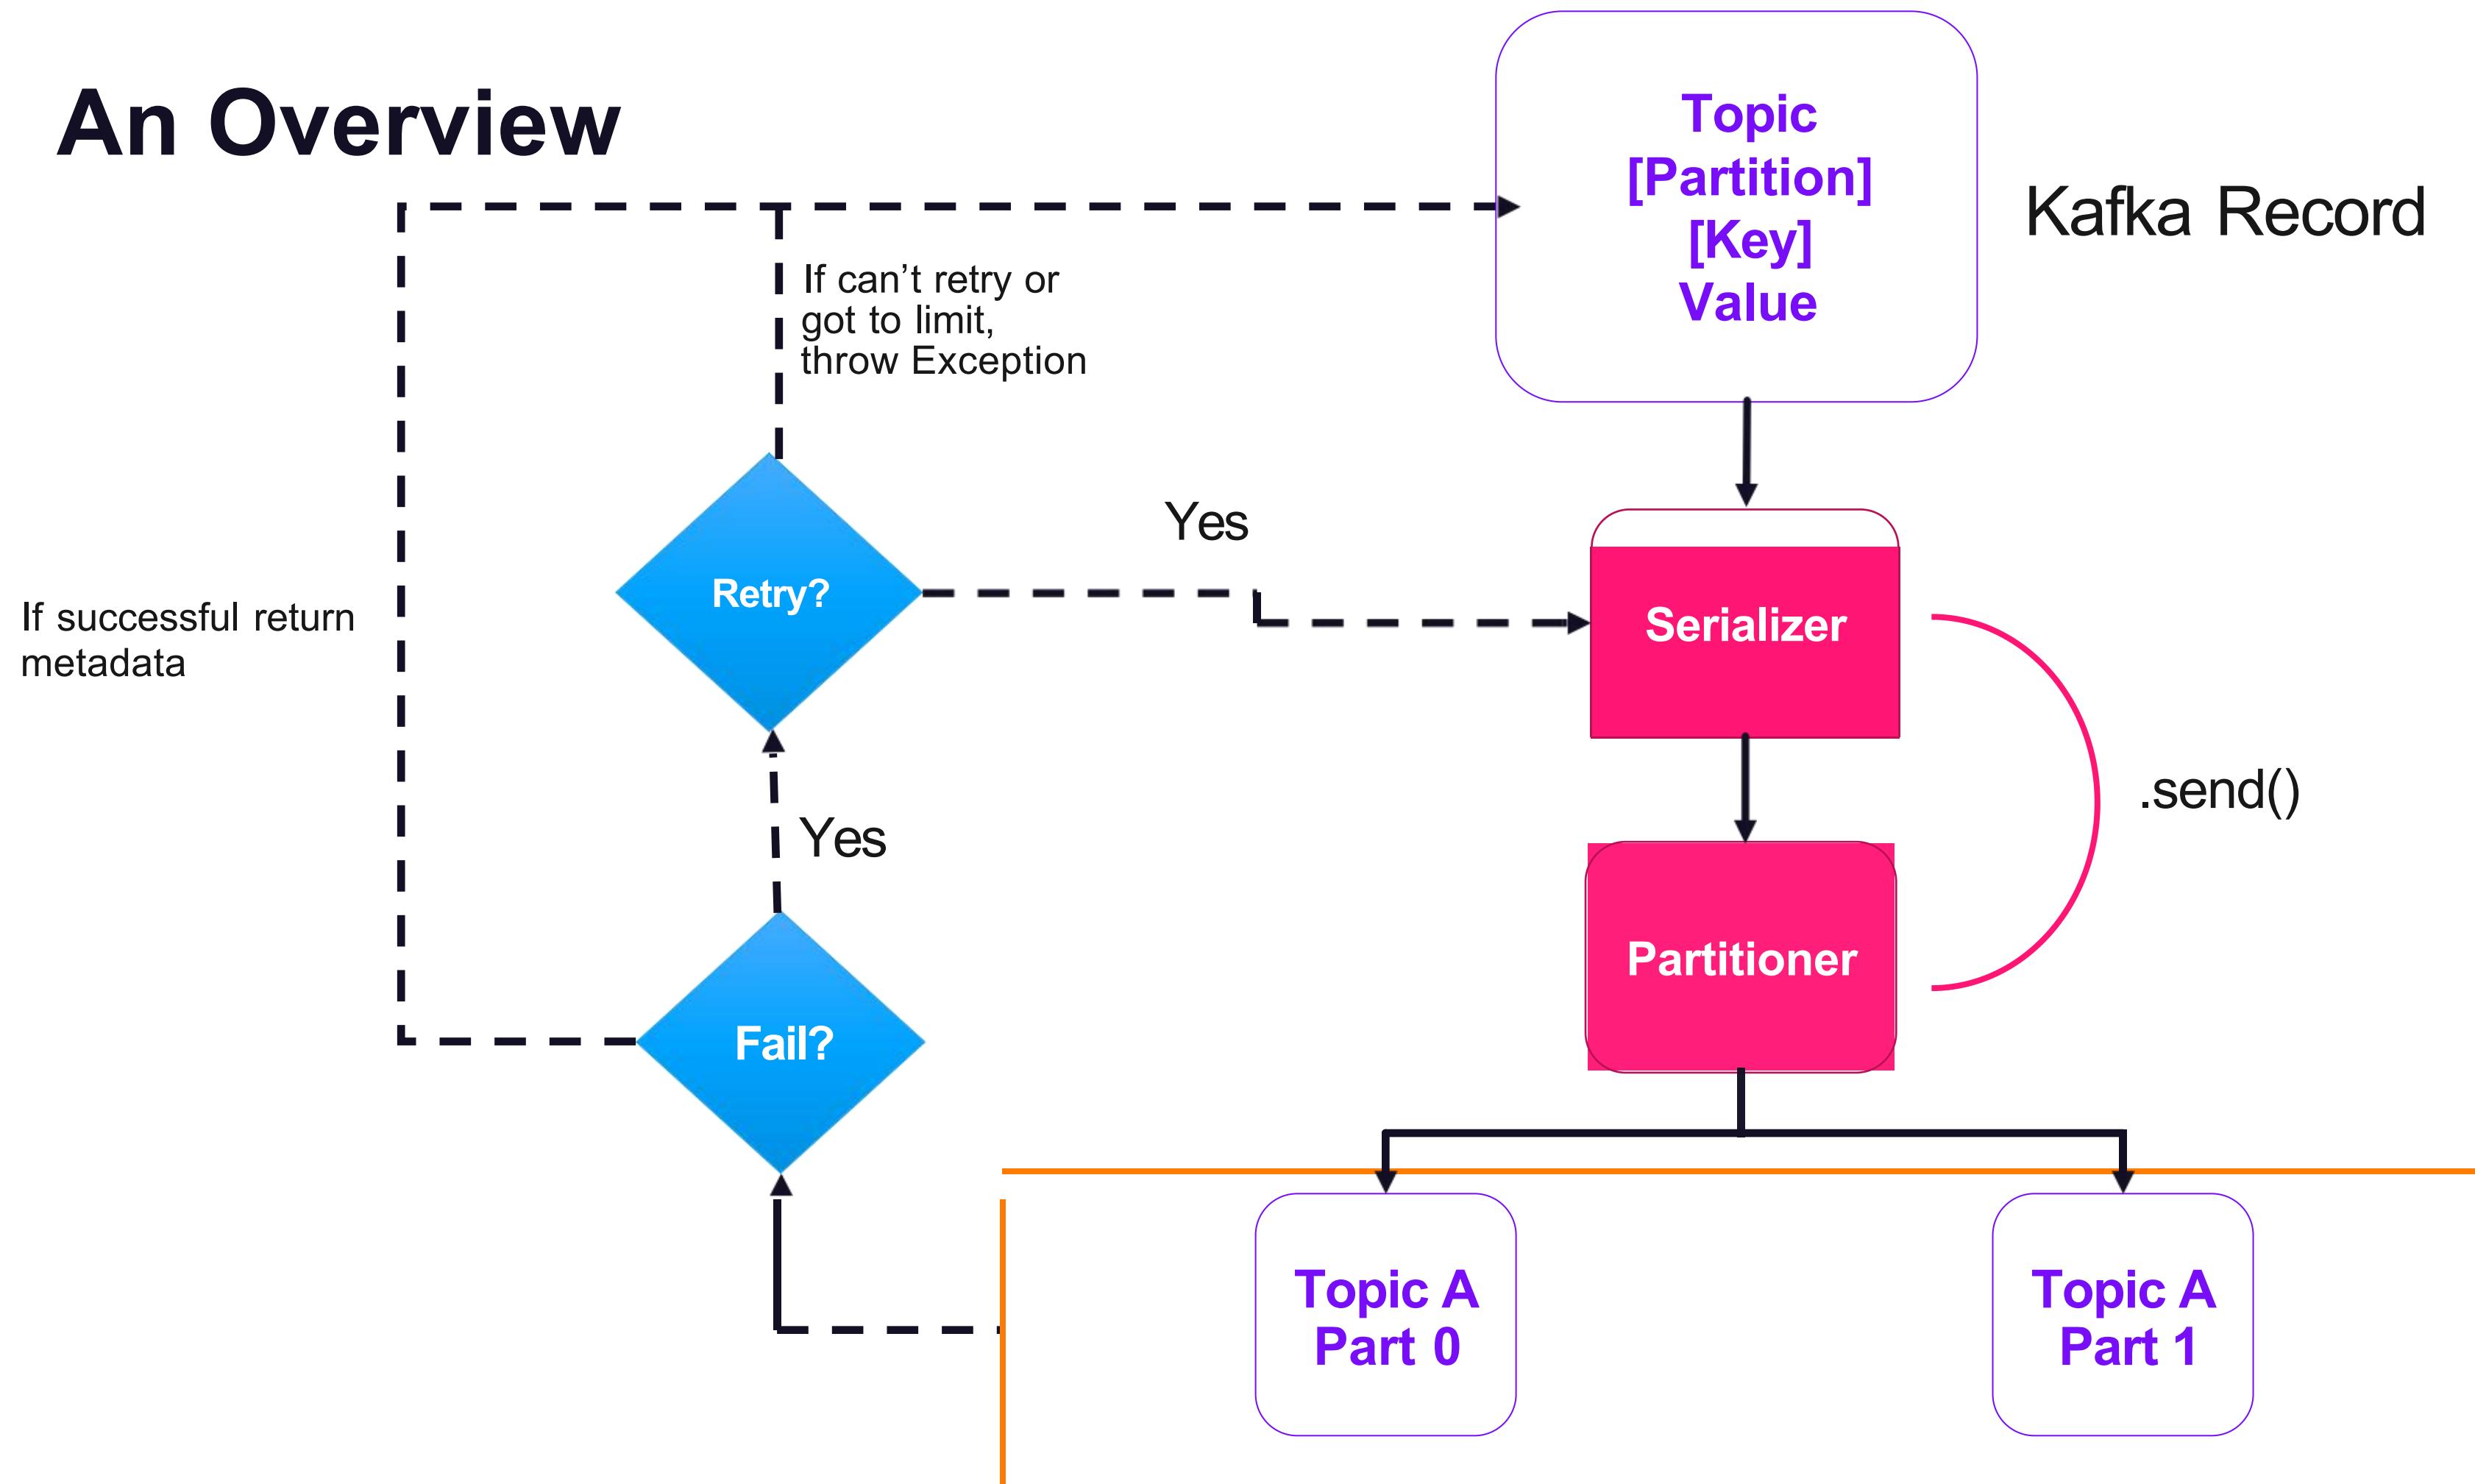
Therefore, more partitions enable more consumers, therefore more scalability

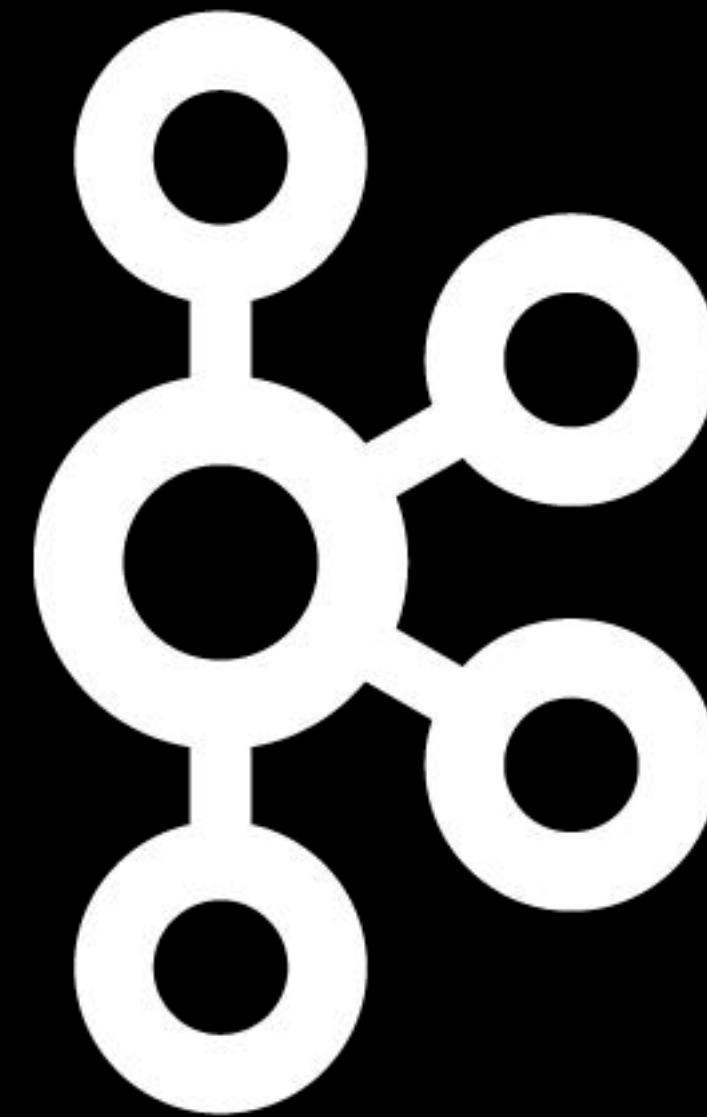


**Deploy
Kafka**

Kafka Producers

An Overview





Producing Messages with Kafka CLI

Serializers and Producer Configuration

Construct a `java.util.Properties` object

```
Properties properties = new Properties();  
  
properties.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG,  
"localhost:9092");  
  
properties.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,  
StringSerializer.class);  
  
properties.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,  
IntegerSerializer.class);
```

Provide two or more locations where the Bootstrap servers are located

```
Properties properties = new Properties();

properties.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG,
"localhost:9092");

properties.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
StringSerializer.class);

properties.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
IntegerSerializer.class);
```

Provide a Serializer for the key

```
Properties properties = new Properties();

properties.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG,
"localhost:9092");

properties.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
StringSerializer.class);

properties.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
IntegerSerializer.class);
```

Provide a Serializer for the value

```
Properties properties = new Properties();

properties.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG,
"localhost:9092");

properties.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
StringSerializer.class);

properties.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
IntegerSerializer.class);
```

Producer Object

```
KafkaProducer producer = new KafkaProducer<>(properties);
```

Create a Record

```
ProducerRecord<T> producerRecord = new  
ProducerRecord<T>("my_orders", state, amount);
```

Sending a Message

```
Future send = producer.send(producerRecord);
```

**Contains information
about your send
including the messages**

Record Metadata

```
if (metadata.hasOffset()) {  
    System.out.format("offset: %d\n",  
        metadata.offset()),  
}  
  
System.out.format("partition: %d\n",  
    metadata.partition());  
  
System.out.format("timestamp: %d\n",  
    metadata.timestamp());  
  
System.out.format("topic: %s\n",  
    metadata.topic());  
  
System.out.format("toString: %s\n",  
    metadata.toString());
```

Capturing a Callback

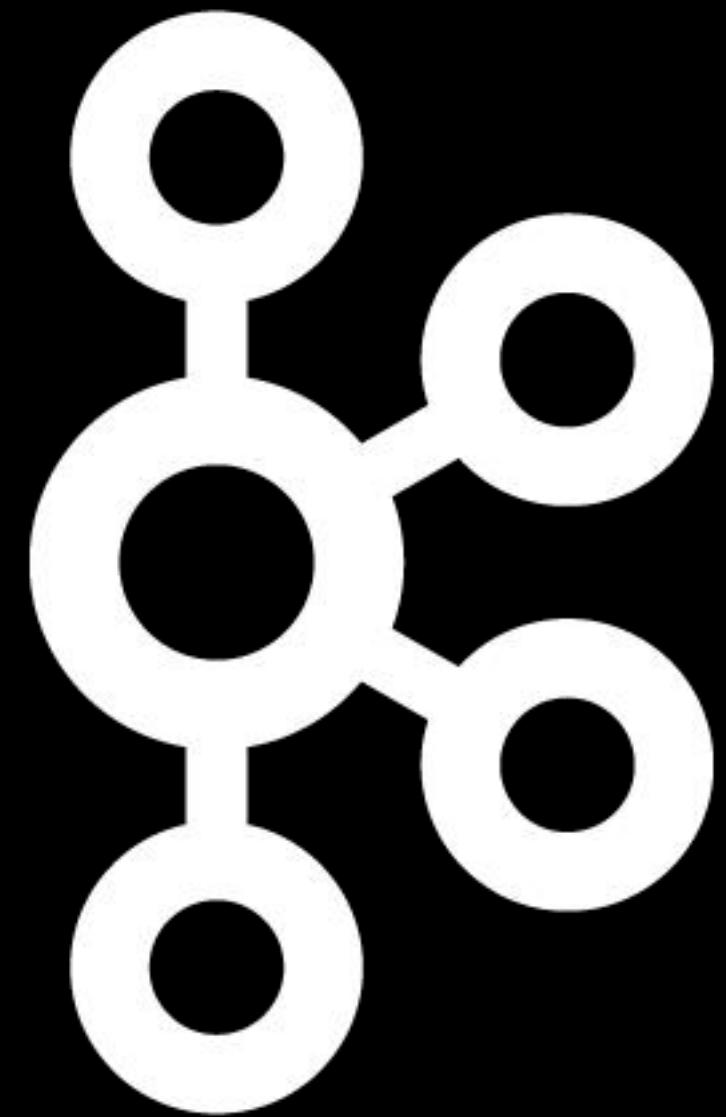
```
producer.send(producerRecord, new Callback() {  
  
    @Override  
  
    public void onCompletion(RecordMetadata metadata,  
                            Exception e){  
  
        ...  
  
    }  
  
})
```

Using Lambdas

```
producer.send(producerRecord,(metadata, e) -> {  
    if(metadata != null) {  
        System.out.println(producerRecord.key());  
        System.out.println(producerRecord.value());  
    }  
}
```

Be a Good Citizen

```
producer.flush();  
producer.close();
```



Processing Messages with Java



Key, Takeaways, and Tips

Takeaways



The producer protocol implies the existence of a Partitioner that redirects messages to the correct partition



The partitions information is caught on the initial instantiation of the Producer Object



The response from the producer.send() method is a Future, but you can capture it in a Callback or Lambda

There are Retryable and non-retryable Exceptions and based on that the Protocol will automatically retry

Keys



Be sure how to configure your own Partitioner as homework



Try to ensure you can write a simple loop of sending messages by yourself

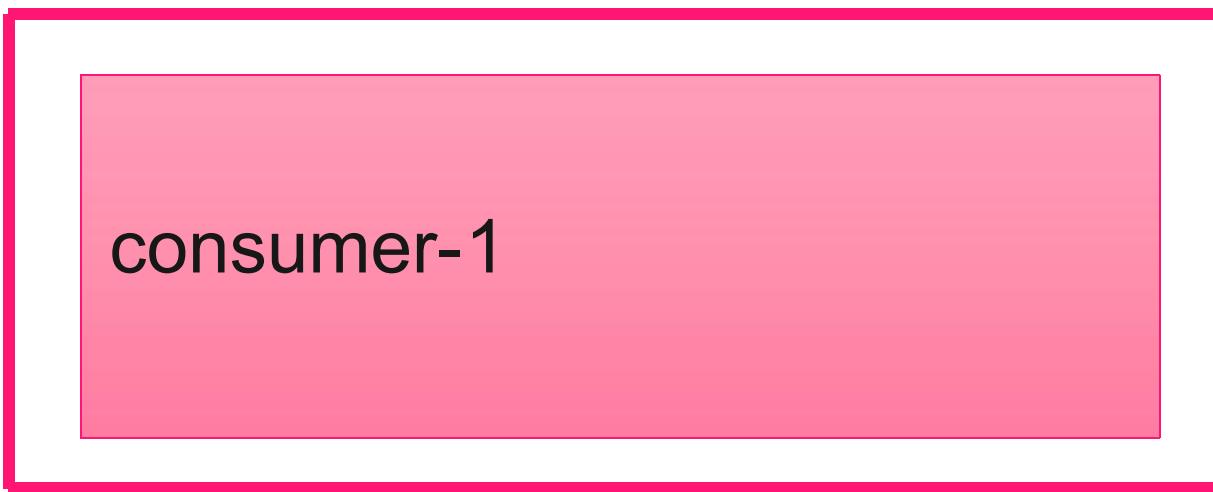
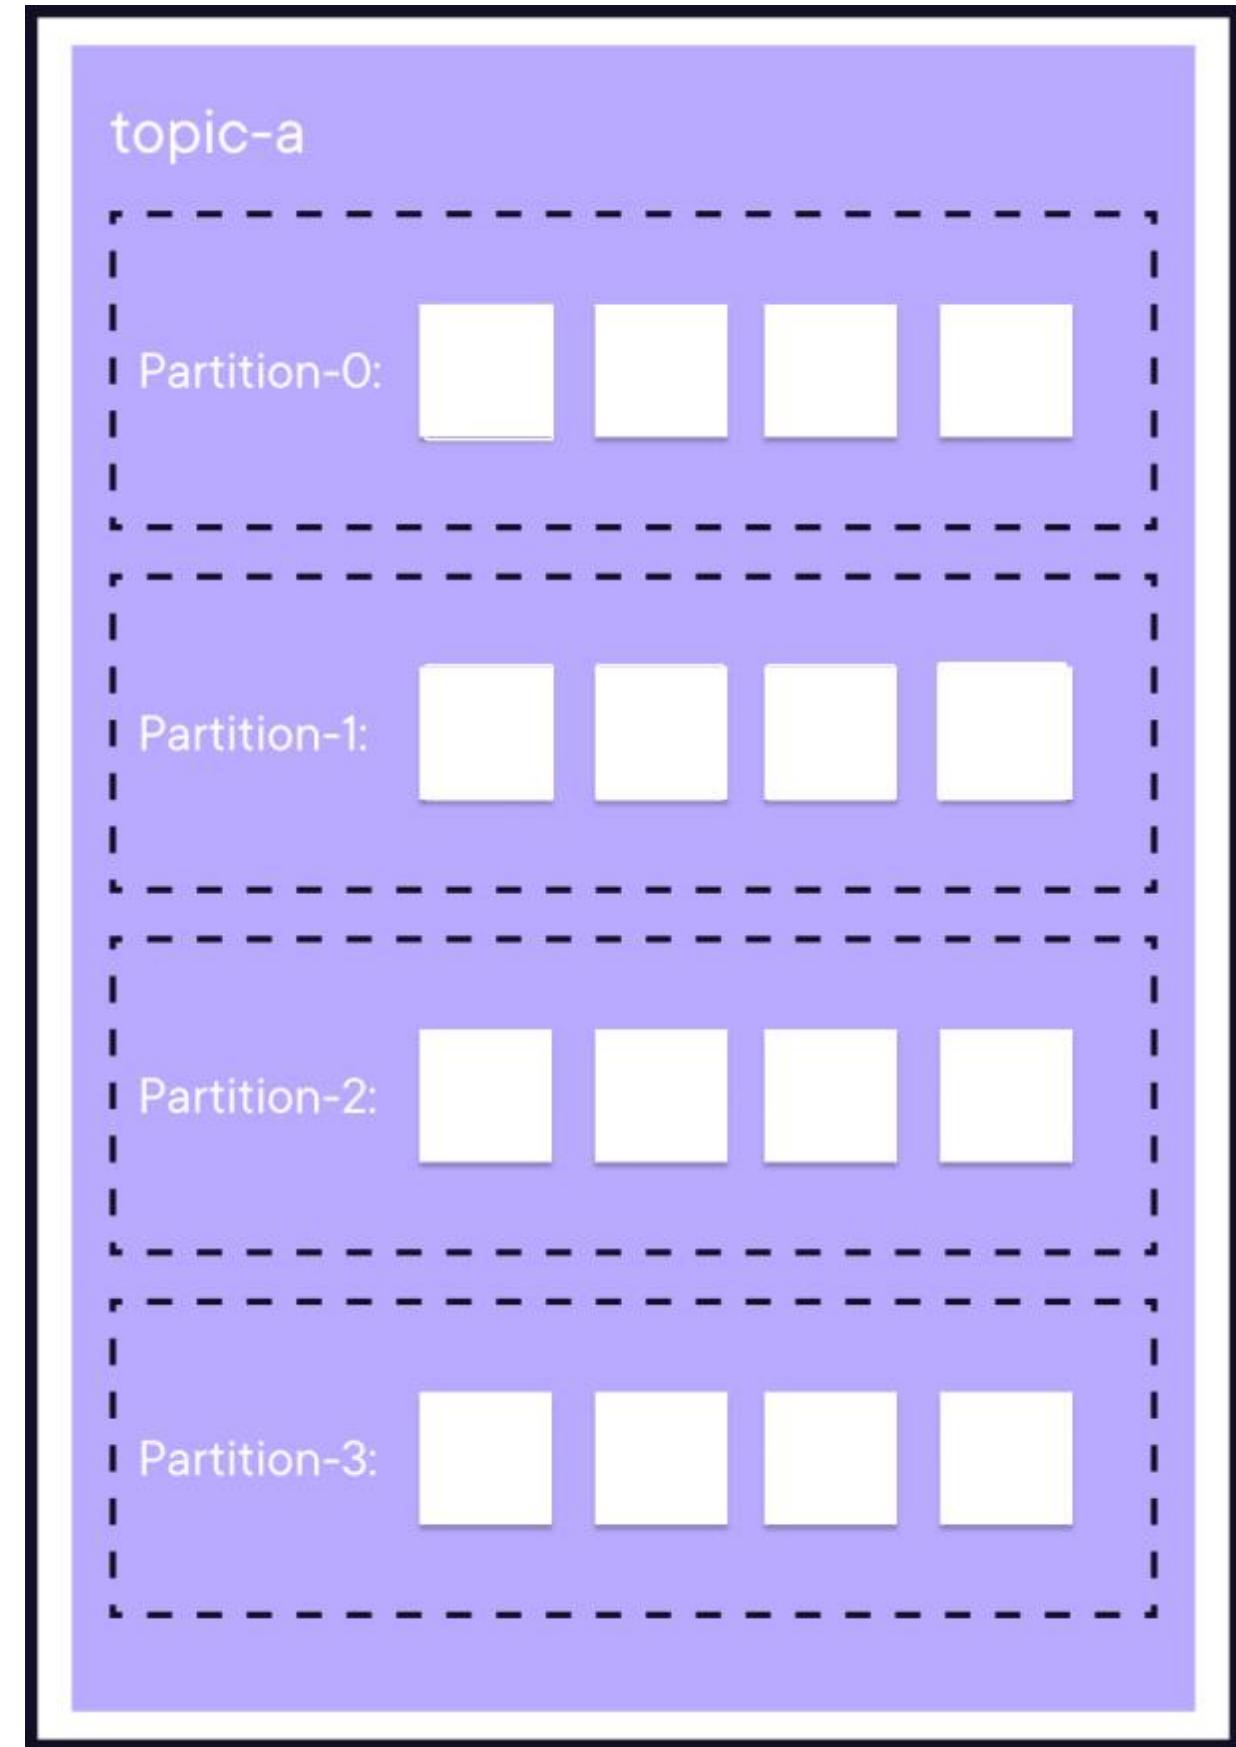


Try playing around what happens with the Producer if a broker is dead

Kafka Consumers

Kafka Consumer Group





Kafka Consumer Group



Consumers are typically done as a group



A single consumer will end up inefficient with large amounts of data

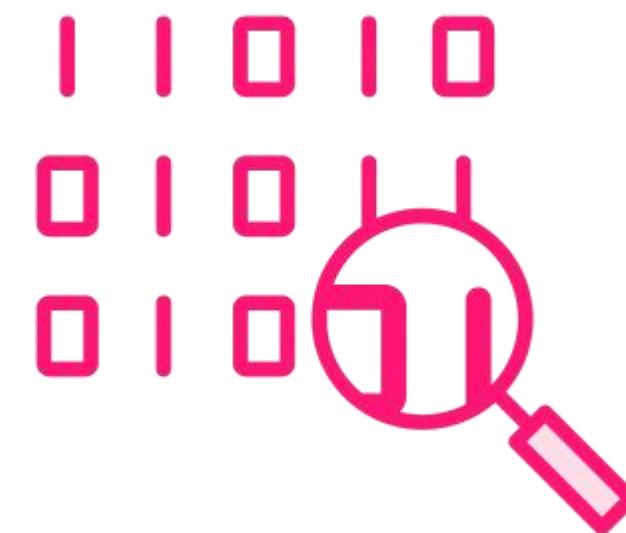
Kafka Consumer Group



Consumers are typically done as a group

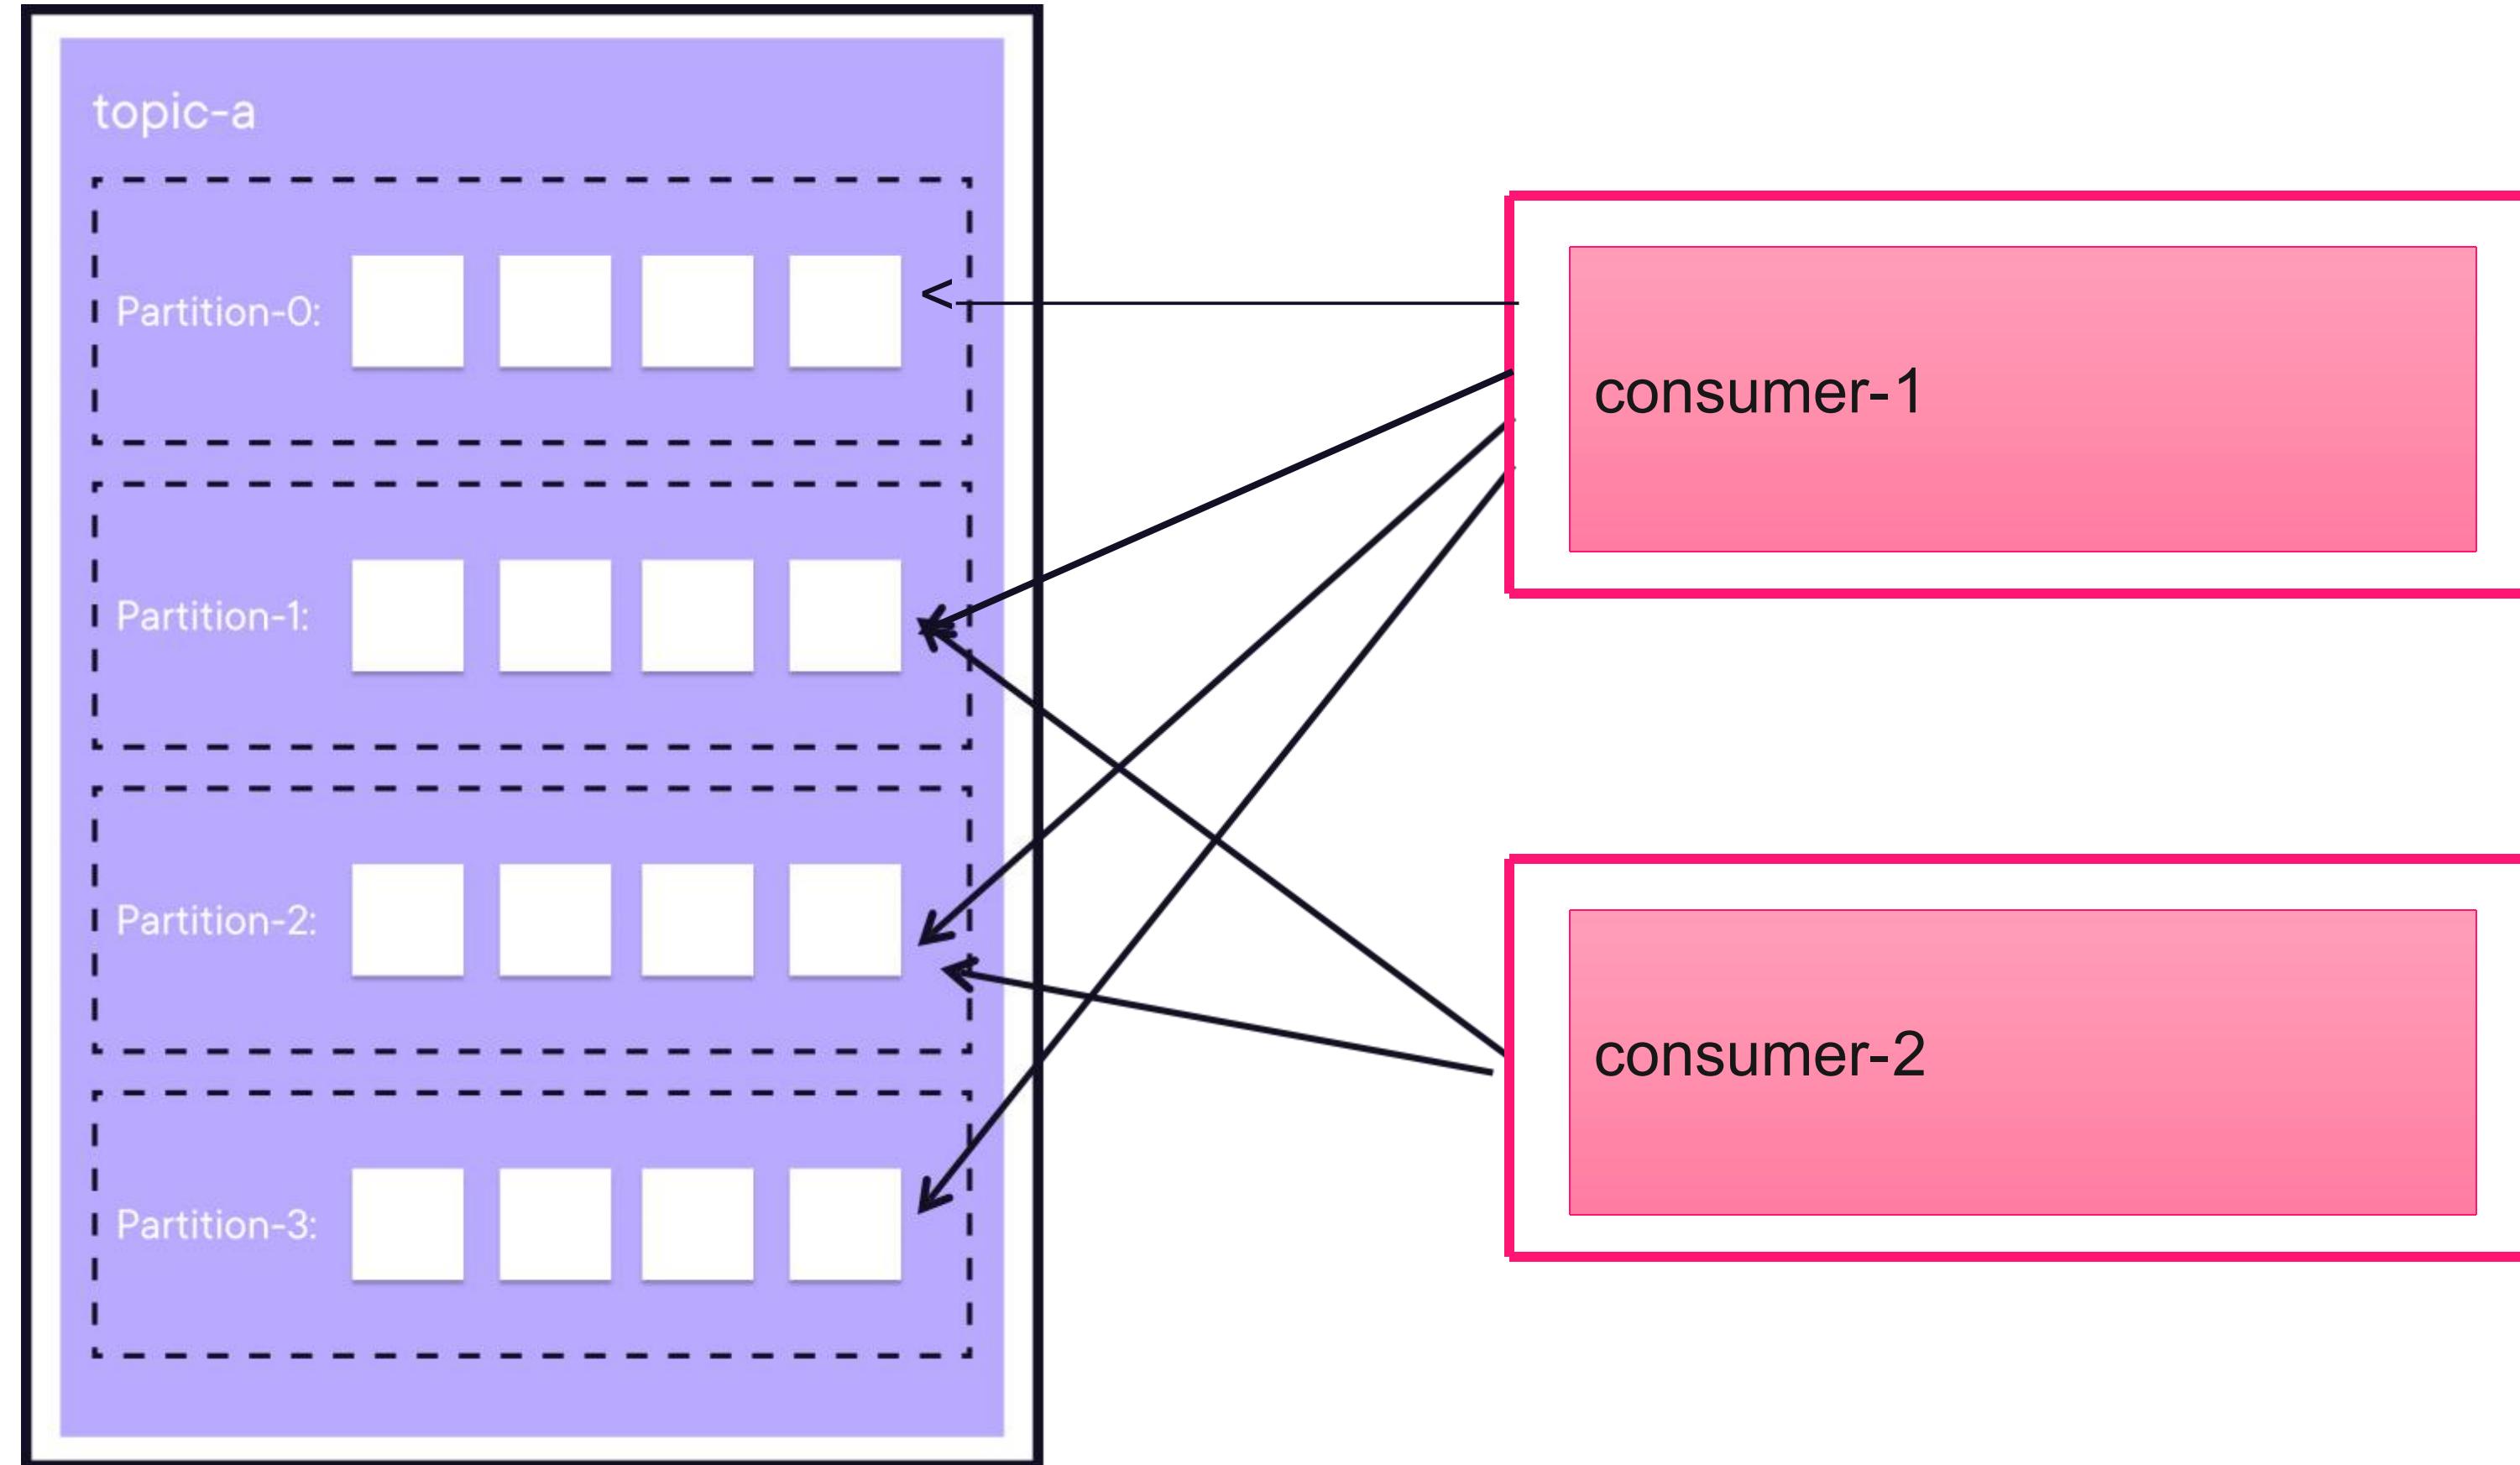


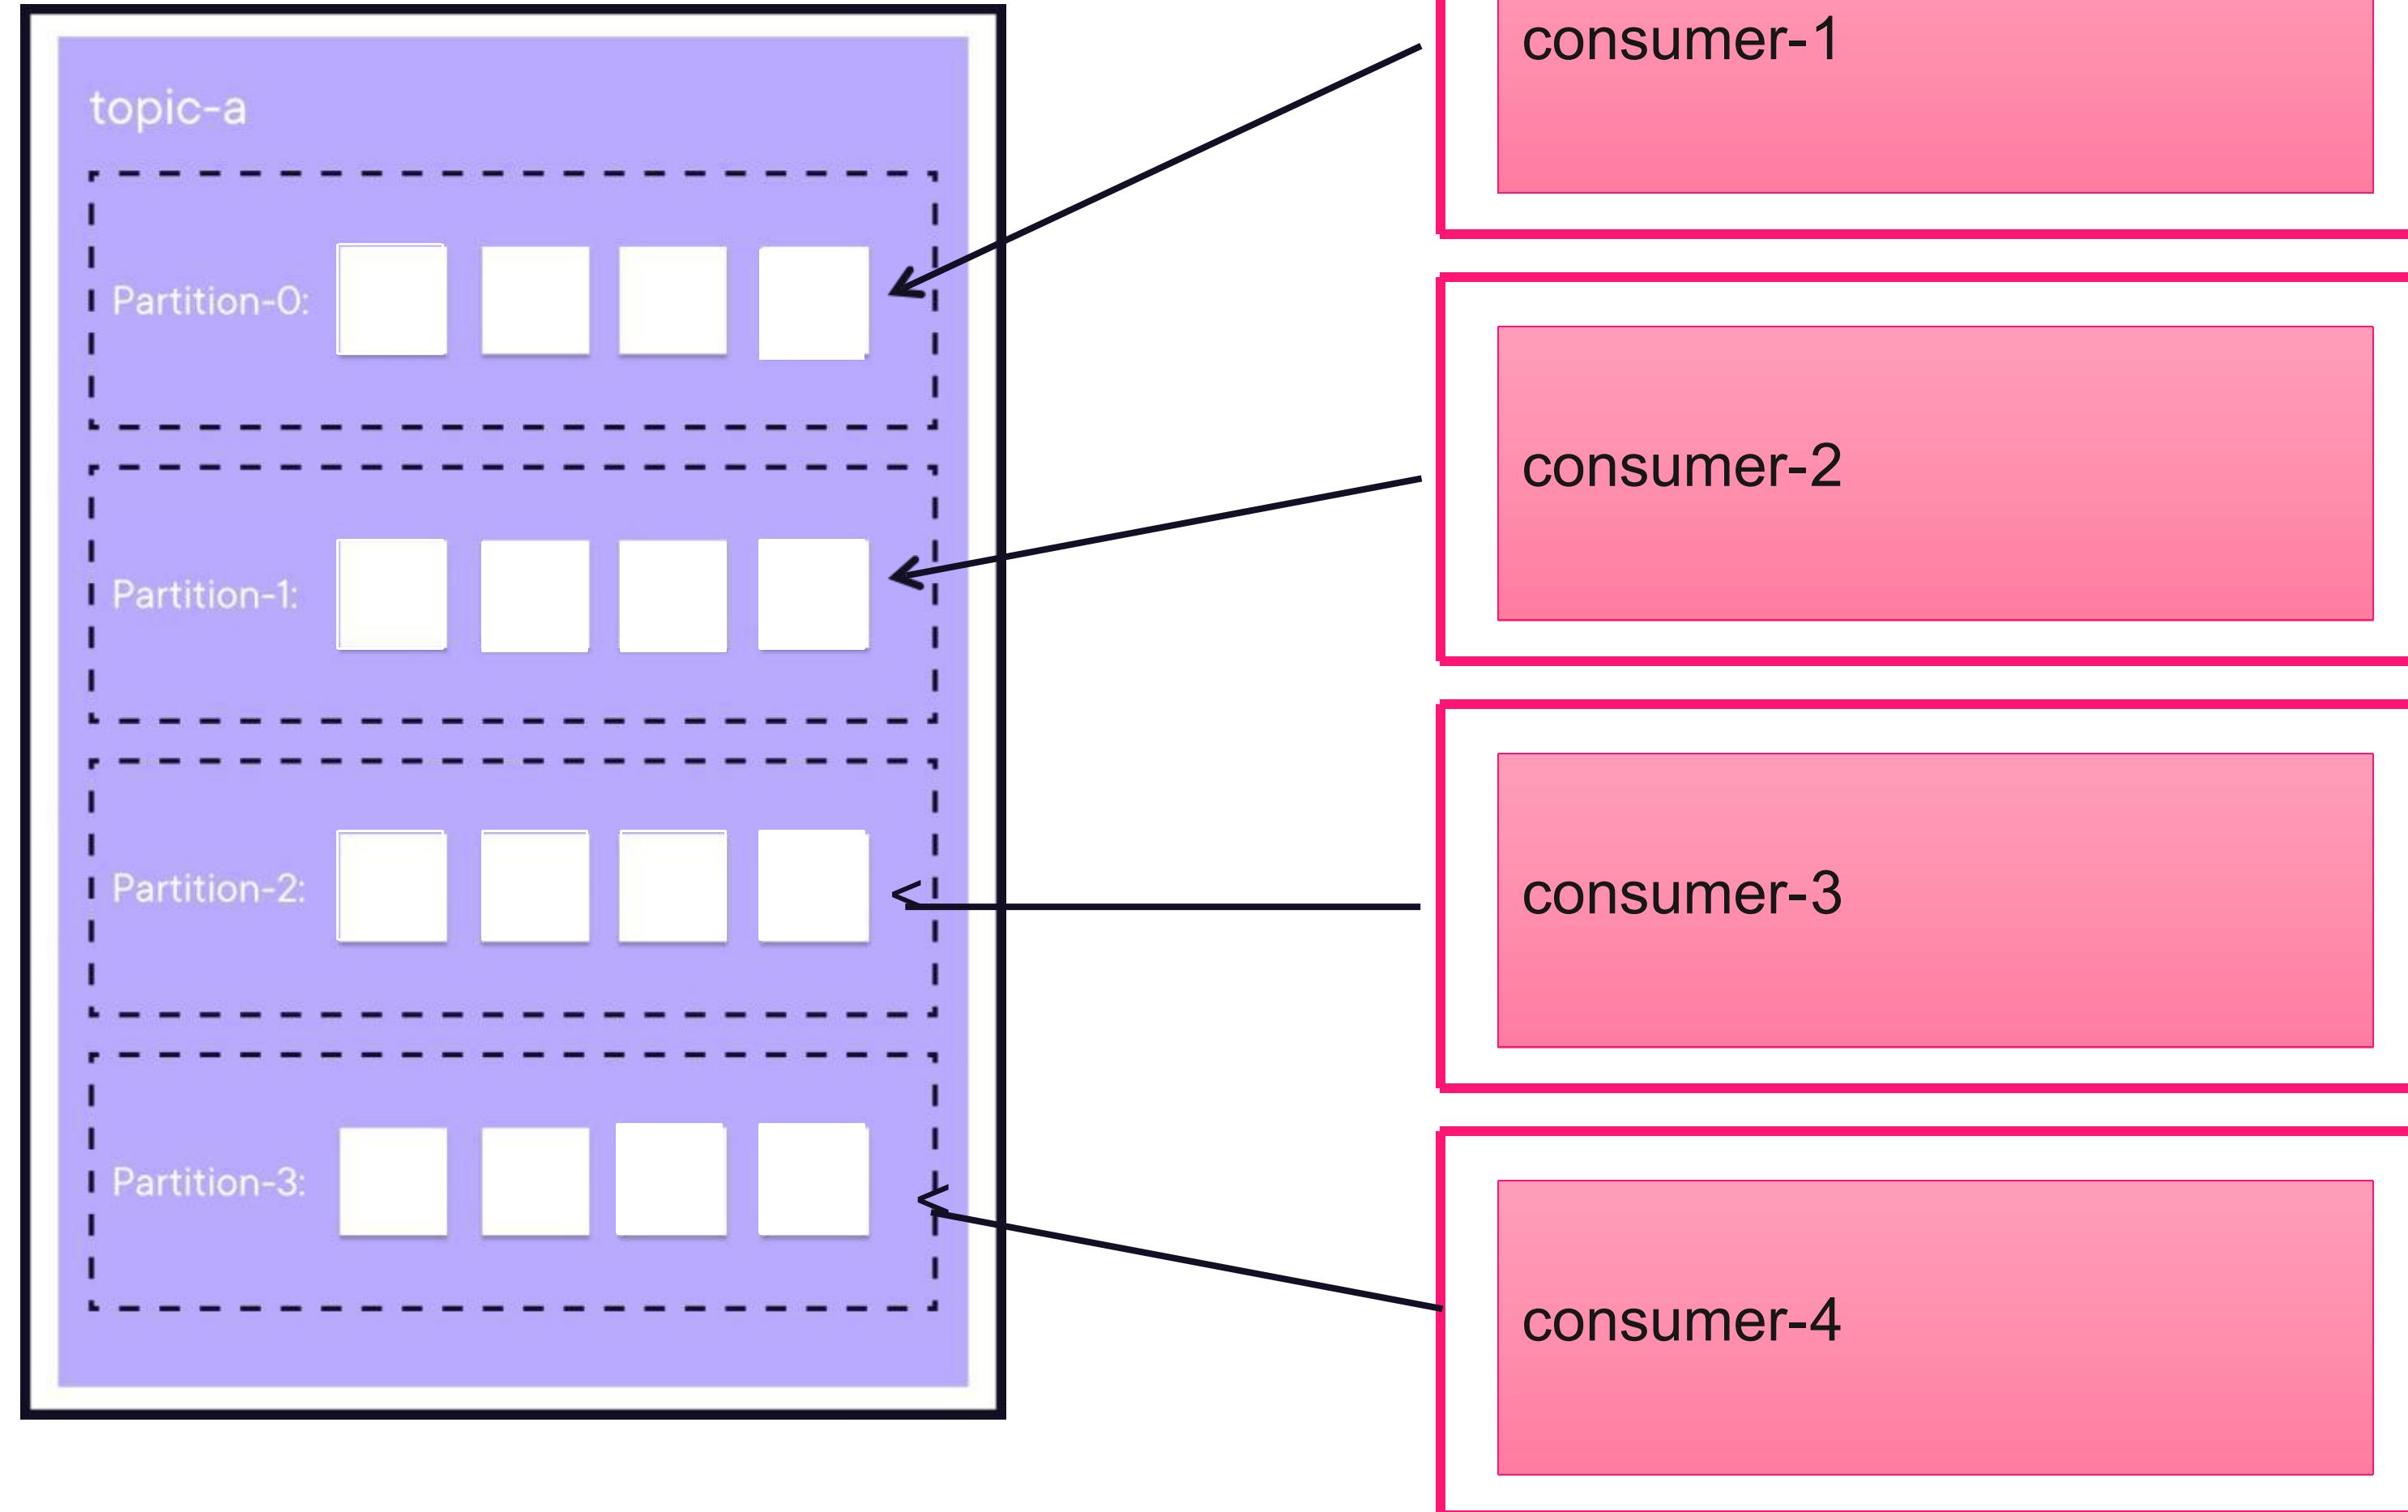
A single consumer will end up inefficient with large amounts of data

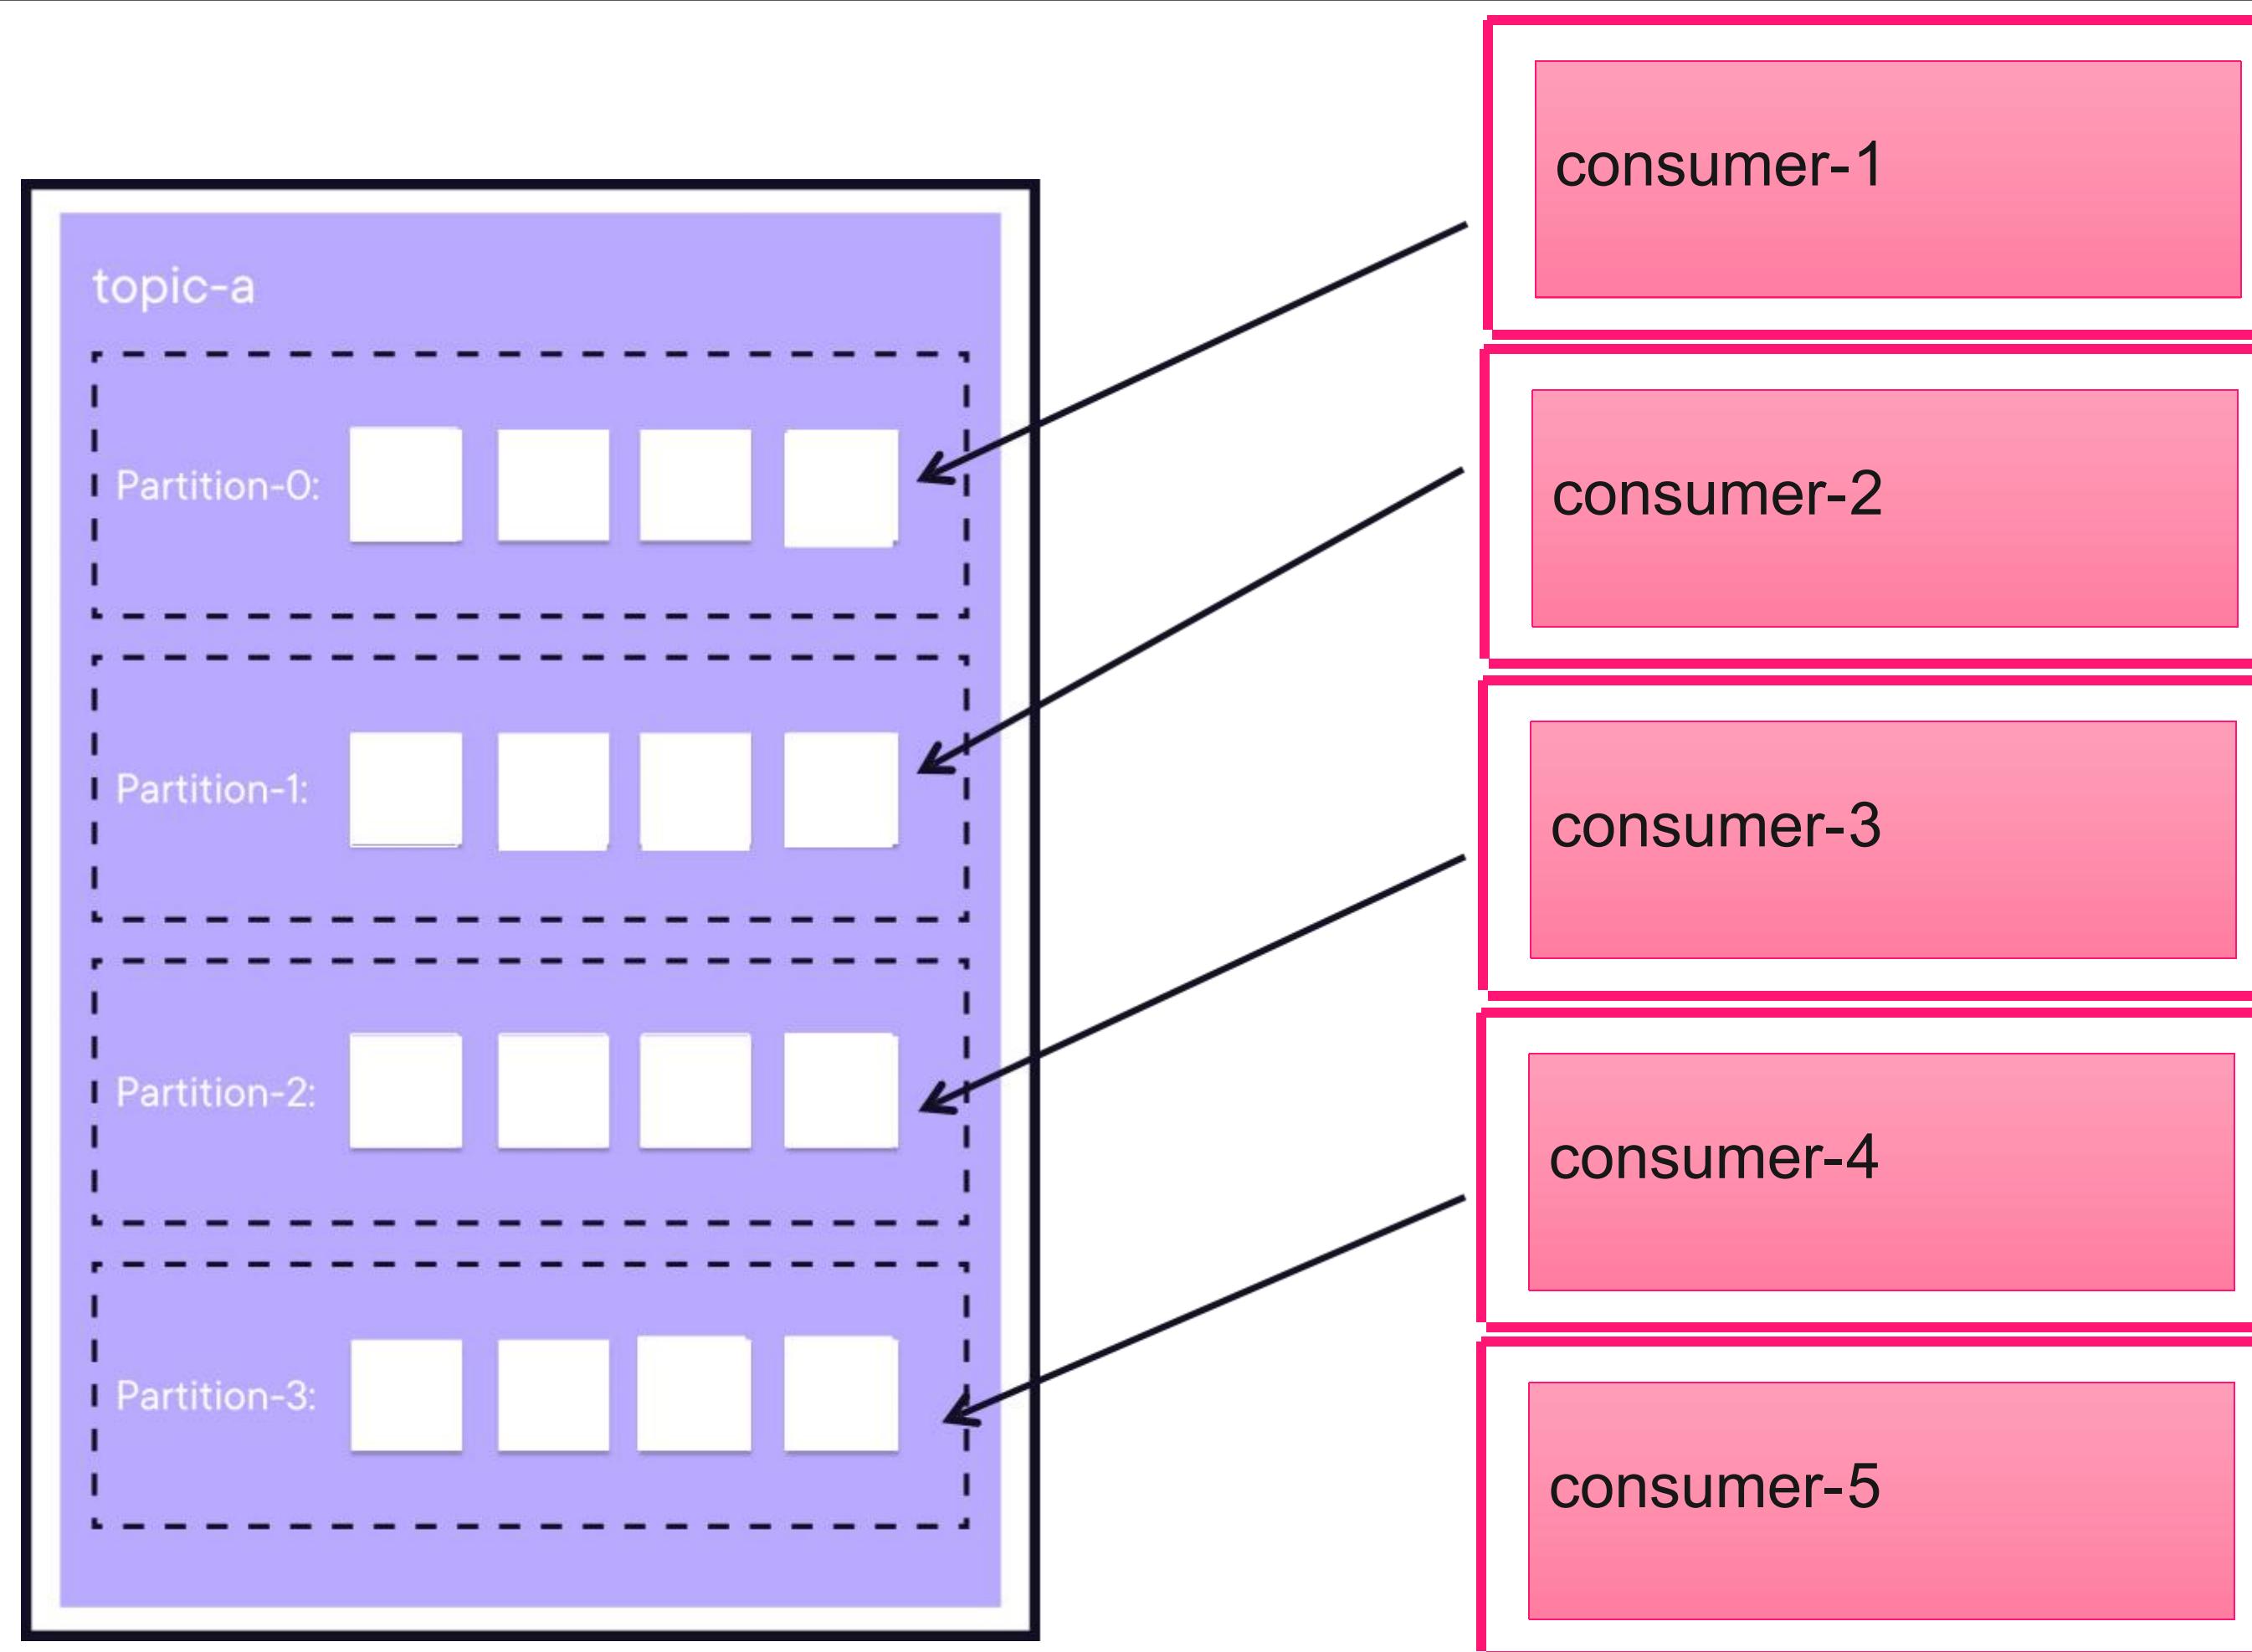


A consumer may never catch up

The consumer group id is key
so Kafka knows that messages
should be distributed to both
consumers without duplicating







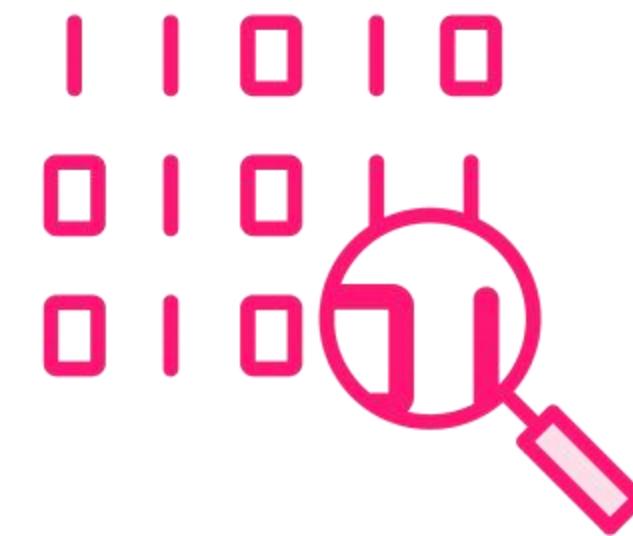
Kafka Consumer Group



Consumers are typically done as a group



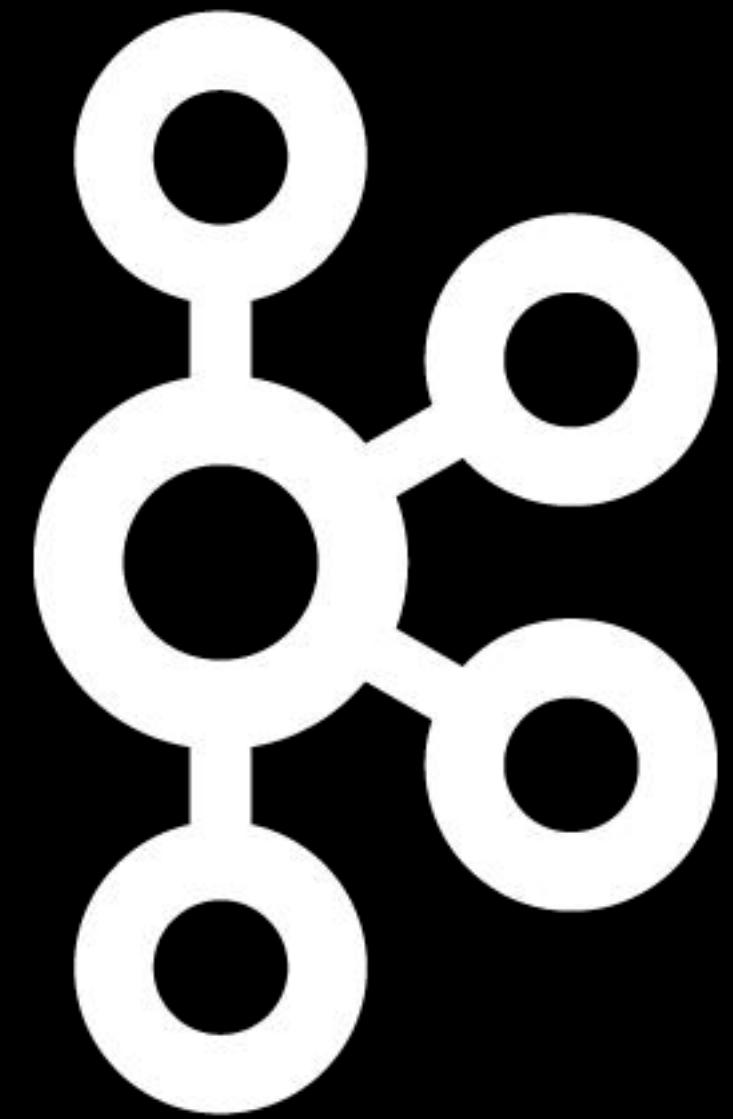
A single consumer will end up inefficient with large amounts of data



A consumer may never catch up



Every consumer should be on it's own machine, instance, pod



Consuming from Kafka



Deserializers and Consumer Configuration

Deserializers and Consumer Configuration

Construct a java.util.Properties object

```
Properties properties = new Properties();  
  
properties.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");  
  
properties.put(ConsumerConfig.GROUP_ID_CONFIG, "my_group");  
  
properties.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG, "org.apache.kafka.common.serialization.StringDeserializer");  
  
properties.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG, "org.apache.kafka.common.serialization.IntegerDeserializer");  
  
properties.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");  
  
KafkaConsumer<String, String> consumer = new KafkaConsumer<>(properties);
```

Deserializers and Consumer Configuration

Provide two or more locations where the Bootstrap servers are located

```
Properties properties = new Properties();

properties.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");

properties.put(ConsumerConfig.GROUP_ID_CONFIG, "my_group");

properties.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG, "org.apache.kafka.common.serialization.StringDeserializer");

properties.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG, "org.apache.kafka.common.serialization.IntegerDeserializer");

properties.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");

KafkaConsumer<String, String> consumer = new KafkaConsumer<>(properties);
```

Deserializers and Consumer Configuration

Provide a “Team Name”,
officially called a group.id

```
Properties properties = new Properties();

properties.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");

properties.put(ConsumerConfig.GROUP_ID_CONFIG, "my_group");

properties.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG, "org.apache.kafka.common.serialization.StringDeserializer");

properties.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG, "org.apache.kafka.common.serialization.IntegerDeserializer");

properties.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");

KafkaConsumer<String, String> consumer = new KafkaConsumer<>(properties);
```

Deserializers and Consumer Configuration

Provide a DeSerializer for
the key

```
Properties properties = new Properties();

properties.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");

properties.put(ConsumerConfig.GROUP_ID_CONFIG, "my_group");

properties.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG, "org.apache.kafka.common.serialization.StringDeserializer");

properties.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG, "org.apache.kafka.common.serialization.IntegerDeserializer");

properties.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");

KafkaConsumer<String, String> consumer = new KafkaConsumer<>(properties);
```

Deserializers and Consumer Configuration

Provide a DeSerializer for the value

```
Properties properties = new Properties();

properties.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");

properties.put(ConsumerConfig.GROUP_ID_CONFIG, "my_group");

properties.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG, "org.apache.kafka.common.serialization.StringDeserializer");

properties.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG, "org.apache.kafka.common.serialization.IntegerDeserializer");

properties.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");

KafkaConsumer<String, String> consumer = new KafkaConsumer<>(properties);
```

Consumer Object

**Construct a
org.apache.kafka.clie
nts.consumer.KafkaC
onsumer object**

```
KafkaConsumer consumer = new KafkaConsumer<>(properties);
```

Processing

**Use the consumer
that you have
constructed, and call
poll with pulse time.**

```
while (!done.get()) {  
    ConsumerRecords<String, String> records =  
    consumer.poll(Duration.of(500, ChronoUnit.MILLIS));  
    for (ConsumerRecord<String, String> record : records) {  
        System.out.format("offset: %d\n", record.offset());  
        System.out.format("partition: %d\n", record.partition());  
        System.out.format("timestamp: %d\n", record.timestamp());  
        System.out.format("timeStampType: %s\n",  
            record.timestampType());  
        System.out.format("topic: %s\n", record.topic());  
        System.out.format("key: %s\n", record.key());  
        System.out.format("value: %s\n", record.value());  
    }  
}
```

Be a Good Citizen

Processing

```
consumer.close();
```

Speaking of Rebalances



Regardless of what a consumer is doing, at regular intervals, which are configurable, they send a heartbeat to Kafka

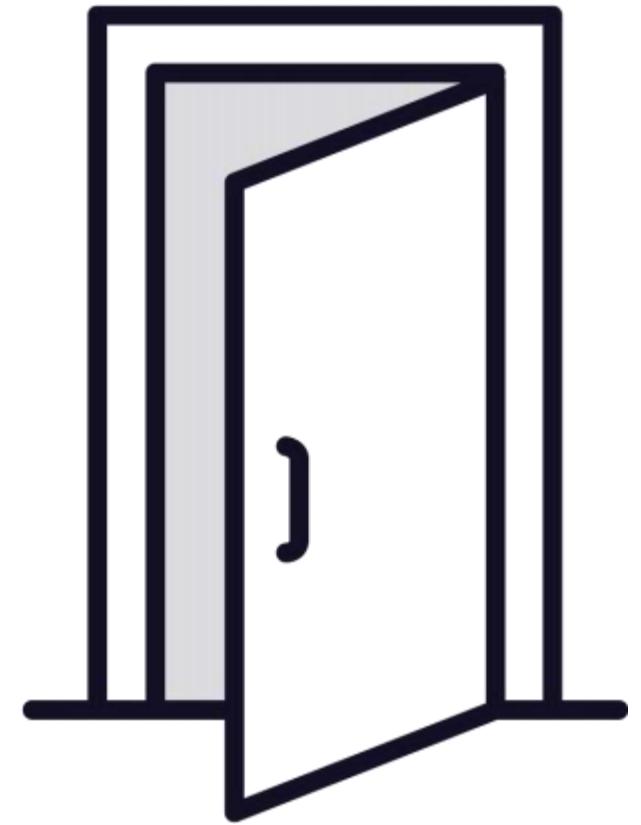


Specifically to the coordinator inside the broker, basically letting them know they are alive

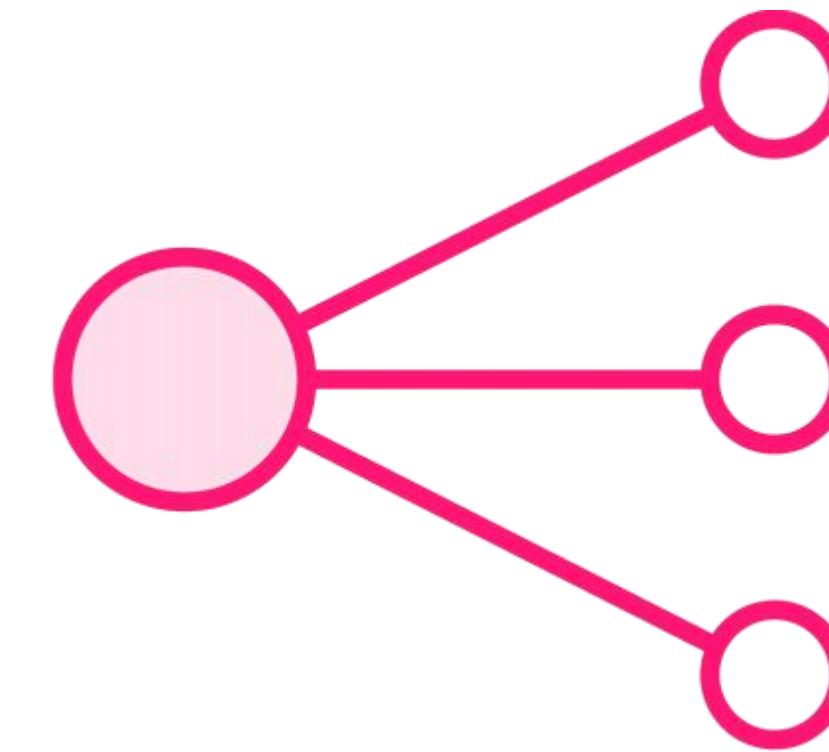


Also when we poll for records that lets Kafka know that consumer is alive

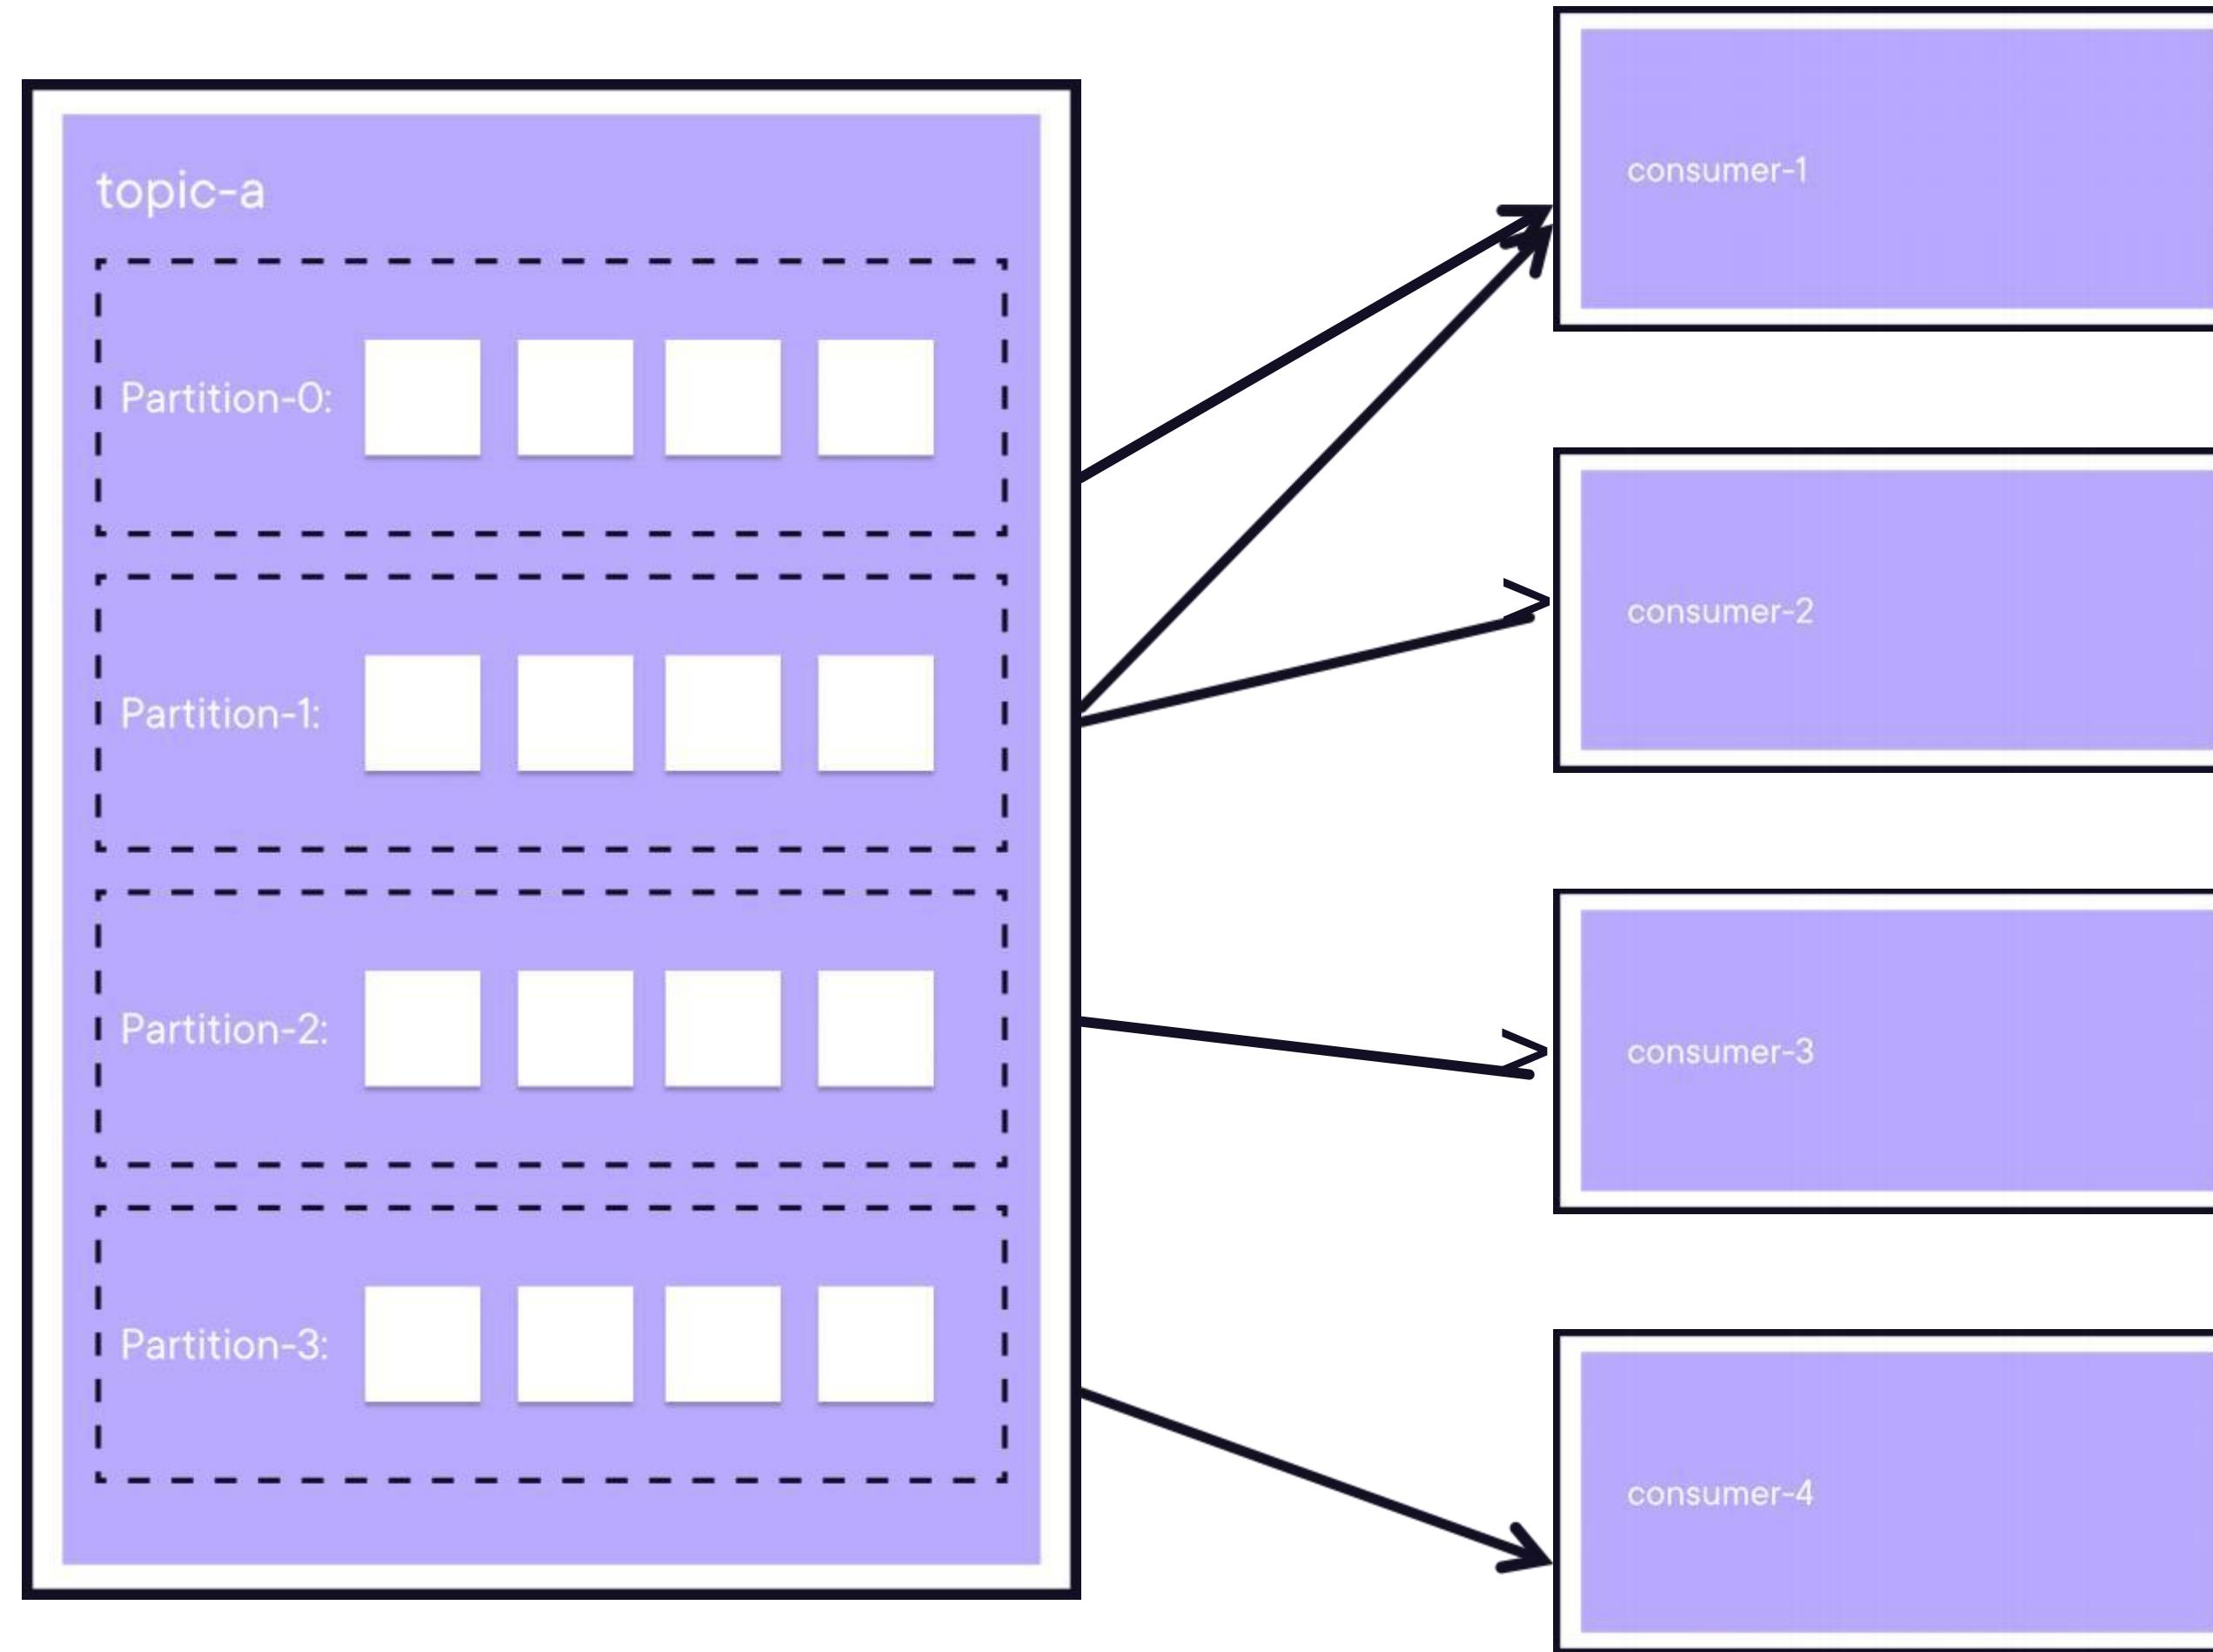
Speaking of Rebalances

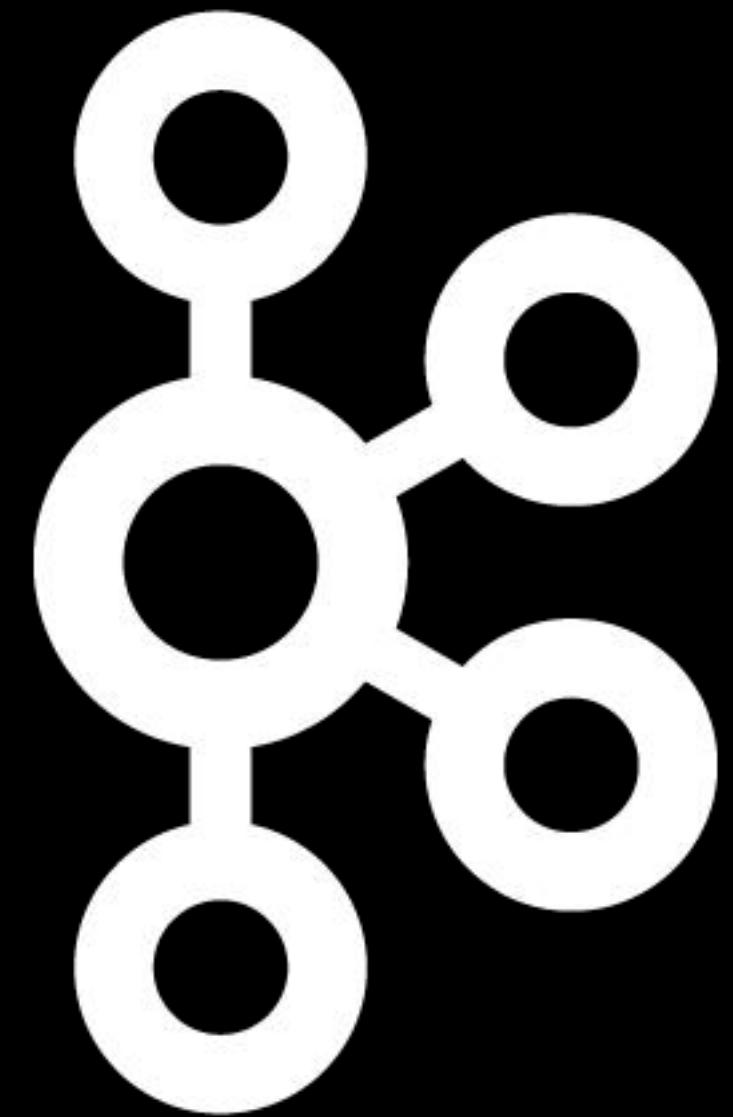


And as we said when you close a consumer that close method lets Kafka know they are not alive anymore



Which triggers a rebalance





Consuming Messages with Java



Key, Takeaways, and Tips

Takeaways



Consumers act as a group via the Consumer group ID



Each consumer in the group gets assigned a partition, and a partition is not shared by two members of a group



A rebalance is triggered when a member leaves the group or they haven't sent a heartbeat in a long time

A rebalance is a stop the world event that ensures all partitions are attended by some consumer in the group

Keys



Be sure to configure your Consumer to ensure you are not duplicating messages



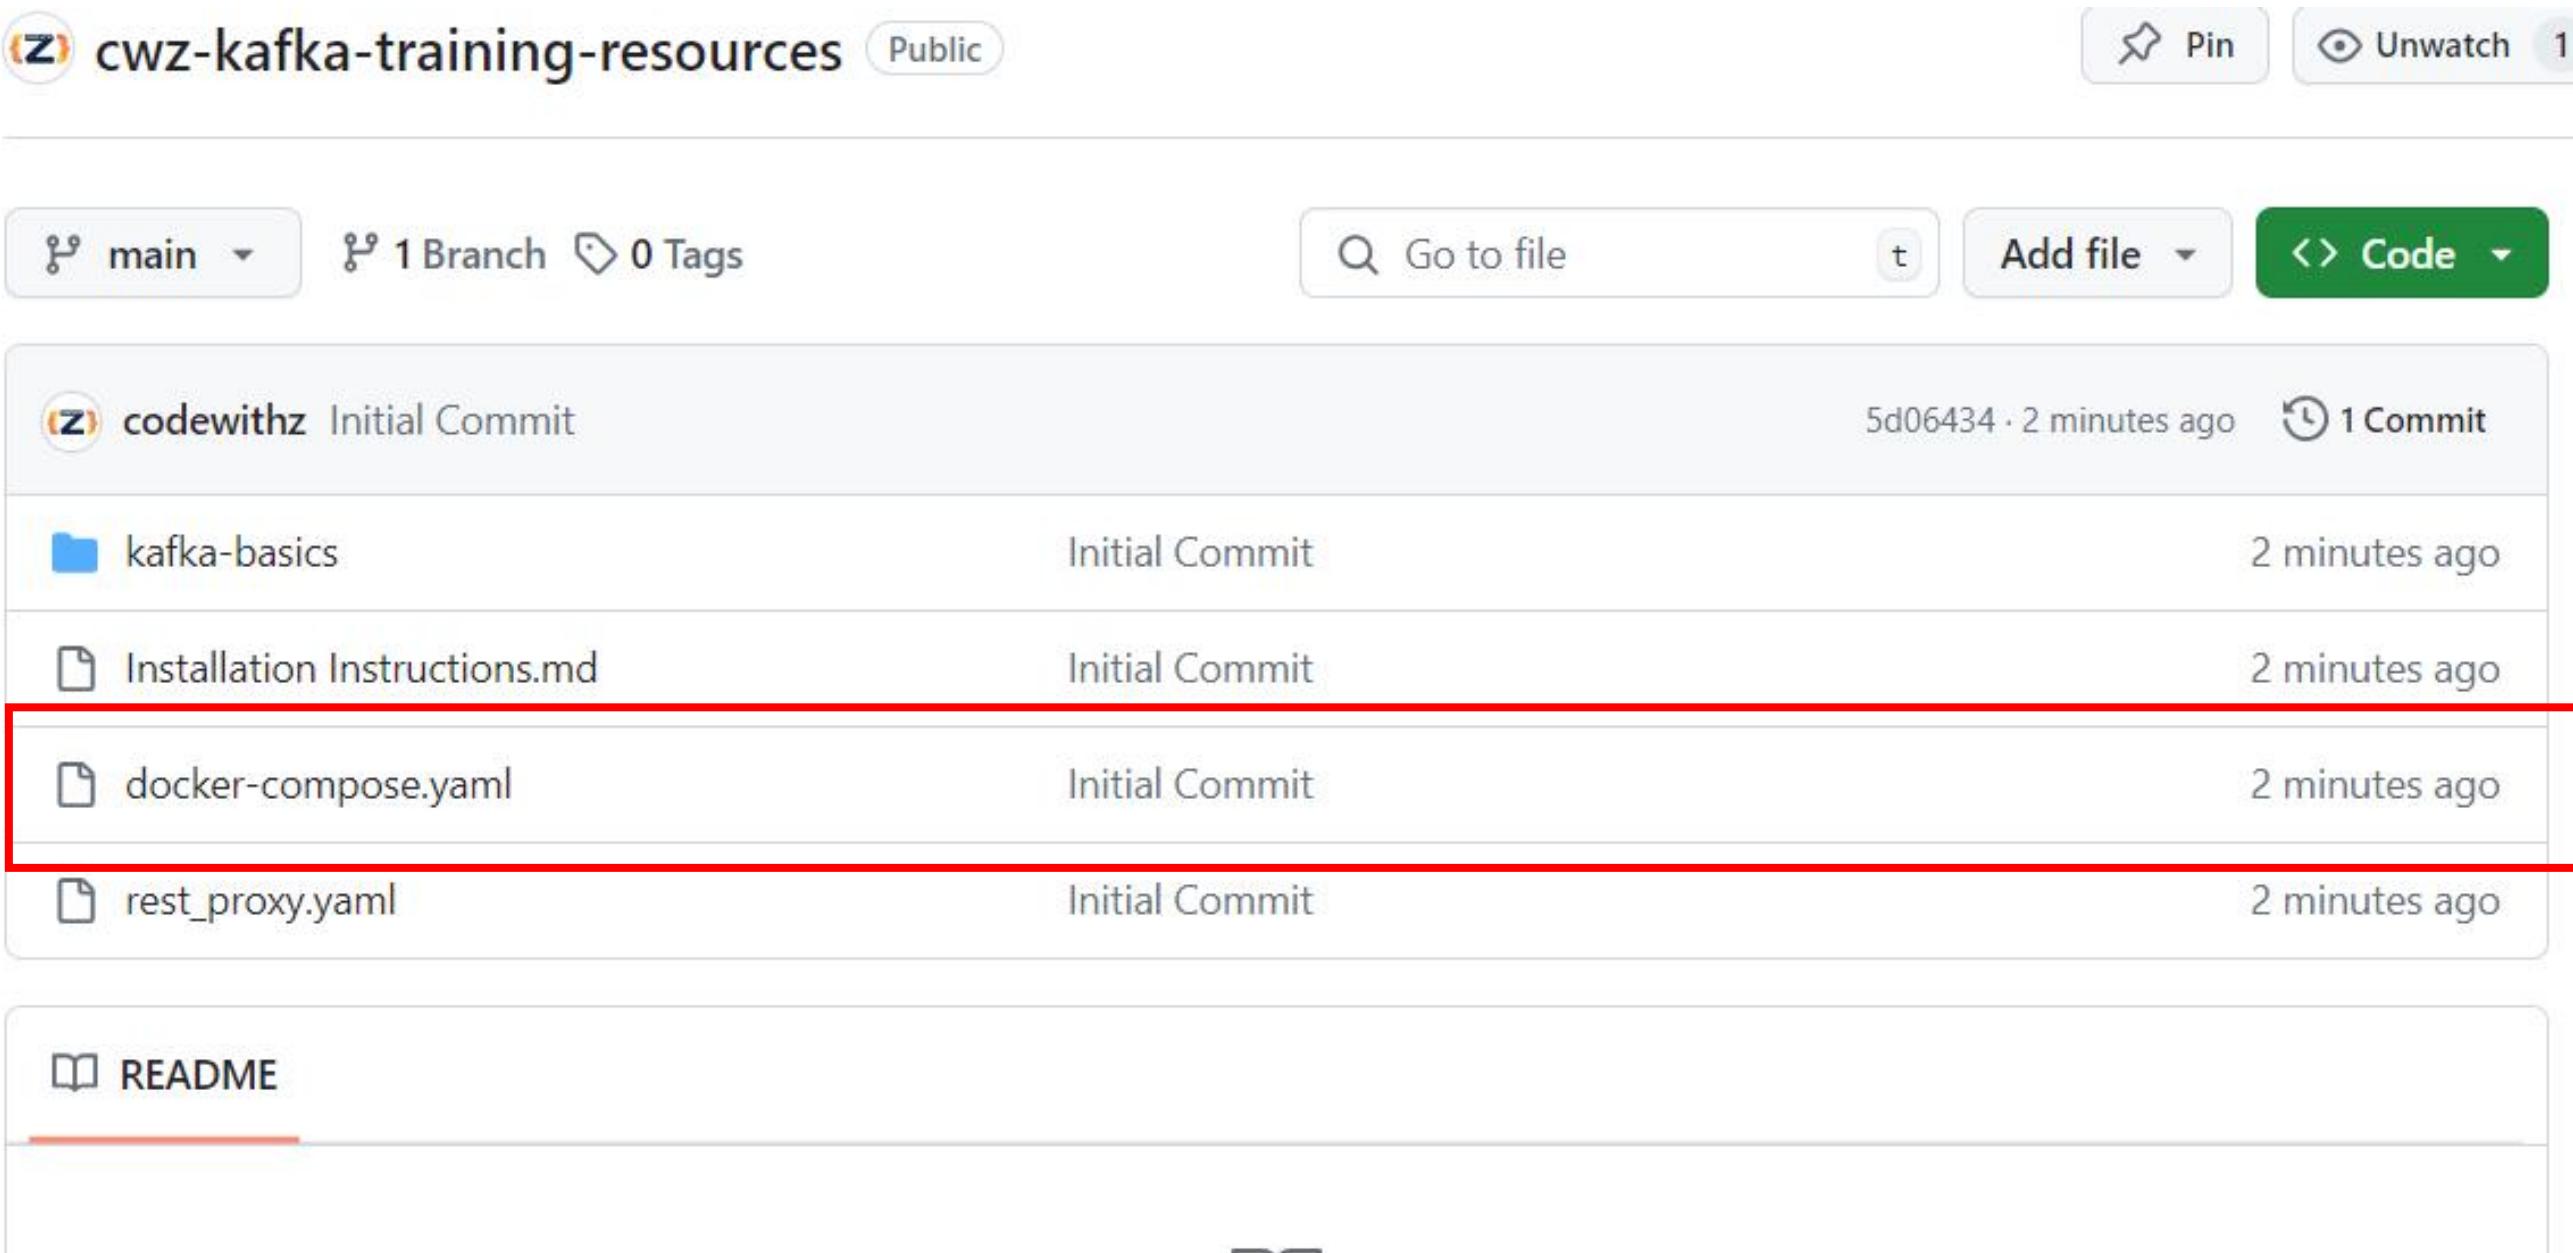
Try to create a consumer loop with two consumers and verify the partition assignment playing with keys in the producer



See what happens in the logs of the broker when a consumer joins or leaves the group

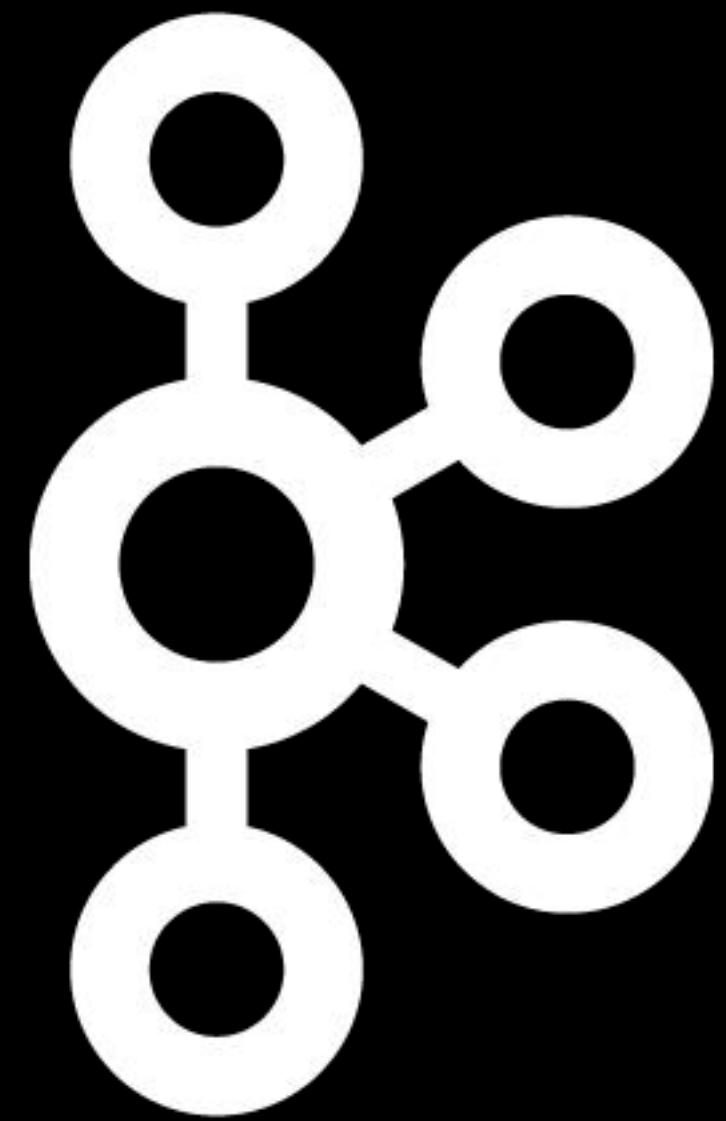
Launch Kafka Using the docker-compose file on following link

<https://github.com/codewithz/cwz-kafka-training-resources>



The screenshot shows a GitHub repository page for 'cwz-kafka-training-resources'. The repository is public. At the top, there are buttons for 'Pin' and 'Unwatch'. Below that, there are buttons for 'main' (with a dropdown arrow), '1 Branch', '0 Tags', 'Go to file' (with a search icon and a 't' button), 'Add file' (with a dropdown arrow), and 'Code' (with a dropdown arrow). The repository has one commit from 'codewithz' labeled 'Initial Commit' made 5d06434 · 2 minutes ago, which contains 1 commit. The commit details show a folder named 'kafka-basics' and two files: 'Installation Instructions.md' and 'docker-compose.yaml'. Both were committed 2 minutes ago. A red box highlights the 'docker-compose.yaml' file. Below these are two more files: 'rest_proxy.yaml' and 'README', both also committed 2 minutes ago.

File	Type	Commit
codewithz Initial Commit		5d06434 · 2 minutes ago
kafka-basics	Folder	Initial Commit 2 minutes ago
Installation Instructions.md	File	Initial Commit 2 minutes ago
docker-compose.yaml	File	Initial Commit 2 minutes ago
rest_proxy.yaml	File	Initial Commit 2 minutes ago
README	File	



Kafka CLI

If using Linux, add .sh after kafka-topics

```
#-----  
#----- To Create a Topic [Zookeper]  
#-----  
  
kafka-topics --zookeeper host:port --topic name_of_topic --create  
--partitions noOfPartitions --replication-factor noOfReplicas  
  
#-----  
#----- To Create a Topic [Bootstrap Server]  
#-----  
  
kafka-topics --bootstrap-server host:port --topic name_of_topic --create  
--partitions noOfPartitions --replication-factor noOfReplicas  
  
kafka-topics --bootstrap-server localhost:9092 --topic bot_first_topic --create --partitions 3  
--replication-factor 1
```

```
#-----  
#----- To List the Topics [Bootstrap Server]  
#-----
```

```
kafka-topics --bootstrap-server host:port --List
```

```
kafka-topics -bootstrap-server localhost:9092 --list
```

```
#-----  
#----- To Describe a Topic [Bootstrap Server]  
#-----
```

```
kafka-topics --bootstrap-server host:port --topic name_of_topic --describe
```

```
kafka-topics --bootstrap-server localhost:9092 --topic bot_first_topic --describe
```

```
#-----  
#----- Delete a Topic [Bootstrap Server]  
#-----
```

```
kafka-topics --bootstrap-server host:port --topic name_of_topic --delete
```

```
kafka-topics --bootstrap-server localhost:9092 --topic bot_first_topic --delete
```

```
#-----  
#----- To Produce Data to a Topic [Bootstrap Server]  
#-----
```

```
kafka-console-producer --bootstrap-server host:port  
--topic name_of_topic
```

```
kafka-console-producer --bootstrap-server localhost:9092  
--topic bot_first_topic
```

```
#-----  
#----- To Produce Data to a Topic with properties  
#-----
```

kafka-console-producer --bootstrap-server host:port
--topic name_of_topic --producer-property key=value

```
kafka-console-producer --bootstrap-server localhost:9092  
--topic bot_first_topic --producer-property acks=all
```

```
#-----  
----- To Produce Data to a Non-Existing Topic [Bootstrap Server]  
#-----
```

```
kafka-console-producer --bootstrap-server host:port  
--topic name_of_topic_which_doesnt_exists
```

```
kafka-console-producer --bootstrap-server localhost:9092  
--topic bot_fourth_topic
```

```
##### ----- ALWAYS CREATE A TOPIC BEFORE PRODUCING DATA TO THEM -----
```

```
#-----  
#----- To Consume Data from a Topic  
#-----
```

```
kafka-console-consumer --bootstrap-server host:port  
--topic name_of_topic
```

```
kafka-console-consumer --bootstrap-server localhost:9092  
--topic bot_first_topic
```

```
#-----  
#----- To Consume Data from a Topic from beginning  
#-----
```

kafka-console-consumer --bootstrap-server host:port
--topic name_of_topic --from-beginning

kafka-console-consumer --bootstrap-server localhost:9092
--topic bot_first_topic --from-beginning

```
kafka-log-dirs.sh --bootstrap-server <kafka-broker> --describe --broker-list <broker-id>
```

```
#-----  
#----- To Consume Data from a Topic  
#-----
```

```
kafka-console-consumer --bootstrap-server host:port  
--topic name_of_topic --group groupName
```

```
kafka-console-consumer --bootstrap-server localhost:9092  
--topic bot_second_topic --group my-first-app
```

```
kafka-console-consumer --bootstrap-server localhost:9092  
--topic bot_second_topic --group my-first-app
```

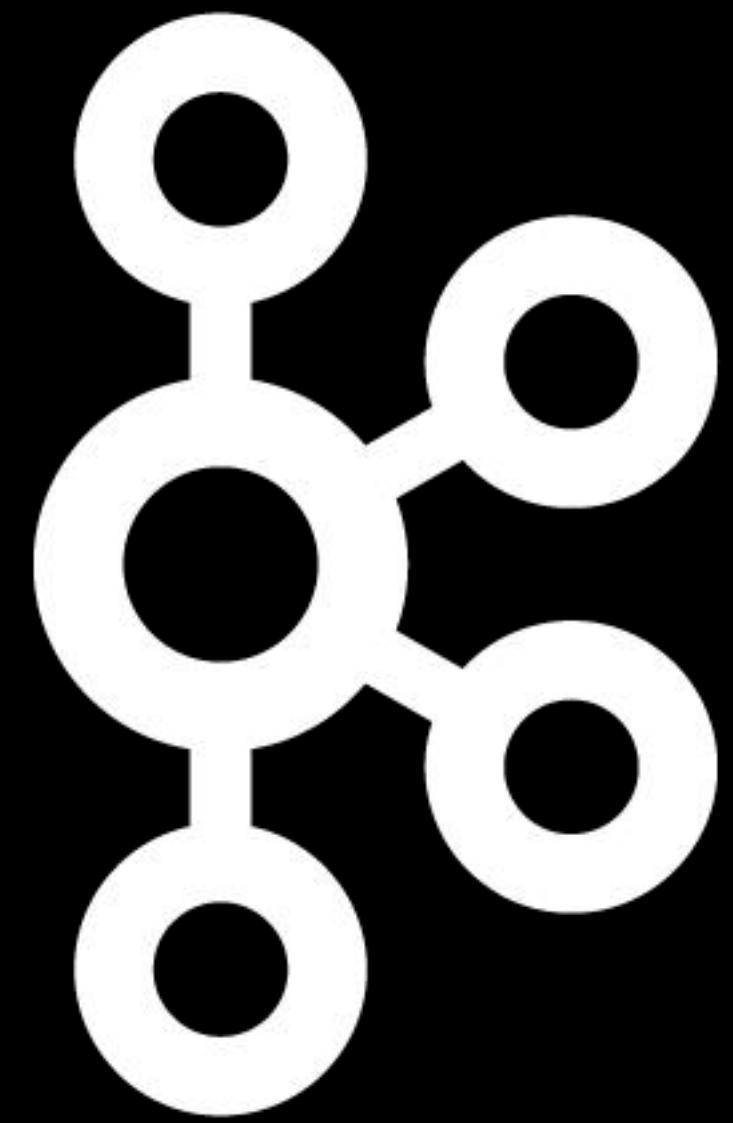
```
kafka-consumer-groups --bootstrap-server localhost:9092 --list
```

```
kafka-consumer-groups --bootstrap-server localhost:9092  
--describe --group my-first-app
```

GROUP	TOPIC	PARTITION	CURRENT-OFFSET	LOG-END-
my-first-app	bot_second_topic	0	8	8
my-first-app	bot_second_topic	1	8	8
my-first-app	bot_second_topic	2	8	8

```
#-----  
#----- Resetting Offset|  
#-----
```

```
kafka-consumer-groups --bootstrap-server localhost:9092  
--group my-first-app --reset-offsets  
--to-earliest --execute --topic bot_first_topic
```



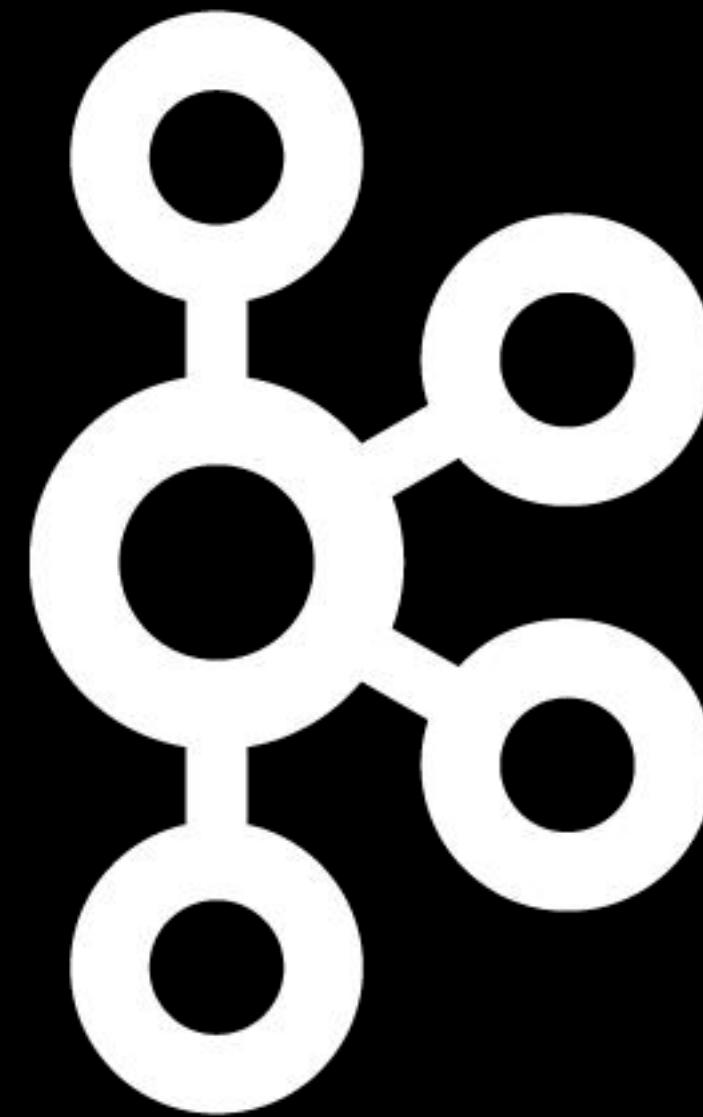
Kafka Rest Proxy

Import `rest_proxy.yaml` from the git repo and import to POSTMAN

<https://github.com/codewithz/cwz-kafka-training-resources>

The screenshot shows a GitHub repository page for "cwz-kafka-training-resources". The repository is public and has one branch ("main") and no tags. The commit history shows an initial commit by "codewithz" 2 minutes ago, which includes several files: "kafka-basics", "Installation Instructions.md", "docker-compose.yaml", and "rest_proxy.yaml". The "rest_proxy.yaml" file is highlighted with a red box. Other files listed are "README" and ".gitignore".

File	Commit	Time
codewithz Initial Commit	Initial Commit	5d06434 · 2 minutes ago
kafka-basics	Initial Commit	2 minutes ago
Installation Instructions.md	Initial Commit	2 minutes ago
docker-compose.yaml	Initial Commit	2 minutes ago
rest_proxy.yaml	Initial Commit	2 minutes ago



Kafka Java Producer and Consumer

Download the Skeleton Project from following link.

<https://github.com/codewithz/cwz-kafka-training-resources>

The screenshot shows a GitHub repository page for 'cwz-kafka-training-resources'. The repository is public. At the top, there are buttons for 'Pin' and 'Unwatch'. Below that, there are buttons for 'main' (with a dropdown arrow), '1 Branch', '0 Tags', 'Go to file' (with a dropdown arrow), 'Add file' (with a dropdown arrow), and 'Code' (with a dropdown arrow). The main content area shows a commit history:

File / Commit Type	Description	Time Ago
Initial Commit	codewithz Initial Commit	5d06434 · 2 minutes ago
kafka-basics	Initial Commit	2 minutes ago
Installation Instructions.md	Initial Commit	2 minutes ago
docker-compose.yaml	Initial Commit	2 minutes ago
rest_proxy.yaml	Initial Commit	2 minutes ago
README		

Create a Java File by name ProducerDemo.java

```
package com.codewithz;

import org.apache.kafka.clients.producer.KafkaProducer;
import org.apache.kafka.clients.producer.ProducerConfig;
import org.apache.kafka.clients.producer.ProducerRecord;
import org.apache.kafka.common.serialization.StringSerializer;

import java.util.Properties;

public class ProducerDemo {

    Run | Debug
    public static void main(String[] args) {
        String bootstrapServer="127.0.0.1:9092";
        String topic="codewithz_first_topic";
        //Create Producer Properties
    }
}
```

```
//Create Producer Properties
```

```
Properties properties=new Properties();
properties.setProperty(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG,bootstrapServer);
properties.setProperty(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG, StringSerializer.class.getName());
properties.setProperty(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG, StringSerializer.class.getName());
```

```
// Create the Producer
```

```
KafkaProducer<String, String> producer=new KafkaProducer<String, String>(properties);
```

```
//Create a ProducerRecord
```

```
ProducerRecord<String, String> record=new ProducerRecord<>(topic,value:"Hello from Java Code");
```

```
// Send the data -- Synchronous Way  
producer.send(record);  
  
//FLush the Data  
producer.flush();  
}  
}
```

Create a Java File by name ProducerDemoWithCallback.java

```
package com.codewithz;

import org.apache.kafka.clients.producer.*;
import org.apache.kafka.common.serialization.StringSerializer;
import org.slf4j.Logger;
import org.slf4j.LoggerFactory;

import java.util.Properties;

public class ProducerDemoWithCallback {

    Run | Debug
    public static void main(String[] args) {
        String bootstrapServer="127.0.0.1:9092";

        Logger logger= LoggerFactory.getLogger(ProducerDemoWithCallback.class);
        String topic ="codewithz_third_topic";
    }
}
```

//Create Producer Properties

```
Properties properties=new Properties();
properties.setProperty(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG,
bootstrapServer);
properties.setProperty(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
StringSerializer.class.getName());
properties.setProperty(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
StringSerializer.class.getName());
```

```
// Create the Producer
```

```
KafkaProducer<String, String> producer=new KafkaProducer<String, String>(properties);
```

```
for (int i=1;i<=10;i++){
    //Create a Producer Record
    ProducerRecord<String, String> record=
        new ProducerRecord<>(topic,"Hello from Java Code: "+Integer.toString(i));
```

Sending Data in Asynchronous Way

```
producer.send(record, new callback() {
    @Override
    public void onCompletion(RecordMetadata recordMetadata, Exception e) {
        // Execute everytime a record is successfully sent,
        //if record sending fails, it will throw an exception
        if(e==null){
            //Record is sucessfully sent
            logger.info("Recieved new MetaData");
            logger.info("Topic:"+recordMetadata.topic()+"\n"+
                        "Partition:"+recordMetadata.partition()+"\n"+
                        "Offset:"+recordMetadata.offset()+"\n"+
                        "Timestamp:"+recordMetadata.timestamp());
        }
        else{
            //Exception have occurred
            logger.error("Error while producing :",e);
        }
    }
});
```

```
    );
}

//Flush the Data
producer.flush();

//Flush and Close Producer
producer.close();

}
```

Create a Java File by name ProducerDemoWithKeys.java

```
package com.codewithz;

import org.apache.kafka.clients.producer.*;
import org.apache.kafka.common.serialization.StringSerializer;
import org.slf4j.Logger;
import org.slf4j.LoggerFactory;
```

```
import java.util.Properties;
```

```
public class ProducerDemoWithKeys {
```

Run | Debug

```
    public static void main(String[] args) {
```

```
        String bootstrapServer="127.0.0.1:9092";
```

```
        Logger logger= LoggerFactory.getLogger(ProducerDemoWithKeys.class);
```

```
        String topic ="codewithz_third_topic";
```

//Create Producer Properties

```
Properties properties=new Properties();
properties.setProperty(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG,
bootstrapServer);
properties.setProperty(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
StringSerializer.class.getName());
properties.setProperty(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
StringSerializer.class.getName());
```

```
KafkaProducer<String, String> producer=new KafkaProducer<String, String>(properties);
```

```
for (int i=1;i<=10;i++){
    //key and Value
    String key="id_"+Integer.toString(i);
    String value="Hello from Java Code: "+Integer.toString(i);
    //Create a Producer Record
    ProducerRecord<String, String> record=new ProducerRecord<>(topic,key,value);

    logger.info("Key:"+key);
    //id_1 -- P0
    //id_2 -- P2
    //id_3 -- P0
    //id_4 -- P2
    //id_5 -- P2
    //id_6 -- P0
    //id_7 -- P2
    //id_8 -- P1
```

```
// Send the data -- Asynchronous way
```

```
producer.send(record, new Callback() {
    @Override
    public void onCompletion(RecordMetadata recordMetadata, Exception e) {
        // Execute everytime a record is successfully sent,
        // if record sending fails, it will throw an exception
    }
})
```

```
if(e==null){  
    //Record is successfully sent  
    logger.info("Recieved new MetaData");  
    logger.info("-----"+'\n'+  
    "Topic:"+recordMetadata.topic()+"\n"+  
    "Partition:"+recordMetadata.partition()+"\n"+  
    "Offset:"+recordMetadata.offset()+"\n"+  
    "Timestamp:"+recordMetadata.timestamp()+"\n"+  
    "Key:"+key+'\n'+  
    "-----"  
);  
}
```

```
        }  
    }  
});  
}  
};
```

```
//Flush the Data  
producer.flush();  
  
//Flush and Close Producer  
producer.close();
```

Create a Java File by name ConsumerDemo.java

```
package com.codewithz;

import org.apache.kafka.clients.consumer.ConsumerConfig;
import org.apache.kafka.clients.consumer.ConsumerRecord;
import org.apache.kafka.clients.consumer.ConsumerRecords;
import org.apache.kafka.clients.consumer.KafkaConsumer;
import org.apache.kafka.common.serialization.StringDeserializer;
import org.slf4j.Logger;
import org.slf4j.LoggerFactory;

import java.time.Duration;
import java.util.Arrays;
import java.util.Collections;
import java.util.Properties;

public class ConsumerDemo {

    Run | Debug
    public static void main(String[] args) {

        Logger logger= LoggerFactory.getLogger(ConsumerDemo.class);
        String bootstrapServer="localhost:9092";
        String groupId="my-third-app";
```

```
//Create Consumer Configs  
Properties properties=new Properties();  
properties.setProperty(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG,  
bootstrapServer);  
properties.setProperty(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,  
StringDeserializer.class.getName());  
properties.setProperty(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,  
StringDeserializer.class.getName());  
properties.setProperty(ConsumerConfig.GROUP_ID_CONFIG,groupId);  
properties.setProperty(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG,  
value:"earliest"); // earliest/latest/none
```

```
//Create a Consumer  
KafkaConsumer<String, String> consumer=  
new KafkaConsumer<String, String>(properties);
```

```
//Subscribe to the topic  
  
String topic="codewithz_third_topic";  
  
/ consumer.subscribe(Collections.singleton(topic));  
  
consumer.subscribe(Arrays.asList(topic));
```

//Poll for New Data

```
while (true){  
    ConsumerRecords<String, String> records=  
    consumer.poll(Duration.ofMillis(millis:100));  
  
    for(ConsumerRecord<String, String> record:records){  
        logger.info("-----");  
        logger.info("Key:"+record.key()+" | Value:"+record.value());  
        logger.info("Partition:"+record.partition()+" | Offset:"+record.offset());  
    }  
}  
}
```